



Computersysteme

Microarchitecture

Markus Bader

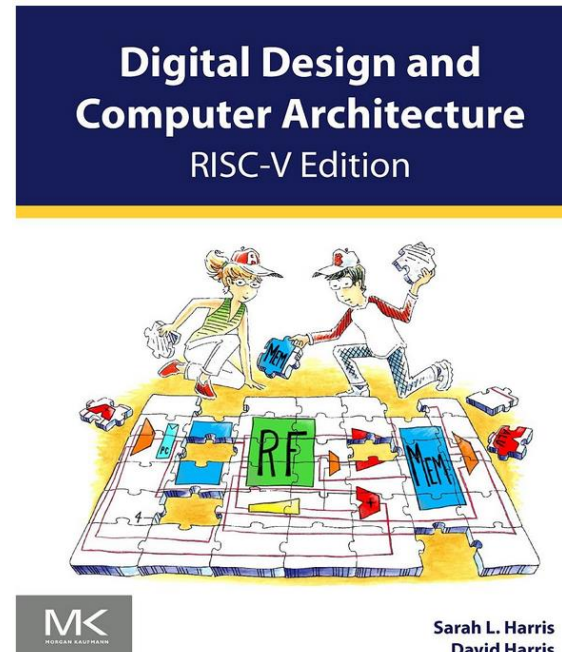
SS2025



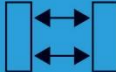
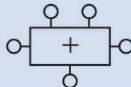

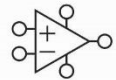


Introduction

DDCA Ch7 - Part 1: Microarchitecture Introduction <https://youtu.be/lrN-uBKooRY?si=QEiy6eyr5c32m31n>

Chapter 7 :: Topics

- Introduction
- Performance Analysis
- Single-Cycle Processor
- Multicycle Processor
- Pipelined Processor
- Advanced Microarchitecture



Application Software	<code>>"hello world!"</code>
Operating Systems	
Architecture	
Micro-architecture	
Logic	
Digital Circuits	
Analog Circuits	
Devices	
Physics	

About these Notes

Digital Design and Computer Architecture Lecture Notes

© 2021 Sarah Harris and David Harris

These notes may be used and modified for educational and/or non-commercial purposes so long as the source is attributed.

- Multiple implementations for a single architecture:
 - **Single-cycle**: Each instruction executes in a single cycle
 - **Multicycle**: Each instruction is broken up into series of shorter steps
 - **Pipelined**: Each instruction broken up into series of steps & multiple instructions execute at once

- Program execution time
Execution Time = (#instructions)(cycles/instruction)(seconds/cycle)
- Definitions:
 - **CPI**: Cycles/instruction
 - **clock period**: seconds/cycle
 - **IPC**: instructions/cycle = IPC
- Challenge is to satisfy constraints of:
 - Cost
 - Power
 - Performance

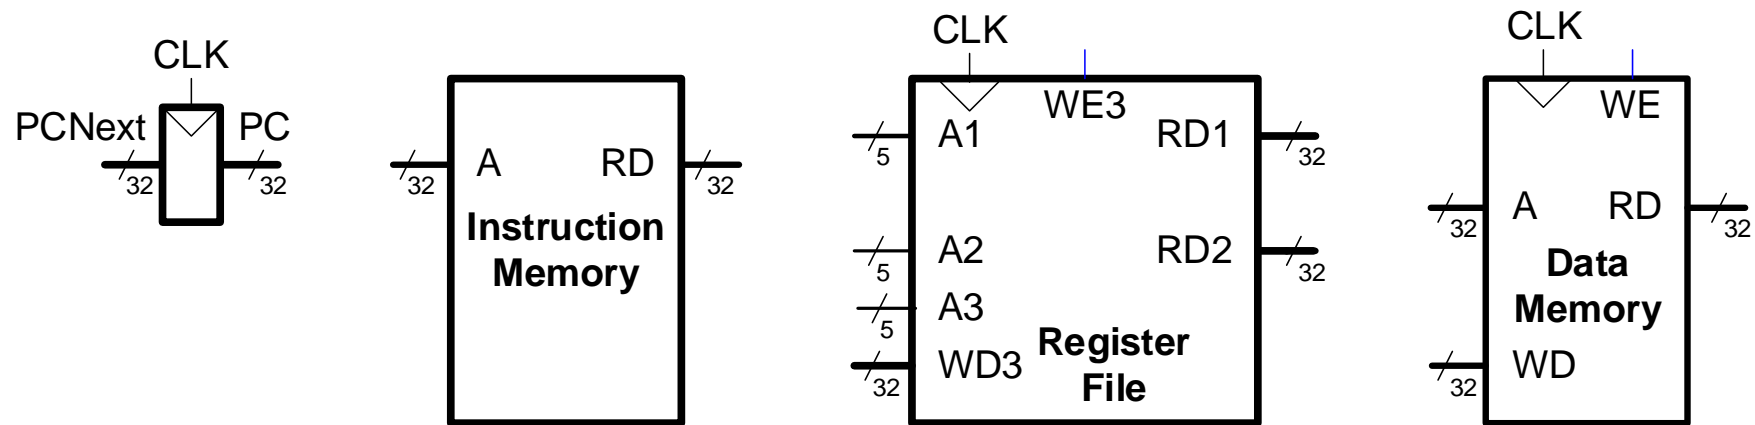
- Consider **subset** of RISC-V instructions:
- R-type ALU instructions:
add, sub, and, or, slt
- Memory instructions:
lw, sw
- Branch instructions:
beq

Architectural State Elements

Determines everything about a processor:

- **Architectural state:**

- 32 registers
- PC
- Memory



Single-Cycle RISC-V Processor

DDCA Ch7 - Part 2: RISC-V Single-Cycle Processor Datapath: lw <https://youtu.be/AoBkibslRBM?si=fgO1anrXwzMCdOrU>

Single-Cycle RISC-V Processor

- Datapath
- Control

Example Program


- Design datapath
- View example program executing

Address	Instruction	Type	Fields					Machine Language	
0x1000	L7: lw x6, -4(x9)	I	imm _{11:0}	rs1	f3	rd	op		
			1111111111100	01001	010	00110	0000011	FFC4A303	
0x1004	sw x6, 8(x9)	S	imm _{11:5}	rs2	rs1	f3	imm _{4:0}	op	
			0000000	00110	01001	010	01000	0100011	0064A423
0x1008	or x4, x5, x6	R	funct7	rs2	rs1	f3	rd	op	
			0000000	00110	00101	110	00100	0110011	0062E233
0x100C	beq x4, x4, L7	B	imm _{12,10:5}	rs2	rs1	f3	imm _{4:1,11}	op	
			1111111	00100	00100	000	10101	1100011	FE420AE3

- **Datapath:** start with `lw` instruction

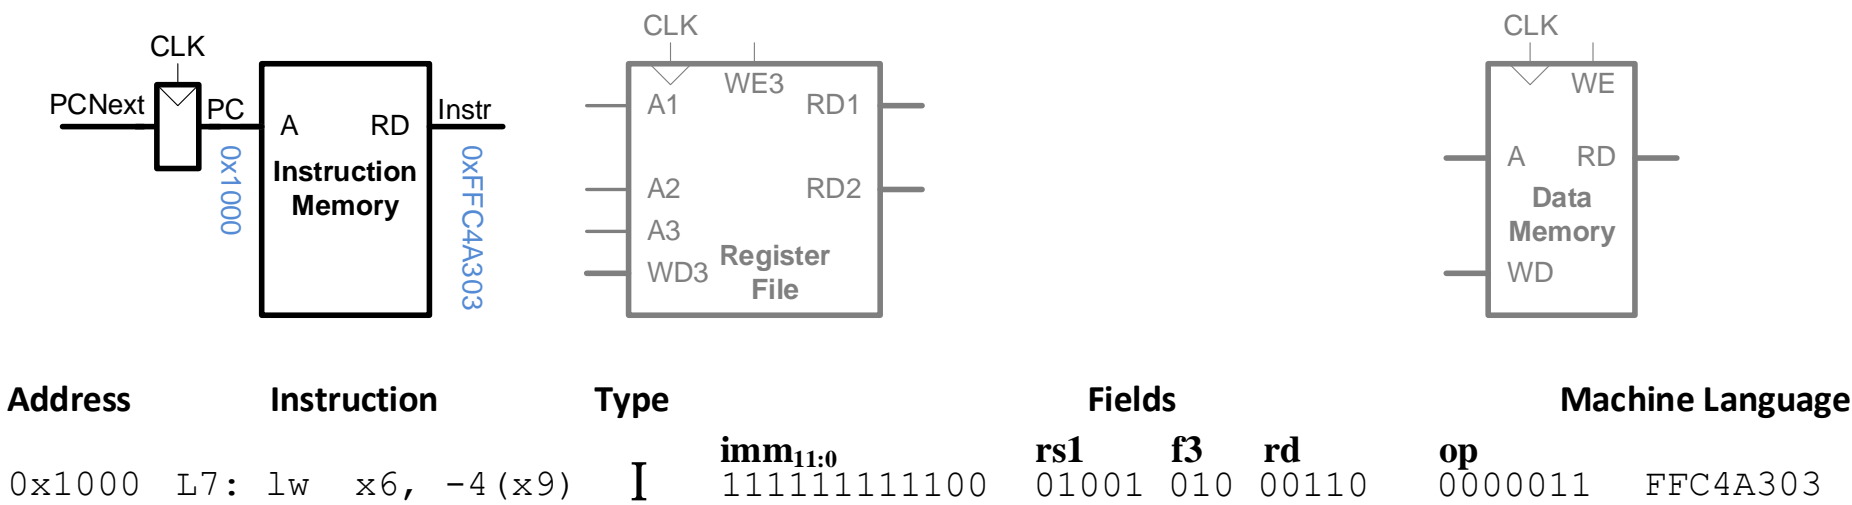
- Example:
`lw x6, -4(x9)`
`lw rd, imm(rs1)`

I-Type



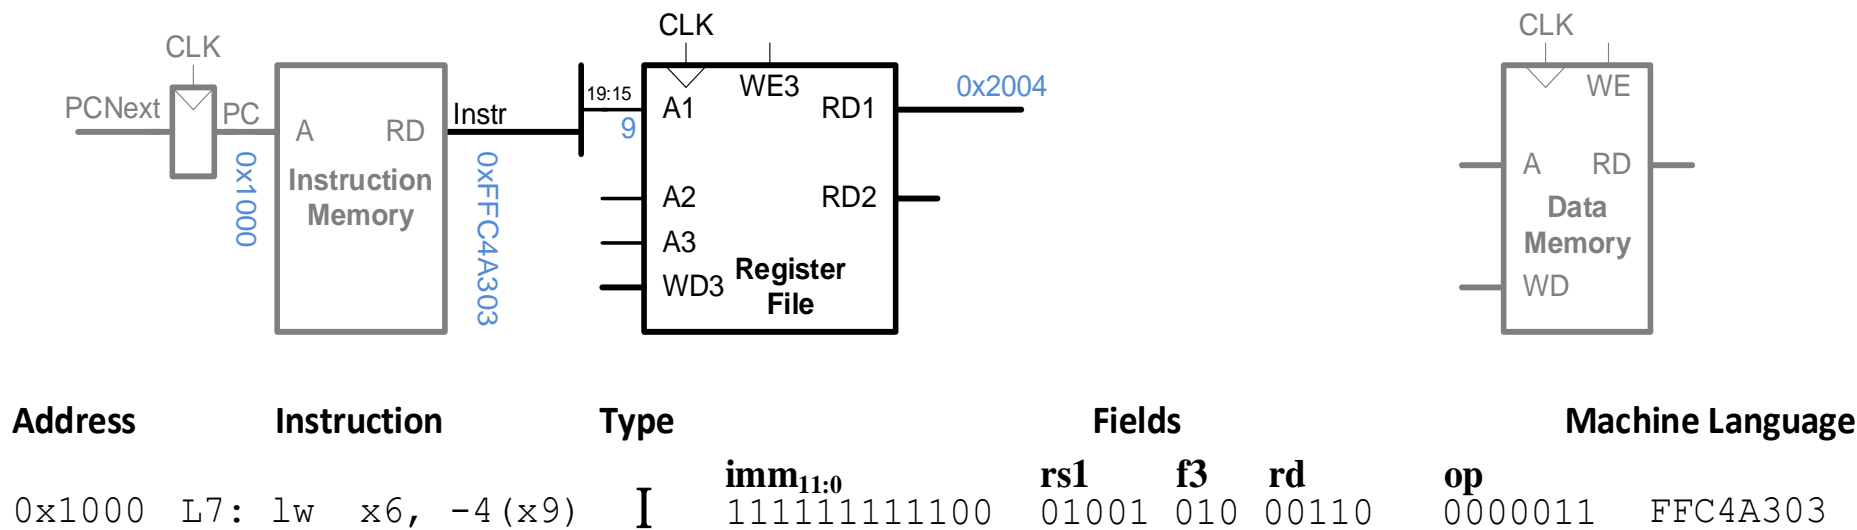
Single-Cycle Datapath: 1_w fetch

- STEP 1: Fetch instruction



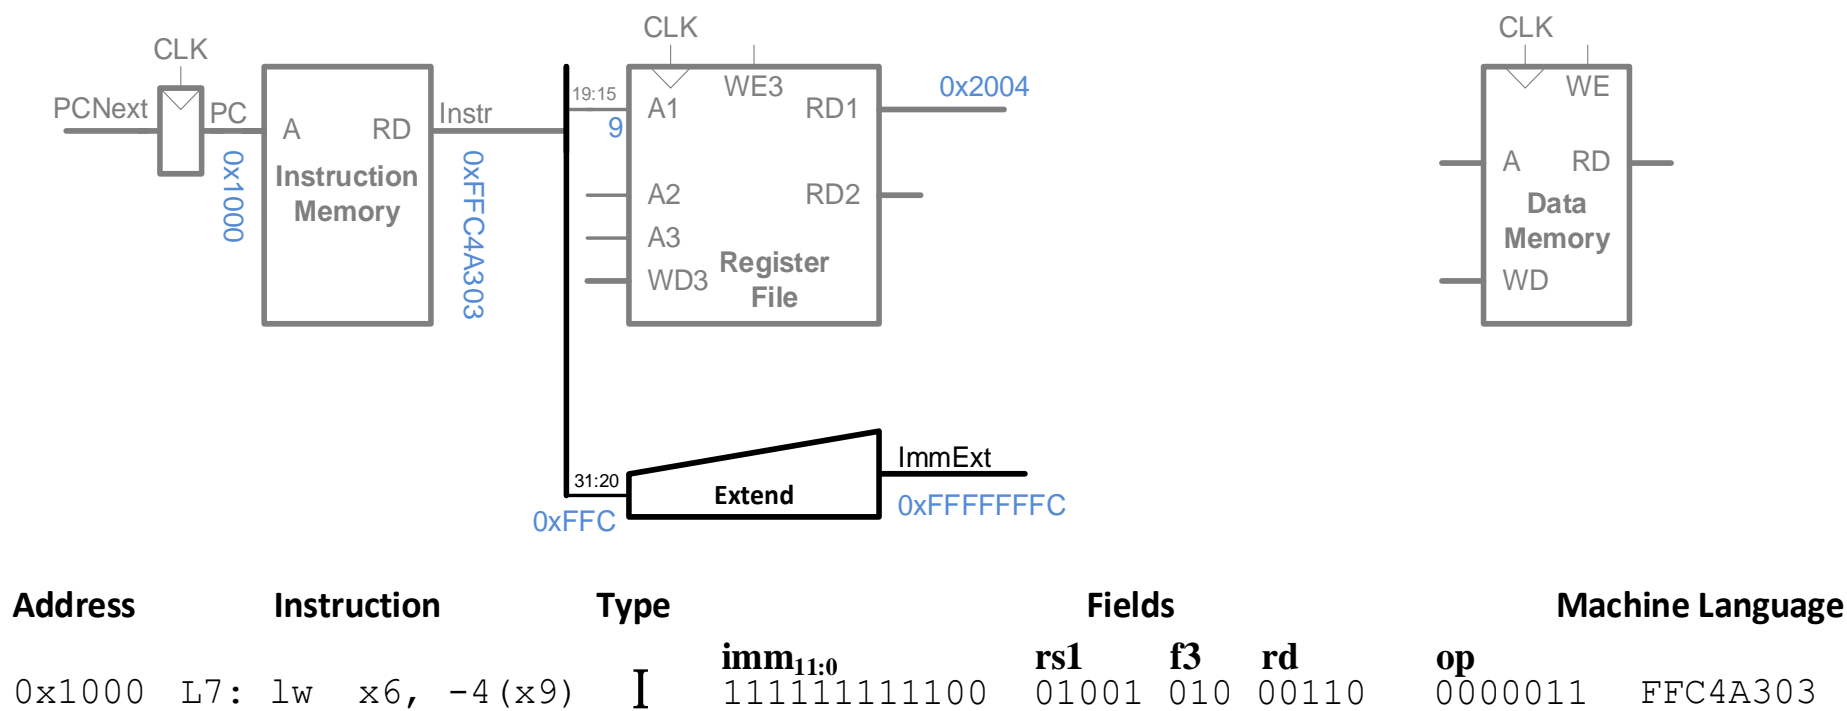
Single-Cycle Datapath: 1_w Reg Read

- **STEP 2:** Read source operand (**rs1**) from RF



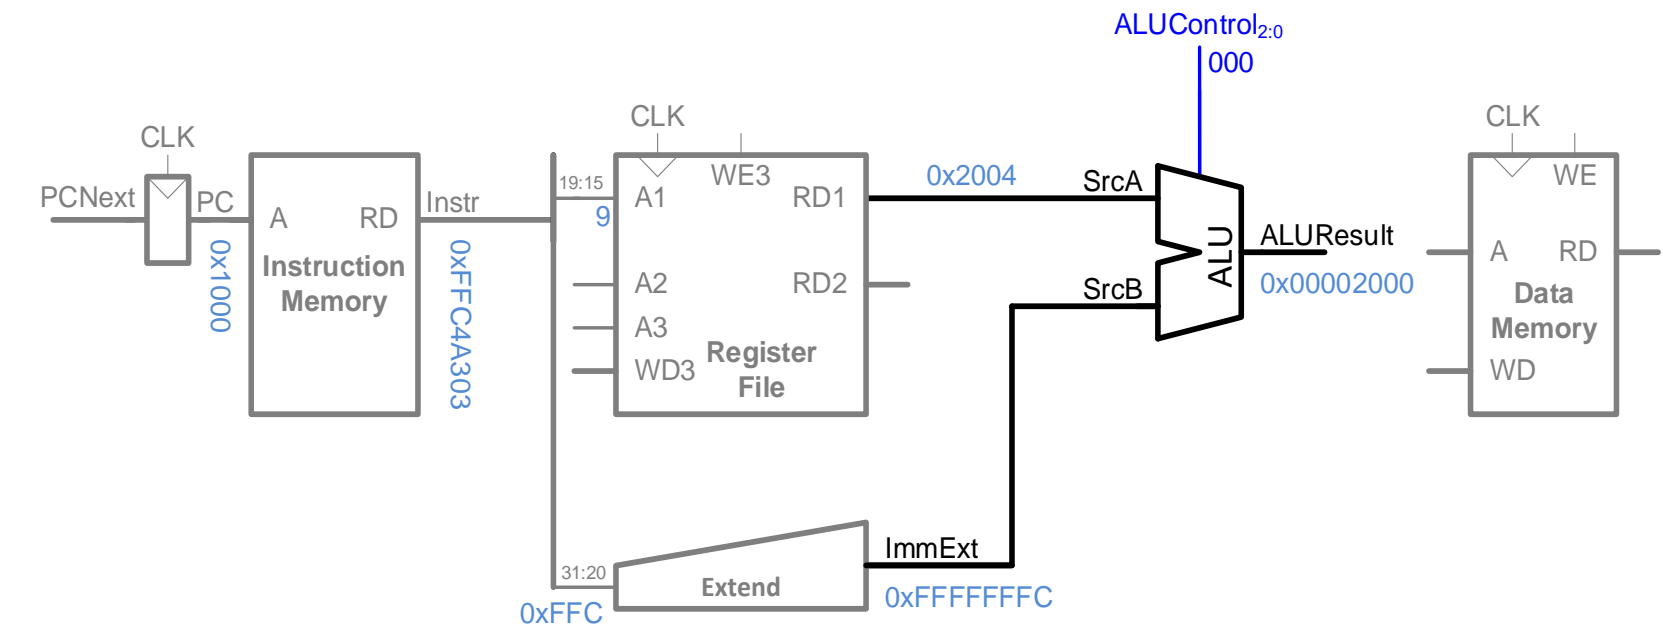
Single-Cycle Datapath: 1_w Immediate

- STEP 3: Extend the immediate



Single-Cycle Datapath: 1_w Address

- **STEP 4:** Compute the memory address

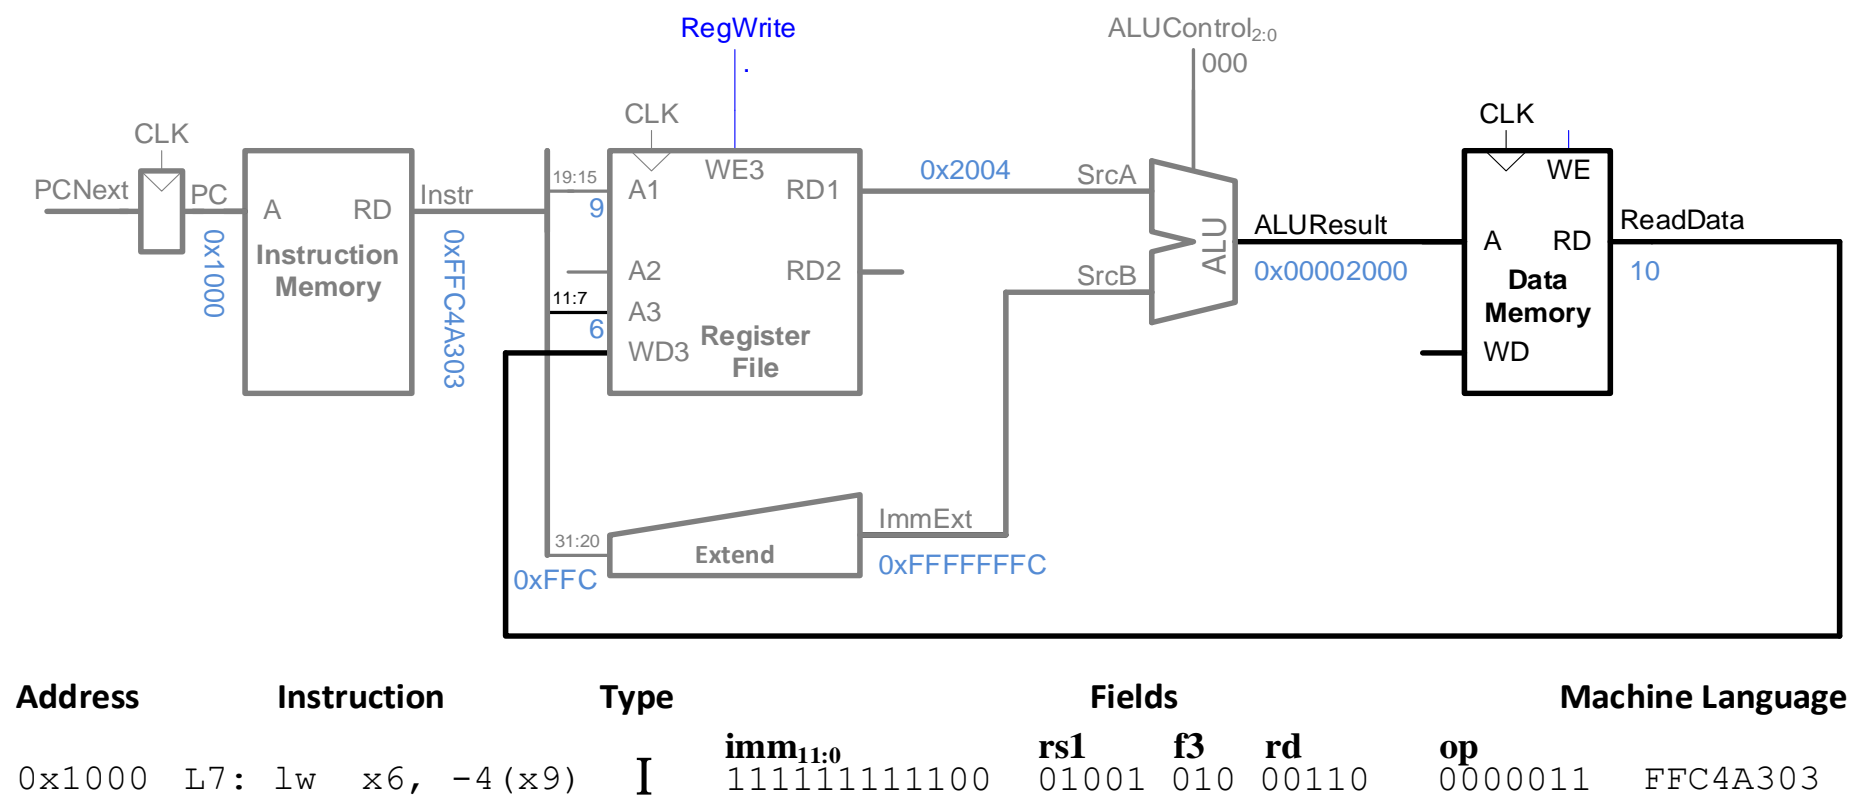


ALUControl _{2:0}	Function
000	add
001	subtract
010	and
011	or
101	SLT

Address	Instruction	Type	Fields	Machine Language										
0x1000	L7: lw x6, -4(x9)	I	<table><tr><th>imm_{11:0}</th><th>rs1</th><th>f3</th><th>rd</th><th>op</th></tr><tr><td>111111111100</td><td>01001</td><td>010</td><td>00110</td><td>0000011</td></tr></table>	imm _{11:0}	rs1	f3	rd	op	111111111100	01001	010	00110	0000011	FFC4A303
imm _{11:0}	rs1	f3	rd	op										
111111111100	01001	010	00110	0000011										

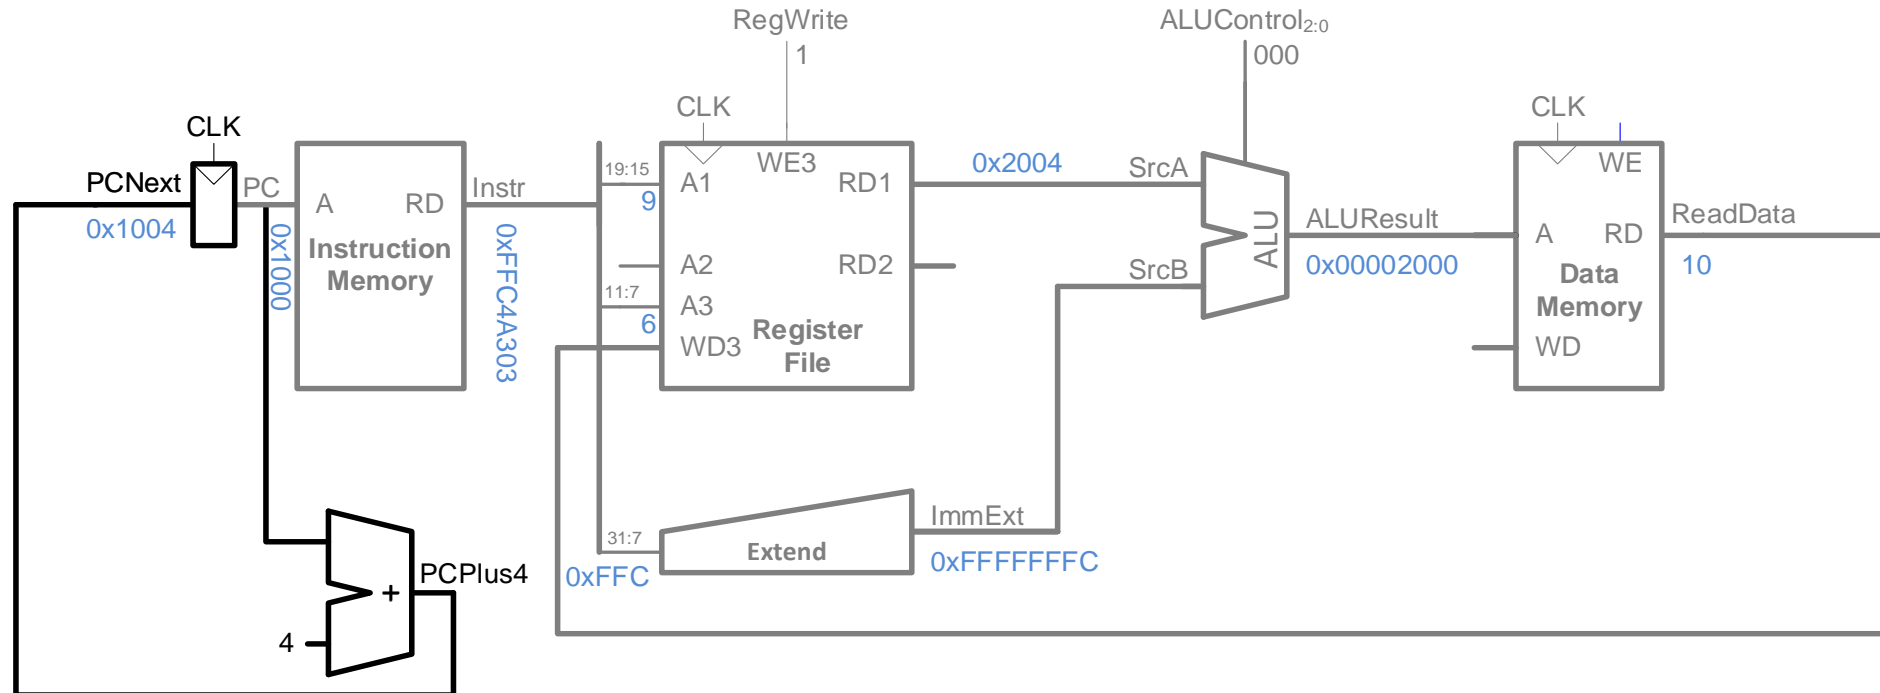
Single-Cycle Datapath: 1_w Mem Read

- **STEP 5:** Read data from memory and write it back to register file



Single-Cycle Datapath: PC Increment

- **STEP 6:** Determine address of next instruction



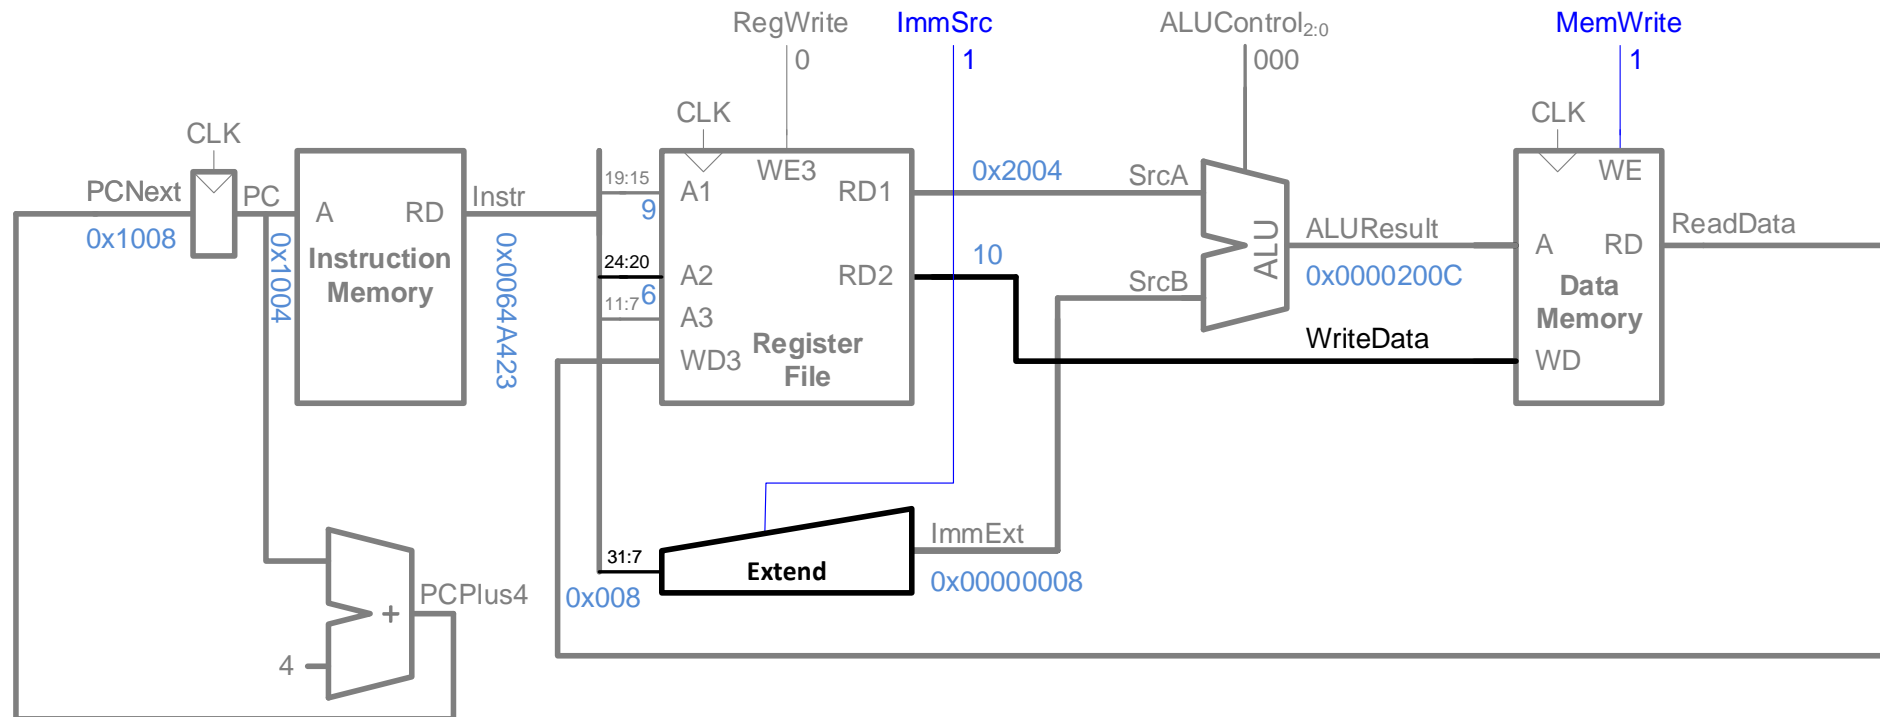
Address	Instruction			Type	Fields			Machine Language	
0x1000	L7:	lw	x6, -4(x9)	I	imm _{11:0}	rs1	f3	rd	op
					111111111100	01001	010	00110	0000011
									FFC4A303

Single-Cycle Datapath: Other Instructions

DDCA Ch7 - Part 3: RISC-V Single-Cycle Processor Datapath https://youtu.be/sVZmqLRkbVk?si=SE_gnswCekVKNuwe

Single-Cycle Datapath: sw

- **Immediate:** now in {instr[31:25], instr[11:7]}
- **Add control signals:** ImmSrc, MemWrite

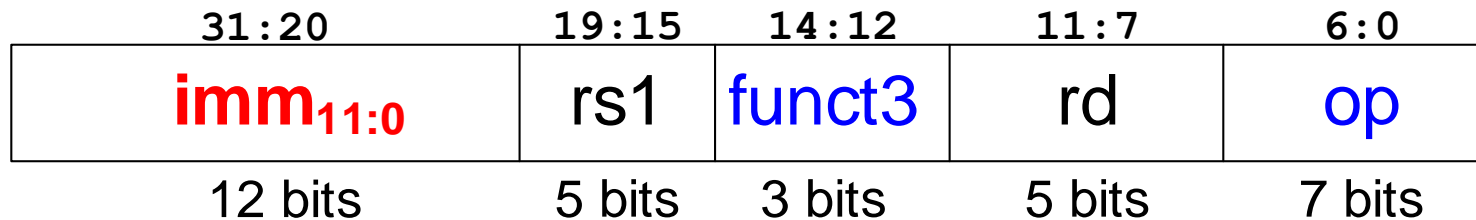


Address	Instruction	Type	Fields					Machine Language
0x1004	sw x6, 8(x9)	S	imm _{11:5}	rs2	rs1	f3	imm _{4:0}	op
0x1004	sw x6, 8(x9)	S	0000000	00110	01001	010	01000	0100011
								0064A423


Single-Cycle Datapath: Immediate

ImmSrc	ImmExt	Instruction Type
0	{{20{instr[31]}}, instr[31:20] }	I-Type
1	{{20{instr[31]}}, instr[31:25] , instr[11:7] }	S-Type

I-Type

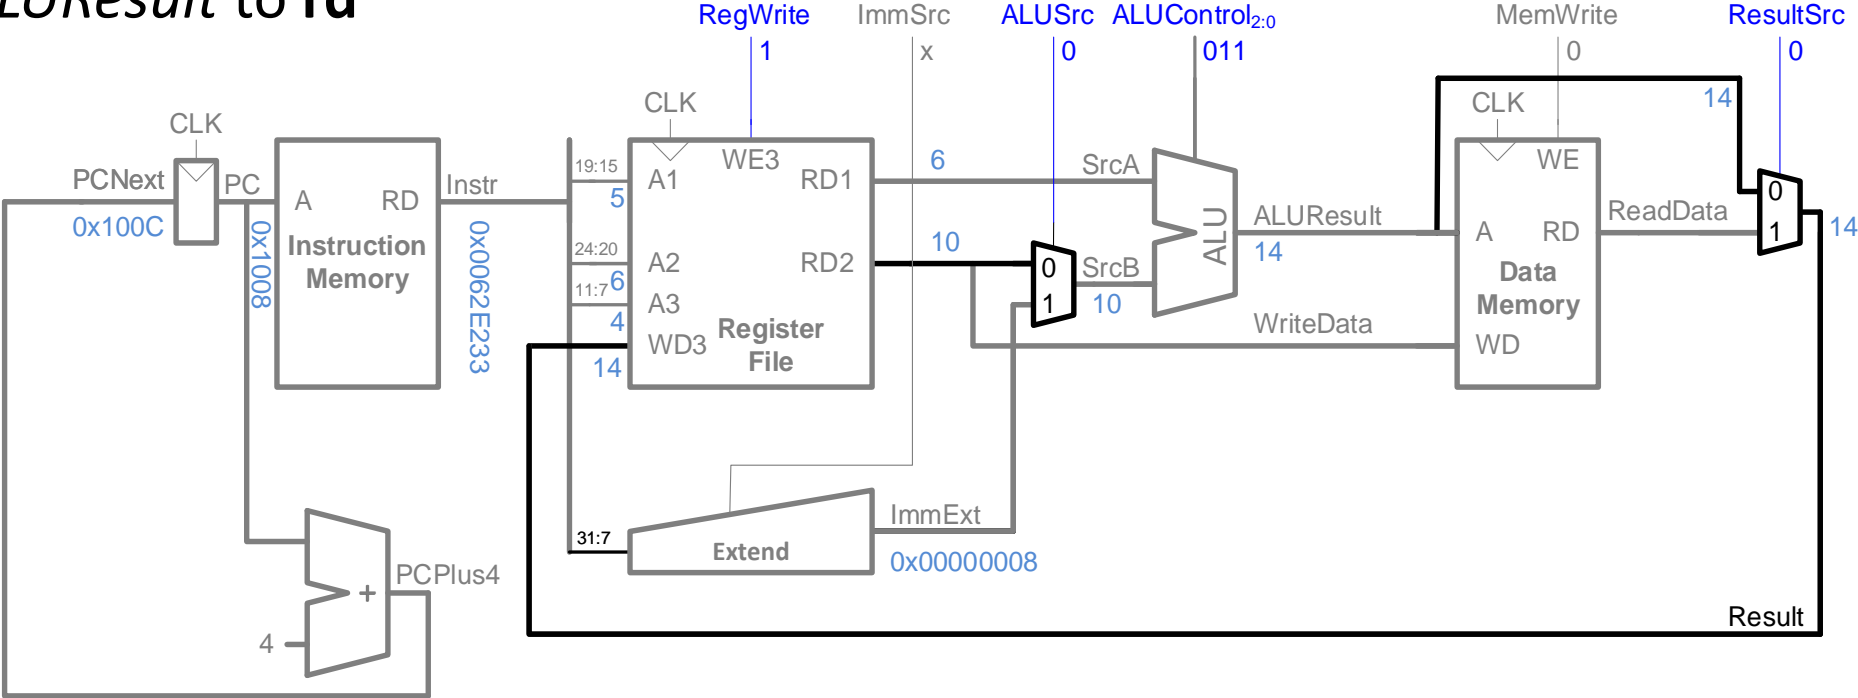


S-Type



Single-Cycle Datapath: R-type

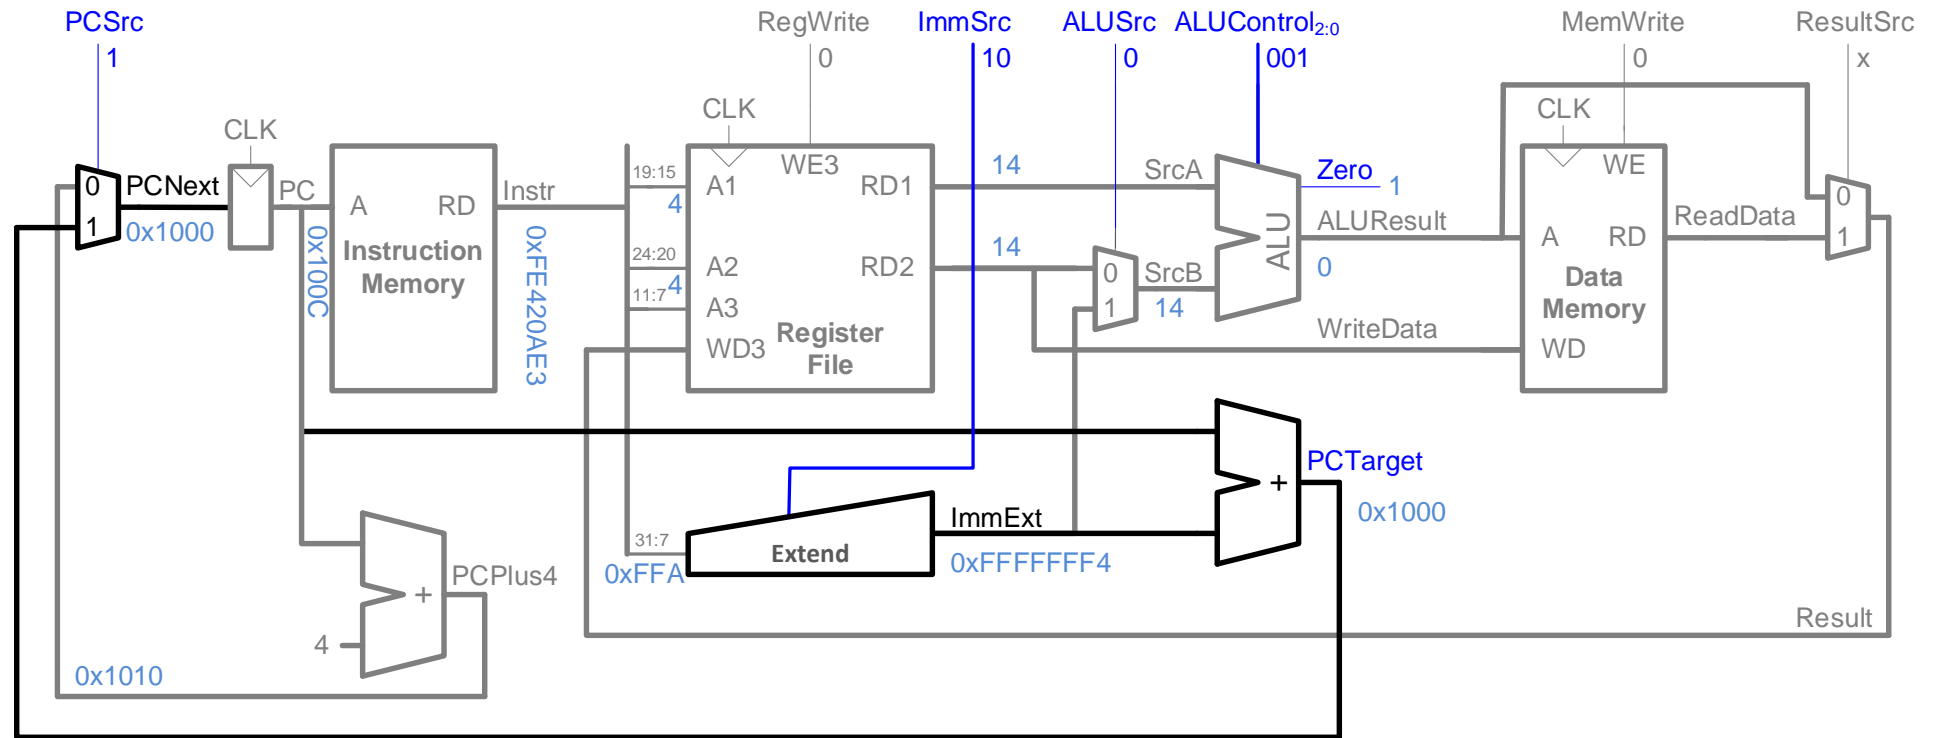
- Read from **rs1** and **rs2** (instead of **imm**)
- Write *ALUResult* to **rd**



Address	Instruction	Type	Fields					Machine Language	
0x1008	or x4, x5, x6	R	funct7	rs2	rs1	f3	rd	op	0062E233
			0000000	00110	00101	110	00100	0110011	

Single-Cycle Datapath: beq

- Calculate **target address**: $PCTarget = PC + imm$




Address	Instruction	Type	Fields						Machine Language
			imm_{12,10:5}	rs2	rs1	f3	imm_{4:1,11}	op	
0x100C	beq x4, x4, L7	B	1111111	00100	00100	000	10101	1100011	FE420AE3


Single-Cycle Datapath: ImmExt

ImmSrc _{1:0}	ImmExt	Instruction Type
00	{{20{instr[31]}}, instr[31:20]}	I-Type
01	{{20{instr[31]}}, instr[31:25], instr[11:7]}	S-Type
10	{{19{instr[31]}}, instr[31], instr[7], instr[30:25], instr[11:8], 1'b0 }	B-Type

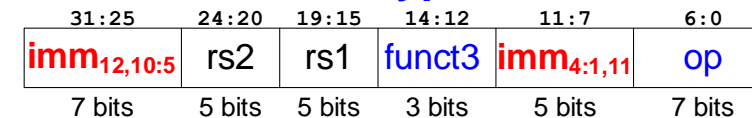
I-Type



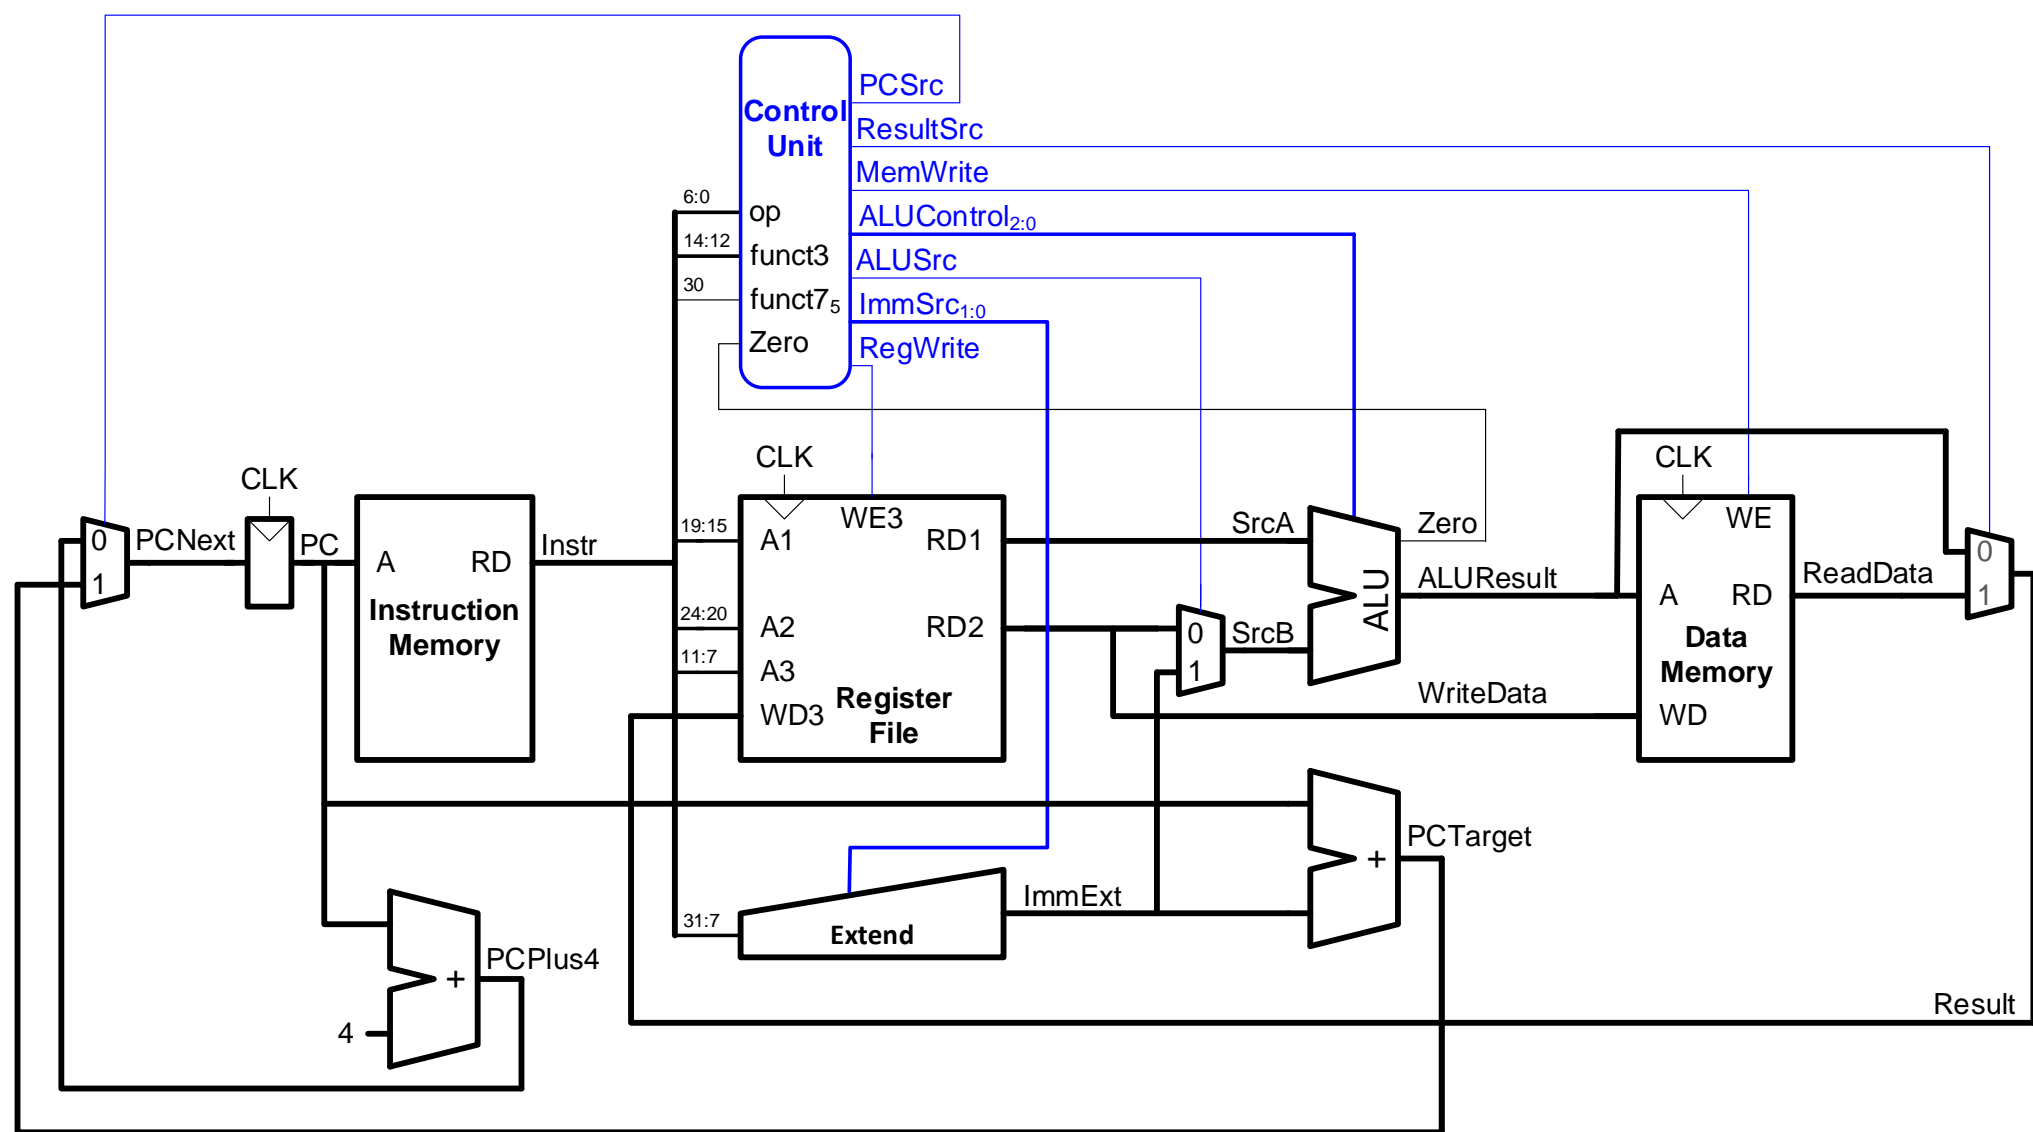
S-Type



B-Type



Single-Cycle RISC-V Processor

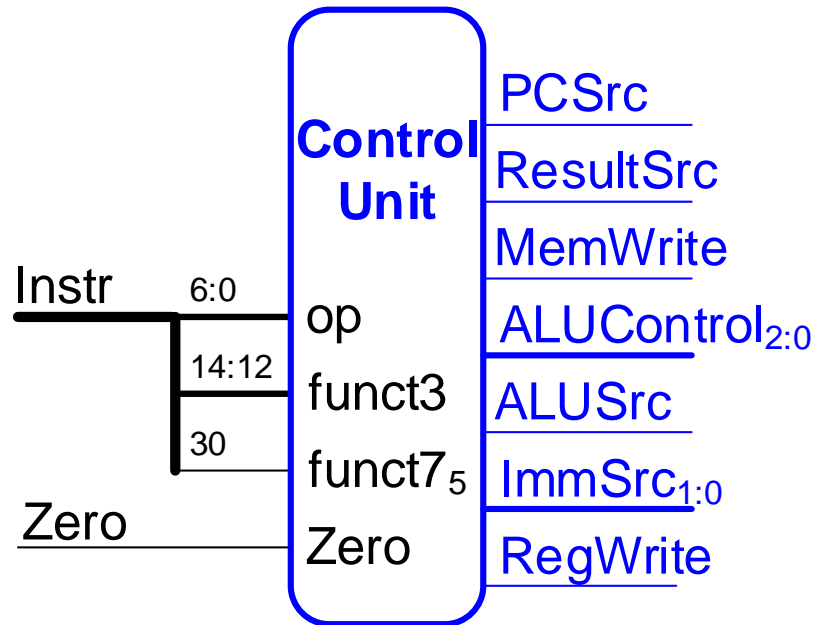


Single-Cycle Control

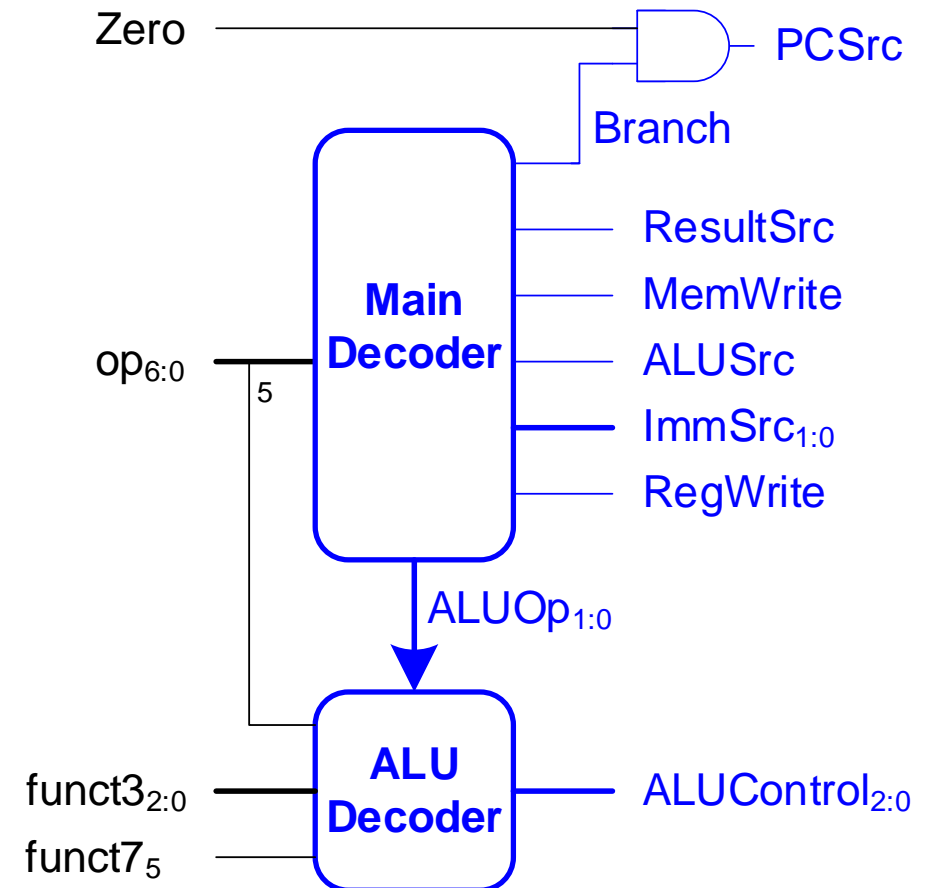
DDCA Ch7 - Part 4: RISC-V Single-Cycle Processor: Control https://www.youtube.com/watch?v=EZb1_VF-yMg

Single-Cycle Control

- High-Level View

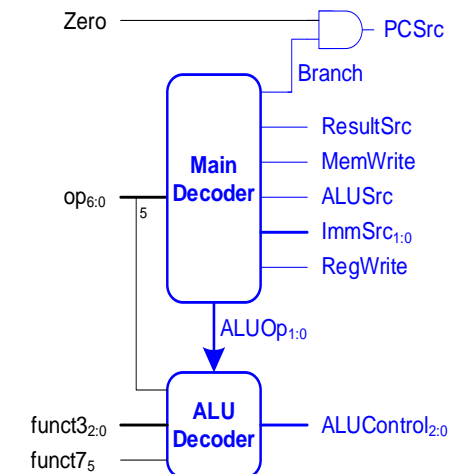
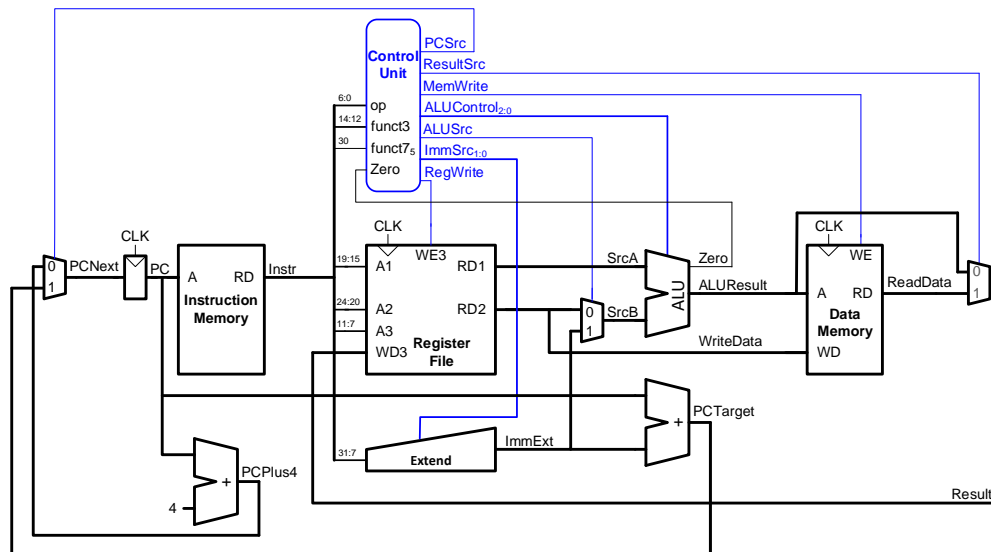


- Low-Level View



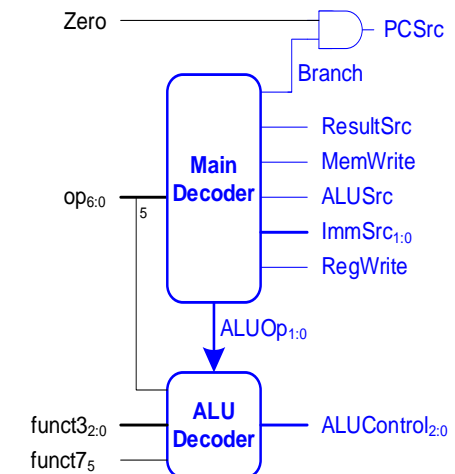
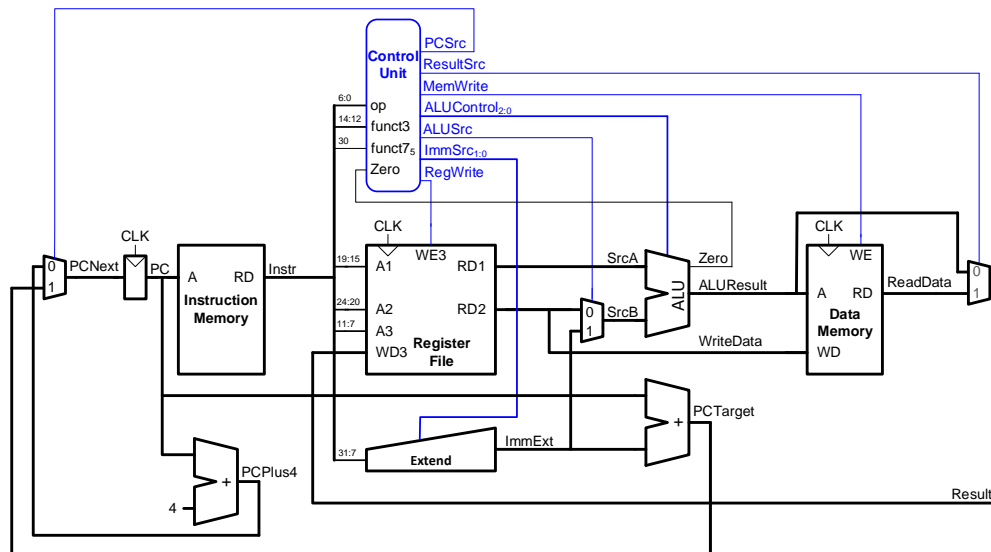
Single-Cycle Control: Main Decoder (Animation)

op	Instr.	RegWrite	ImmSrc	ALUSrc	MemWrite	ResultSrc	Branch	ALUOp
3	lw							
35	sw							
51	R-type							
99	beq							




Single-Cycle Control: Main Decoder

op	Instr.	RegWrite	ImmSrc	ALUSrc	MemWrite	ResultSrc	Branch	ALUOp
3	lw	1	00	1	0	1	0	00
35	sw	0	01	1	1	X	0	00
51	R-type	1	XX	0	0	0	0	10
99	beq	0	10	0	0	X	1	01




Review: ALU

ALUControl _{2:0}	Function
000	add
001	subtract
010	and
011	or
101	SLT




Review: ALU

ALUControl _{2:0}	Function
000	add
001	subtract
010	and
011	or
101	SLT

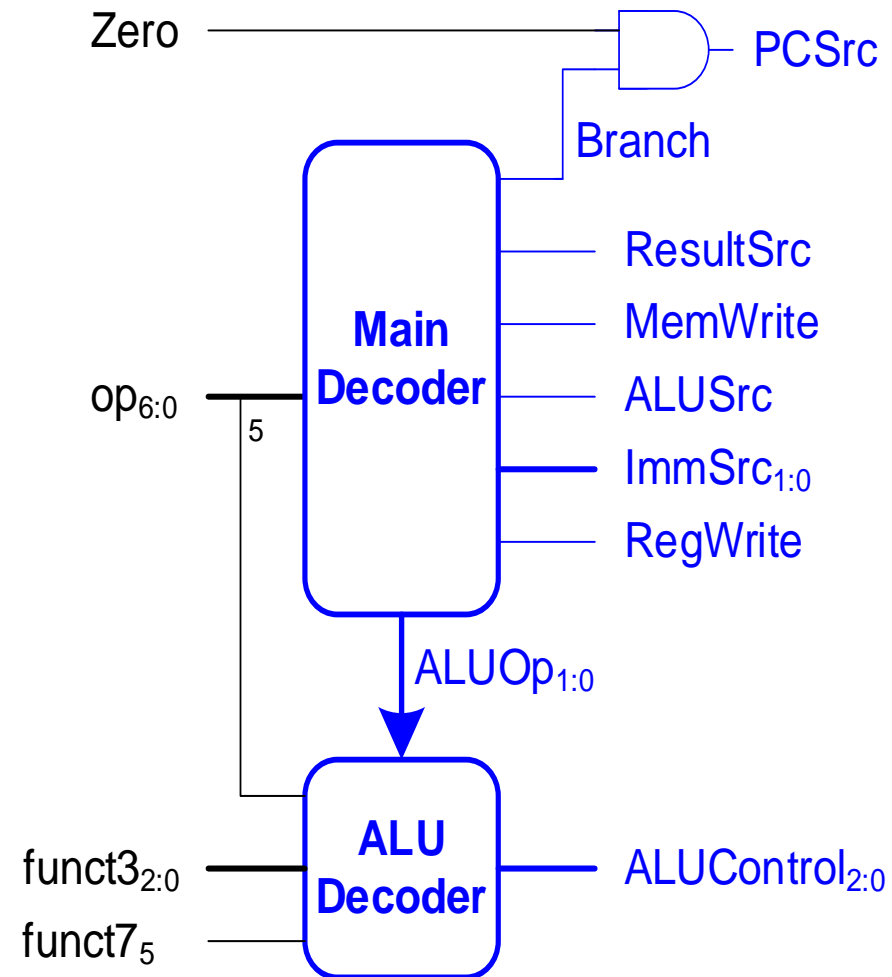


Review: ALU

ALUControl _{2:0}	Function
000	add
001	subtract
010	and
011	or
101	SLT

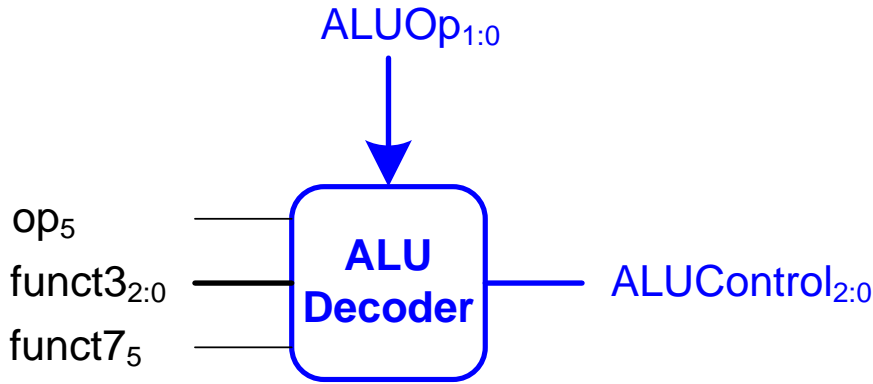


Single-Cycle Control: ALU Decoder



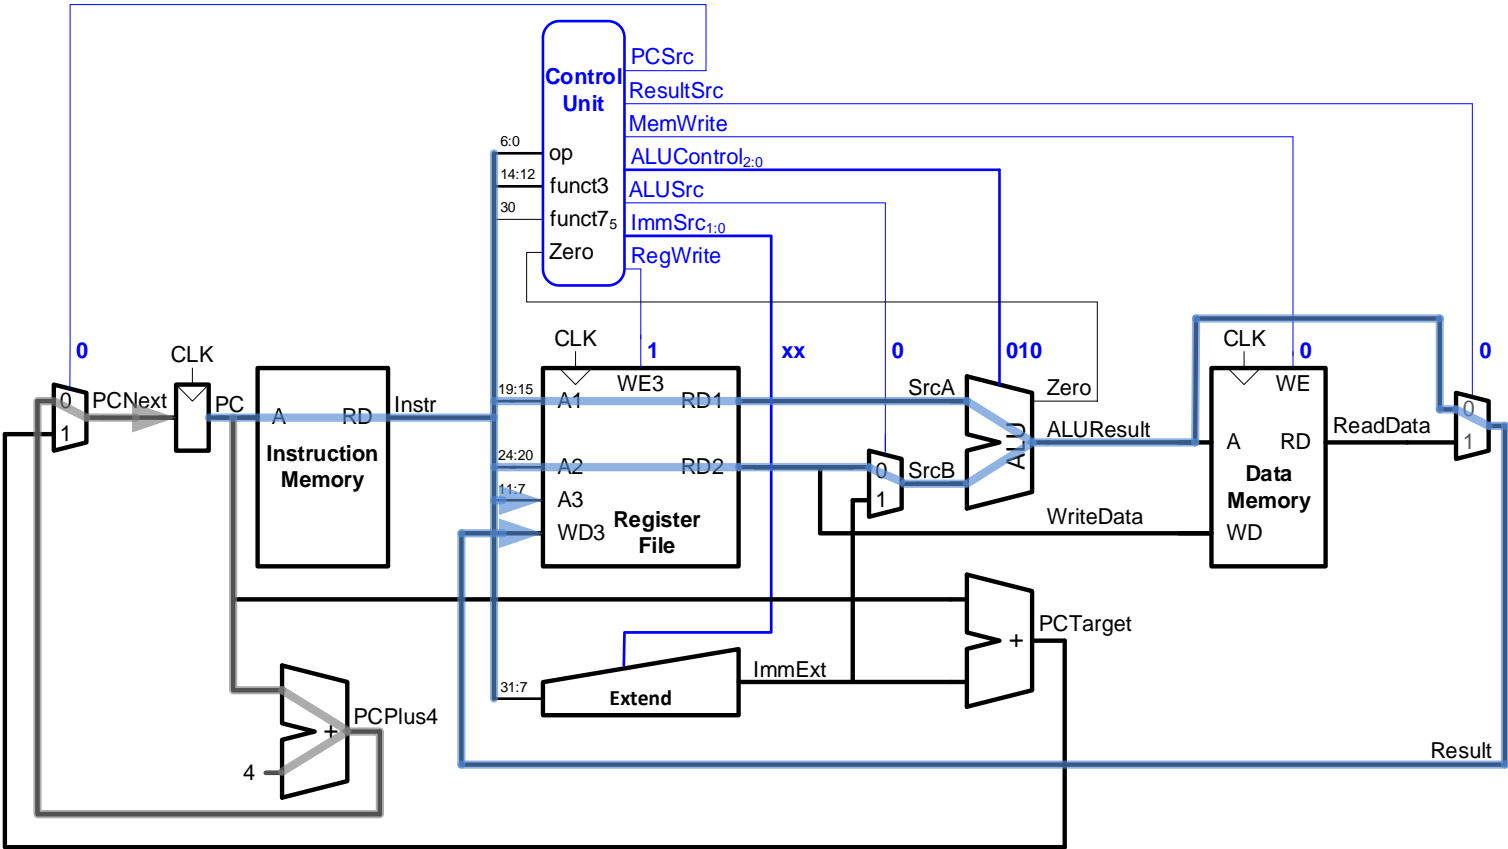
Single-Cycle Control: ALU Decoder

<i>ALUOp</i>	funct3	op ₅ , funct7 ₅	Instruction	<i>ALUControl</i> _{2:0}
00	x	x	lw, sw	000 (add)
01	x	x	beq	001 (subtract)
10	000	00, 01, 10	add	000 (add)
	000	11	sub	001 (subtract)
	010	x	slt	101 (set less than)
	110	x	or	011 (or)
	111	x	and	010 (and)



Example: and

op	Instruct	RegWrite	ImmSrc	ALUSrc	MemWrite	ResultSrc	Branch	ALUOp
51	R-type	1	XX	0	0	0	0	10



and x5, x6, x7

Computer Systems

Extending the Single-Cycle Processor

DDCA Ch7 - Part 5: RISC-V Single-Cycle Processor: https://youtu.be/z6qxMFgNEM4?si=QTFcjiic_Hq3uRfi

Extended Functionality: I-Type ALU

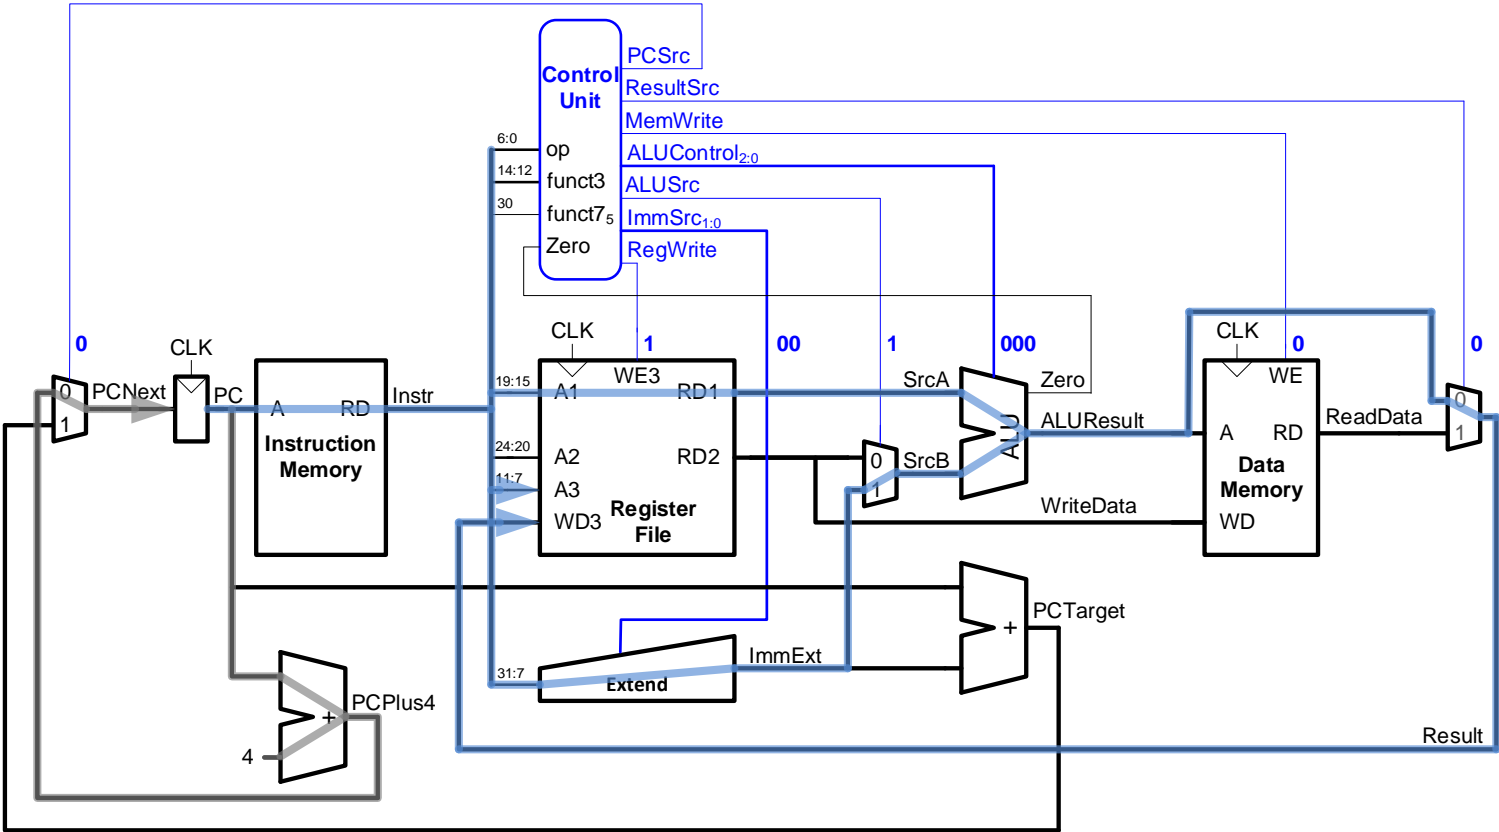
- Enhance the single-cycle processor to handle **I-Type ALU instructions:**
`addi, andi, ori, and slti`
- **Similar to R-type** instructions
- But **second source** comes from **immediate**
- Change **ALUSrc** to select the immediate
- And **ImmSrc** to pick the correct immediate

Extended Functionality: I-Type ALU

op	Instruct.	RegWrite	ImmSrc	ALUSrc	MemWrite	ResultSrc	Branch	ALUOp
3	lw	1	00	1	0	1	0	00
35	sw	0	01	1	1	X	0	00
51	R-type	1	XX	0	0	0	0	10
99	beq	0	10	0	0	X	1	01
19	I-type	1	00	1	0	0	0	10

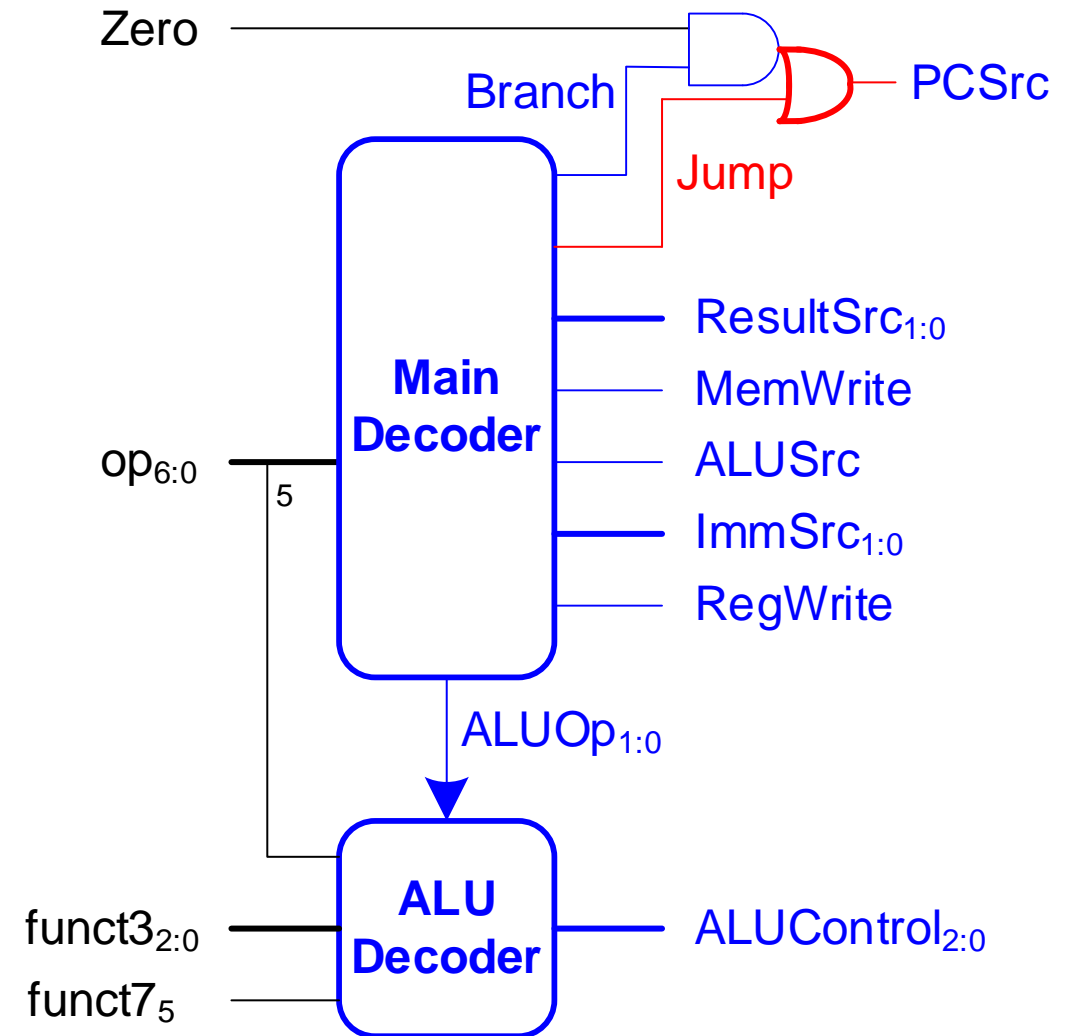
Extended Functionality: addi

op	Instruct.	RegWrite	ImmSrc	ALUSrc	MemWrite	ResultSrc	Branch	ALUOp
19	I-type	1	00	1	0	0	0	10



Extended Functionality: jal


- Enhance the single-cycle processor to handle jal
 - Similar to beq
- But jump is **always** taken
 - *PCSrc* should be 1
- **Immediate format** is different
 - Need a new *ImmSrc* of 11
- And jal must **compute PC+4** and store in rd
 - Take PC+4 from adder through *ResultMux*




Extended Functionality: ImmExt

ImmSrc _{1:0}	ImmExt	Instruction Type
00	{{20{instr[31]}}, instr[31:20]}	I-Type
01	{{20{instr[31]}}, instr[31:25], instr[11:7]}	S-Type
10	{{19{instr[31]}}, instr[31], instr[7], instr[30:25], instr[11:8], 1'b0}	B-Type
11	{{12{instr[31]}}, instr[19:12], instr[20], instr[30:21], 1'b0}	J-Type


I-Type




B-Type



S-Type



J-Type

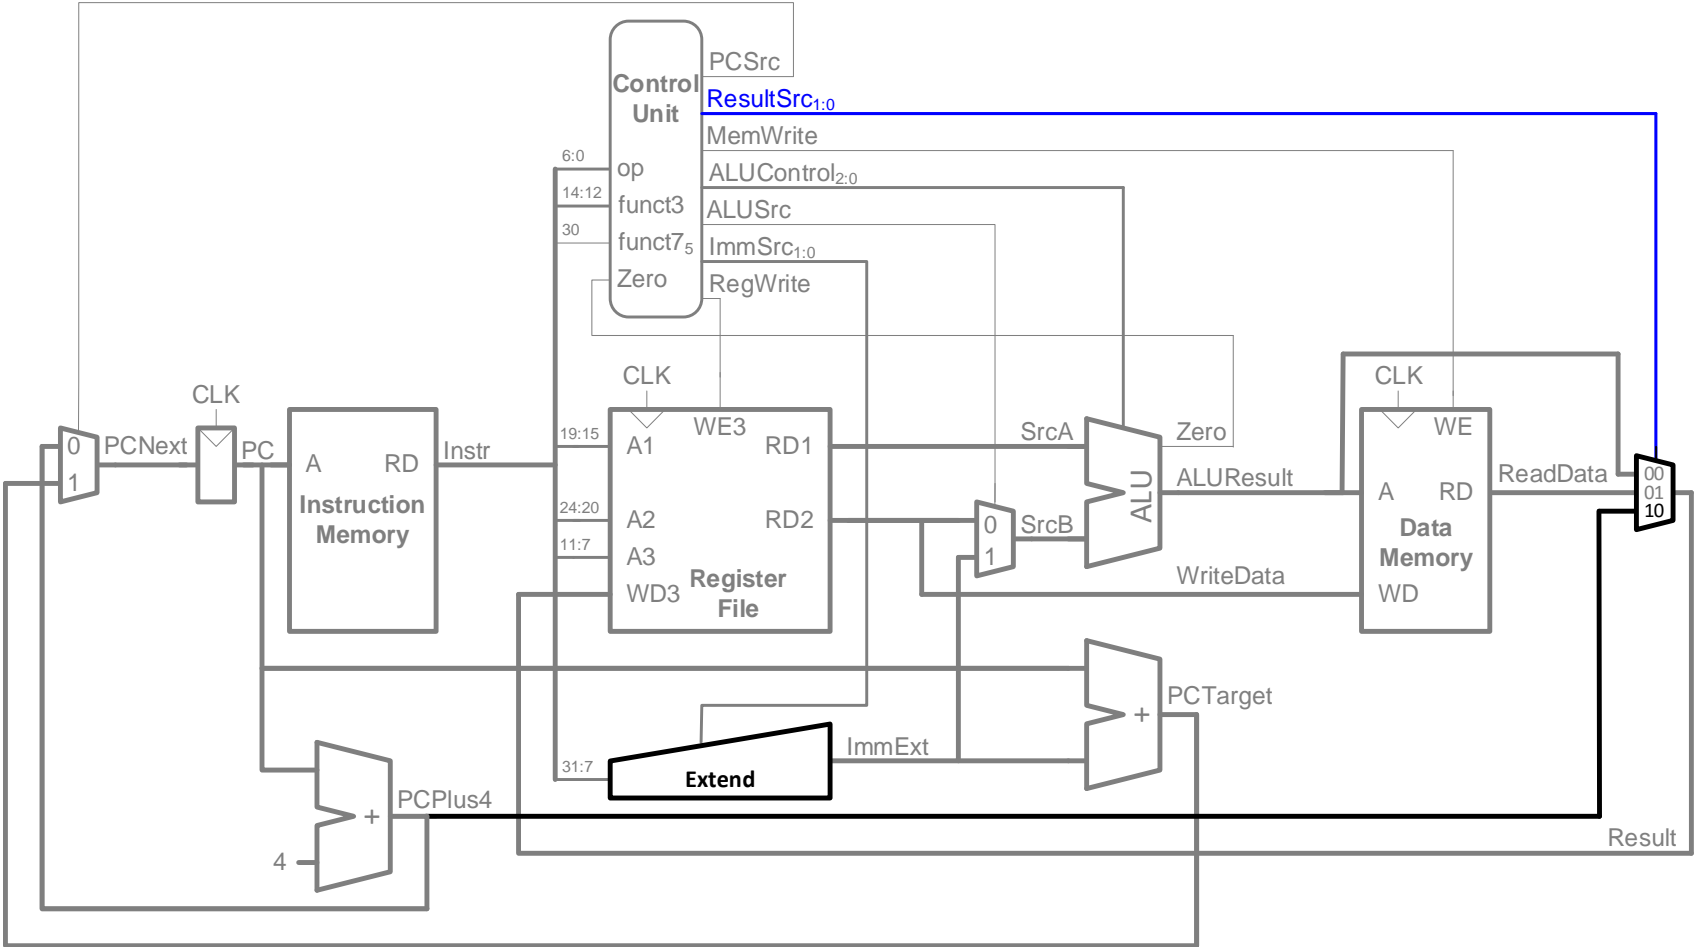


Extended Functionality: jal

op	Instruct.	RegWrite	ImmSrc	ALUSrc	MemWrite	ResultSrc	Branch	ALUOp	Jump
3	lw	1	00	1	0	01	0	00	0
35	sw	0	01	1	1	XX	0	00	0
51	R-type	1	XX	0	0	00	0	10	0
99	beq	0	10	0	0	XX	1	01	0
19	l-type	1	00	1	0	00	0	10	0
111	jal	1	11	X	0	10	0	XX	1

Extended Functionality: jal

op	Instruct.	RegWrite	ImmSrc	ALUSrc	MemWrite	ResultSrc	Branch	ALUOp	Jump
111	jal	1	11	X	0	10	0	XX	1



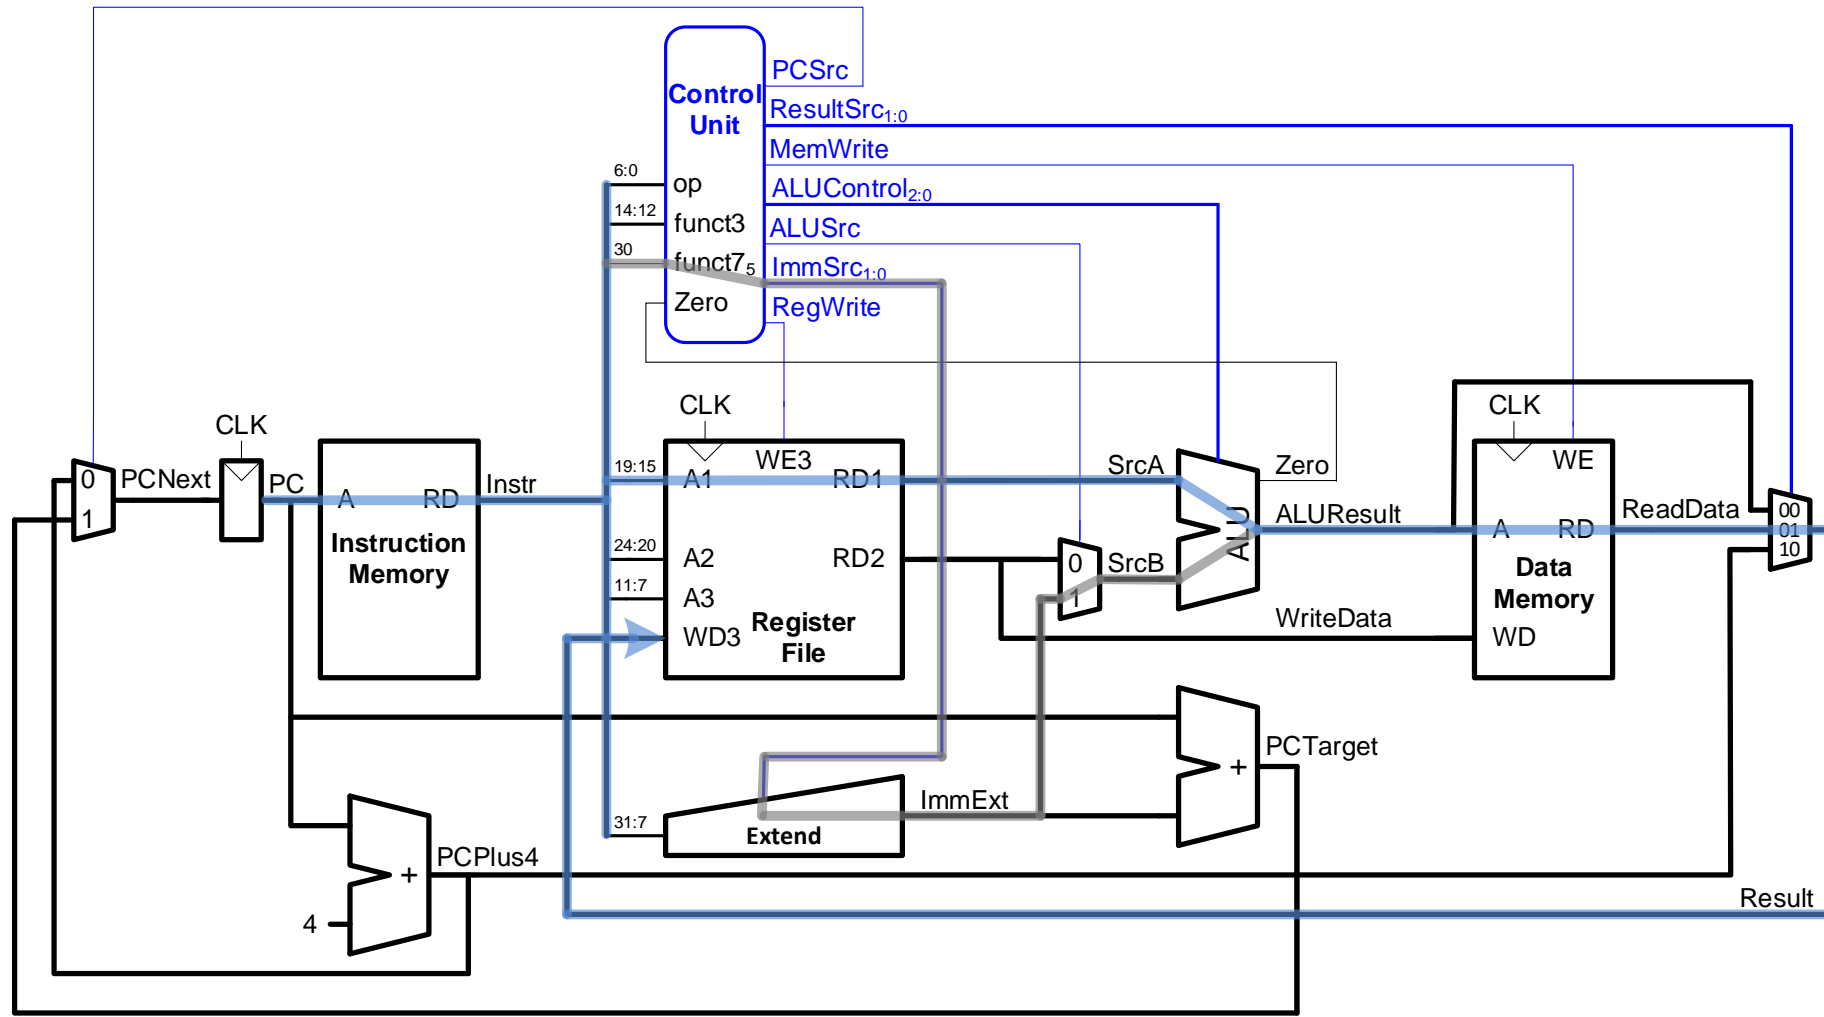
Single-Cycle Performance

DDCA Ch7 - Part 6: RISC-V Single-Cycle Performance <https://www.youtube.com/watch?v=w82mNGranjA>

Program Execution Time

$$\begin{aligned} &= \#instructions \times \frac{\text{cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{cycle}} \\ &= \#instructions \times CPI \times T_c \end{aligned}$$

Single-Cycle Processor Performance

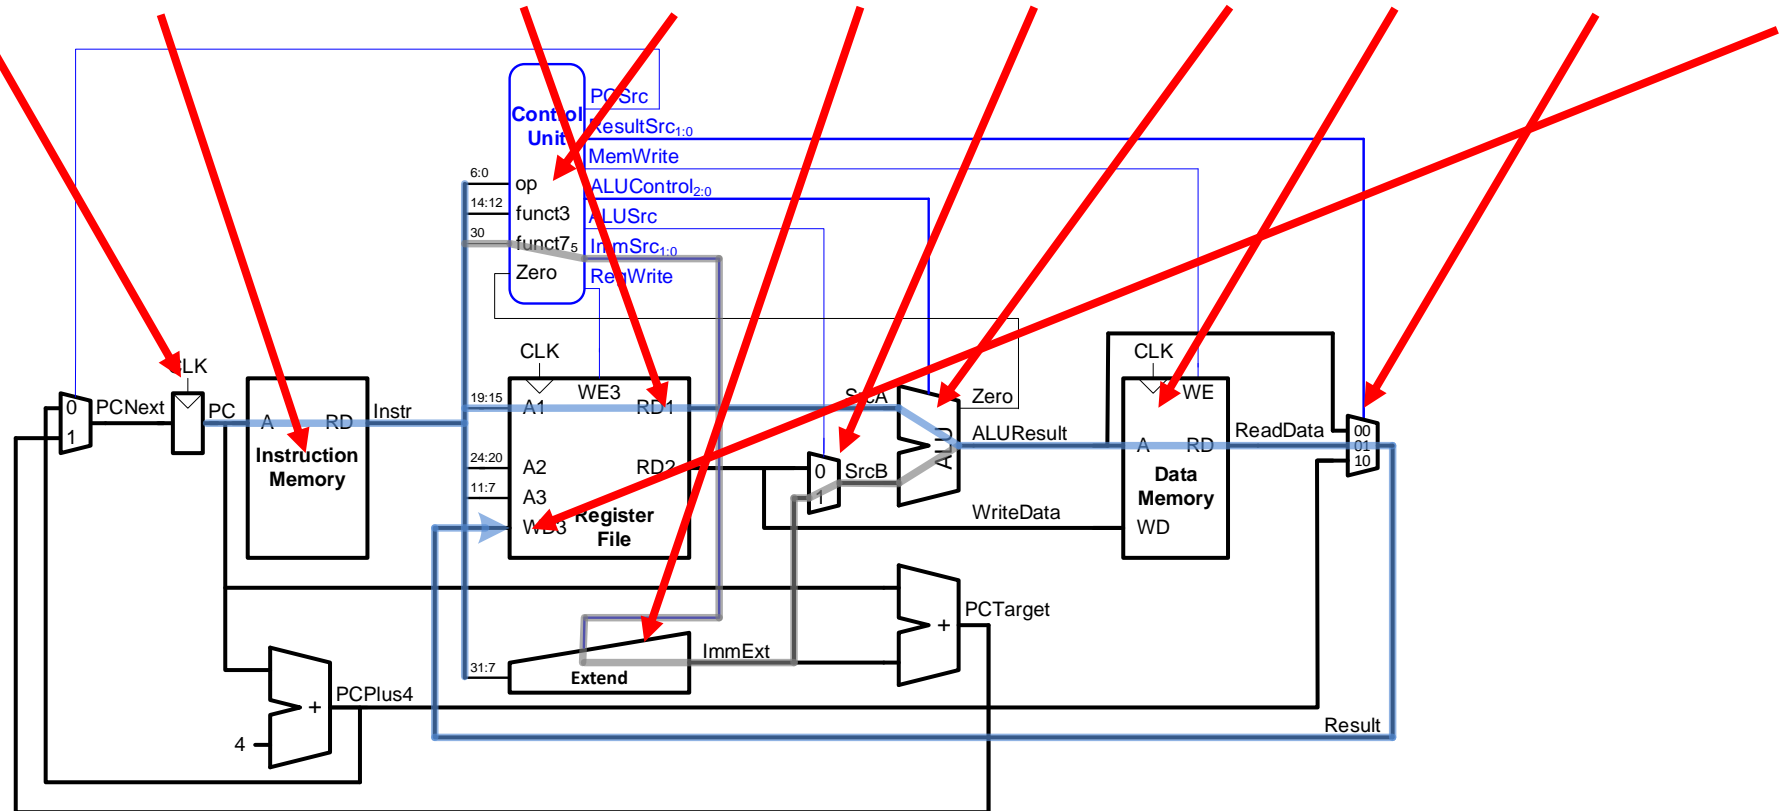


T_c limited by critical path (**1w**)

Single-Cycle Processor Performance

- Single-cycle critical path:

$$T_{c_single} = t_{pcq_PC} + t_{mem} + \max[t_{RFread}, t_{dec} + t_{ext} + t_{mux}] + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$$



- **Single-cycle critical path:**

$$T_{c_single} = t_{pcq_PC} + t_{mem} + \max [t_{RFread}, t_{dec} + t_{ext} + t_{mux}] + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$$

- **Typically, limiting paths are:**

- memory, ALU, register file
- So, $t_{dec} + t_{ext} + t_{mux}$ can be neglected

$$T_{c_single} = t_{pcq_PC} + t_{mem} + t_{RFread} + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$$

$$T_{c_single} = t_{pcq_PC} + 2t_{mem} + t_{RFread} + t_{ALU} + t_{mux} + t_{RFsetup}$$

Single-Cycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	t_{pcq_PC}	40
Register setup	t_{setup}	50
Multiplexer	t_{mux}	30
AND-OR gate	t_{AND-OR}	20
ALU	t_{ALU}	120
Decoder (Control Unit)	t_{dec}	25
Extend unit	t_{ext}	35
Memory read	t_{mem}	200
Register file read	t_{RFread}	100
Register file setup	$t_{RFsetup}$	60

$$\begin{aligned} T_{c_single} &= t_{pcq_PC} + 2t_{mem} + t_{RFread} + t_{ALU} + t_{mux} + t_{RFsetup} \\ &= (40 + 2*200 + 100 + 120 + 30 + 60) \text{ ps} = \mathbf{750 \text{ ps}} \end{aligned}$$

Single-Cycle Performance Example

Program with 100 billion instructions = 10^{11} instructions :

$$\begin{aligned}\text{Execution Time} &= \#instructions \times CPI \times T_c \\ &= (100 \times 10^9)(1)(750 \times 10^{-12} \text{ s}) \\ &= \mathbf{75 \text{ seconds}}\end{aligned}$$

Multicycle RISC-V Processor

DDCA Ch7 - Part 7: Multicycle Processor <https://www.youtube.com/watch?v=sATaQNCC0-g>

- **Single-cycle:**
 - + simple
 - - cycle time limited by longest instruction (l_w)
 - - separate memories for instruction and data
 - - 3 adders/ALUs
- **Multicycle processor** addresses these issues by breaking instruction into shorter steps
 - shorter instructions take **fewer steps**
 - can re-use hardware
 - cycle time is faster

Single- vs. Multicycle Processor

- **Single-cycle:**

- + simple
- - cycle time limited by longest instruction (l_w)
- - separate memories for instruction and data
- - 3 adders/ALUs

- **Multicycle**

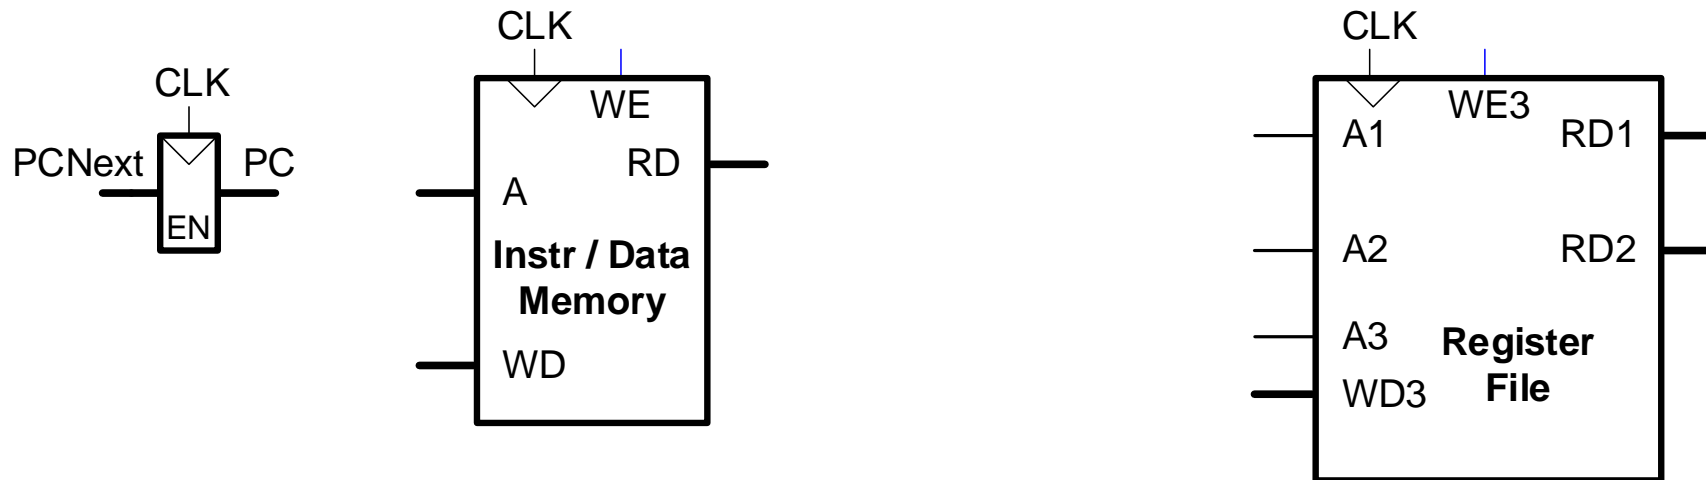
- + higher clock speed
- + simpler instructions run faster
- + reuse expensive hardware on multiple cycles
- - sequencing overhead paid many times

**Same design steps
as single-cycle:**

- **first datapath**
- **then control**

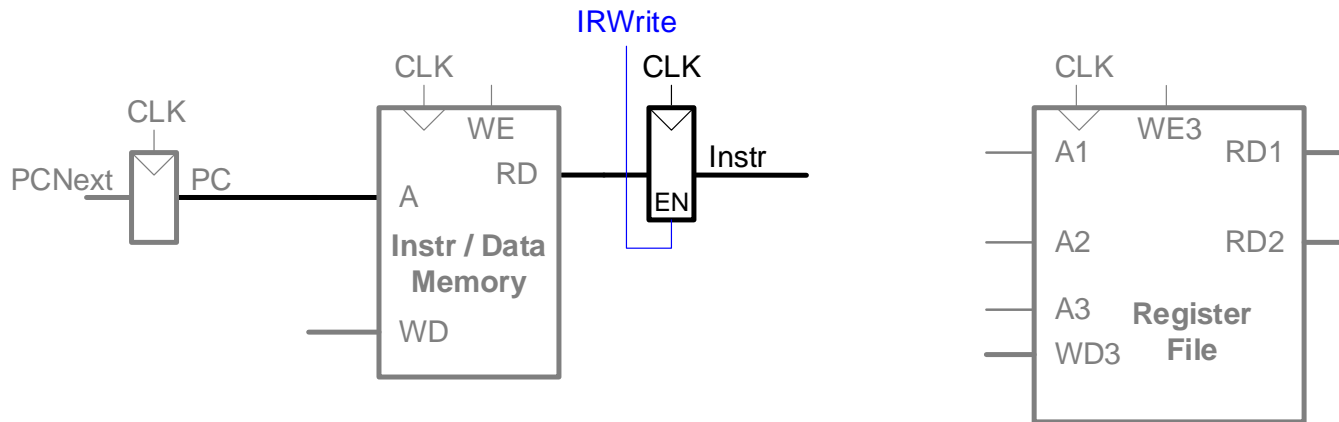
Multicycle State Elements

- Replace separate Instruction and Data memories with a **single unified memory** – more realistic



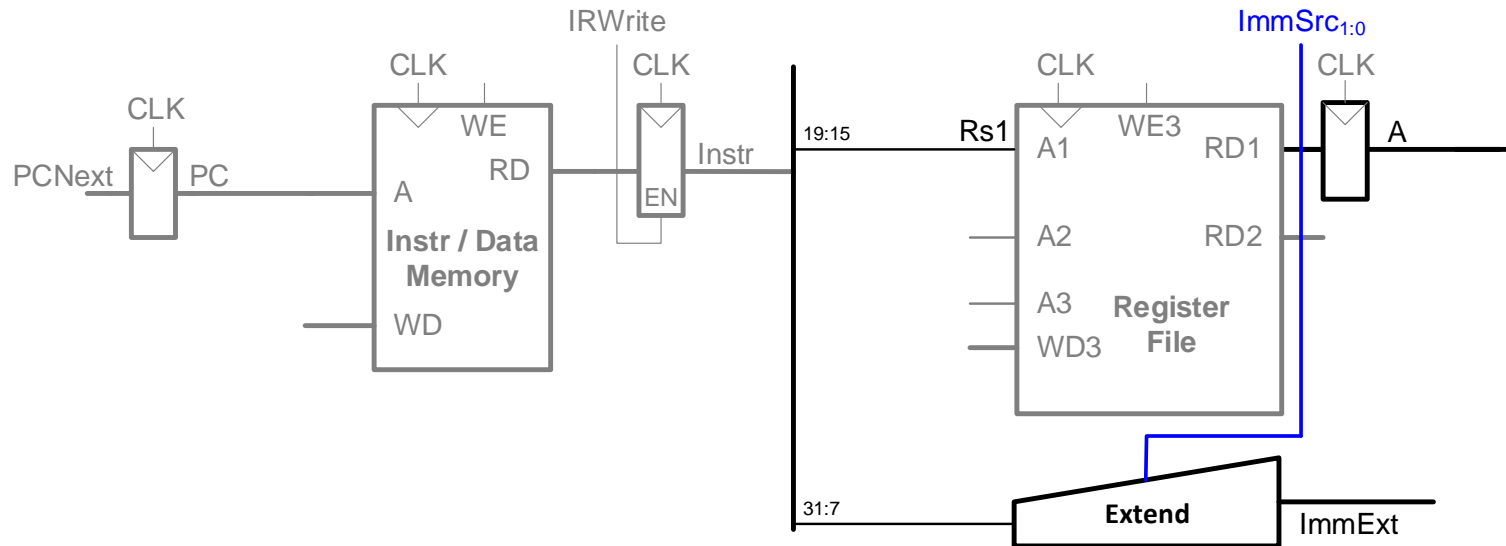
Multicycle Datapath: Instruction Fetch

STEP 1: Fetch instruction



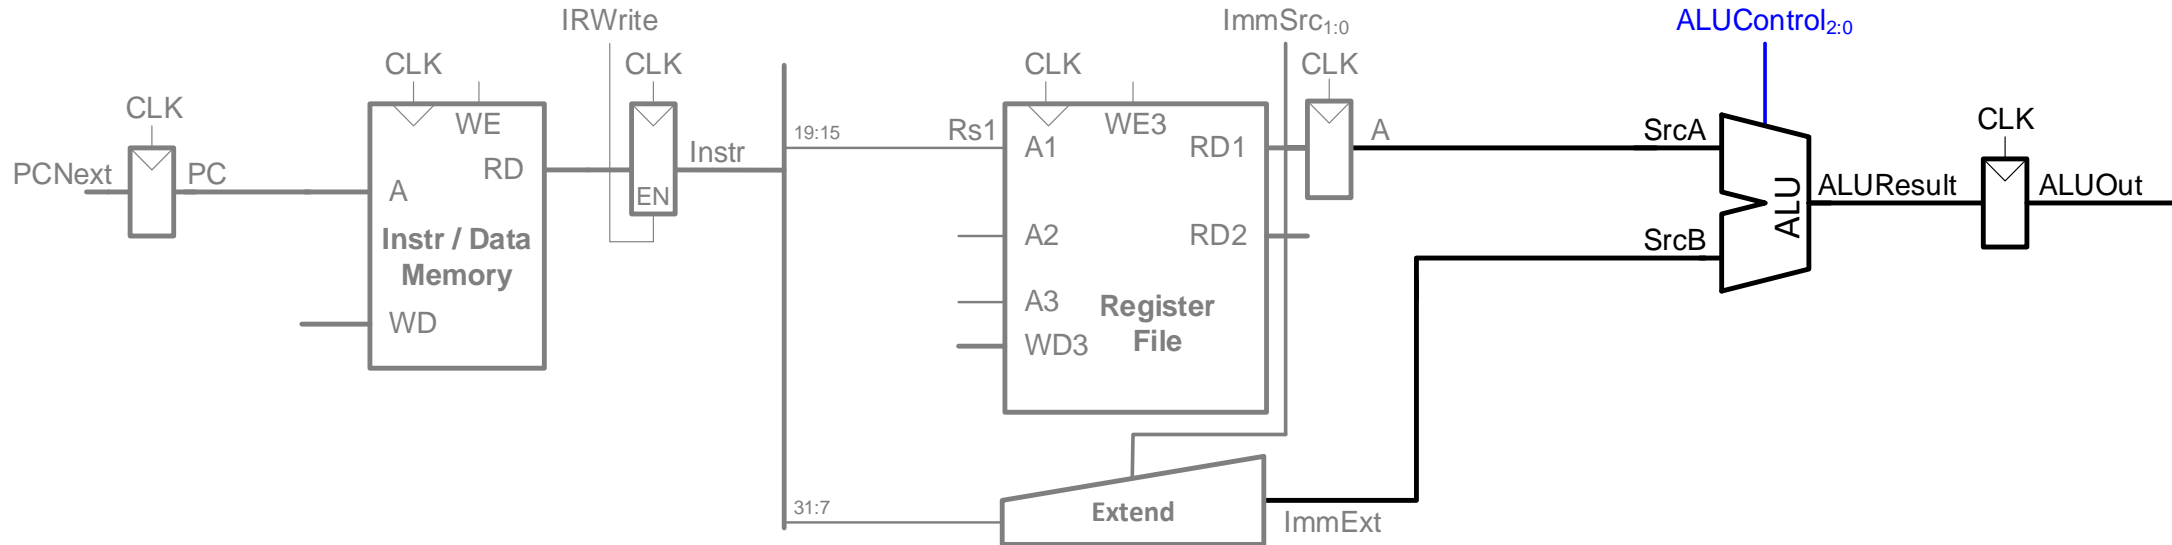
Multicycle Datapath: 1_w Get Sources

STEP 2: Read source operand from RF and extend immediate



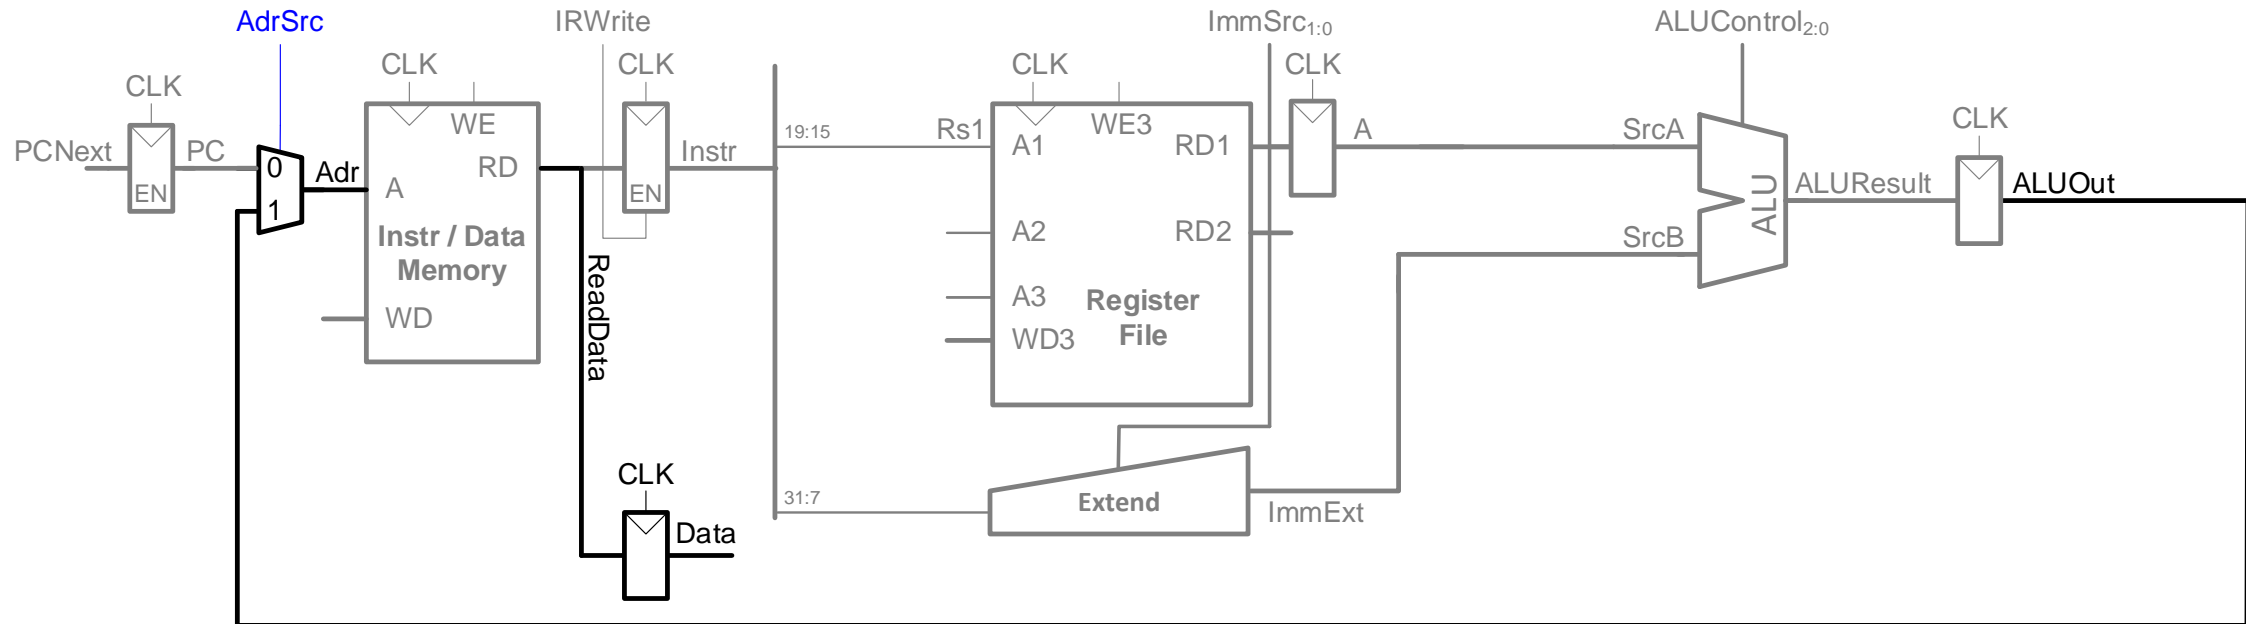
Multicycle Datapath: 1_w Address

STEP 3: Compute the memory address



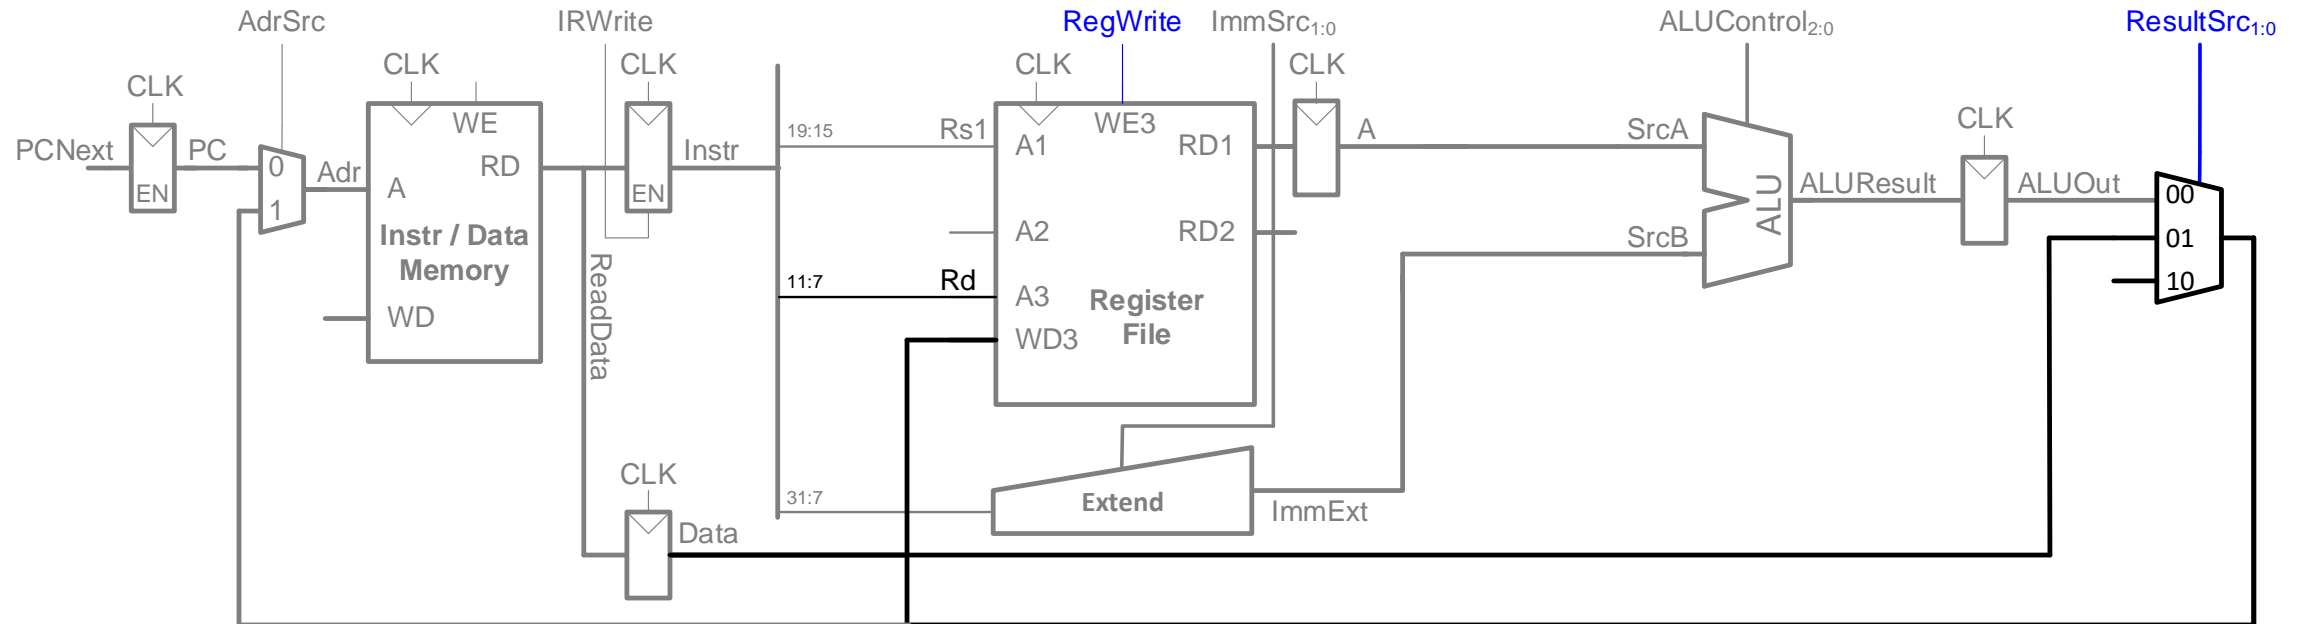
Multicycle Datapath: 1_w Memory Read

STEP 4: Read data from memory



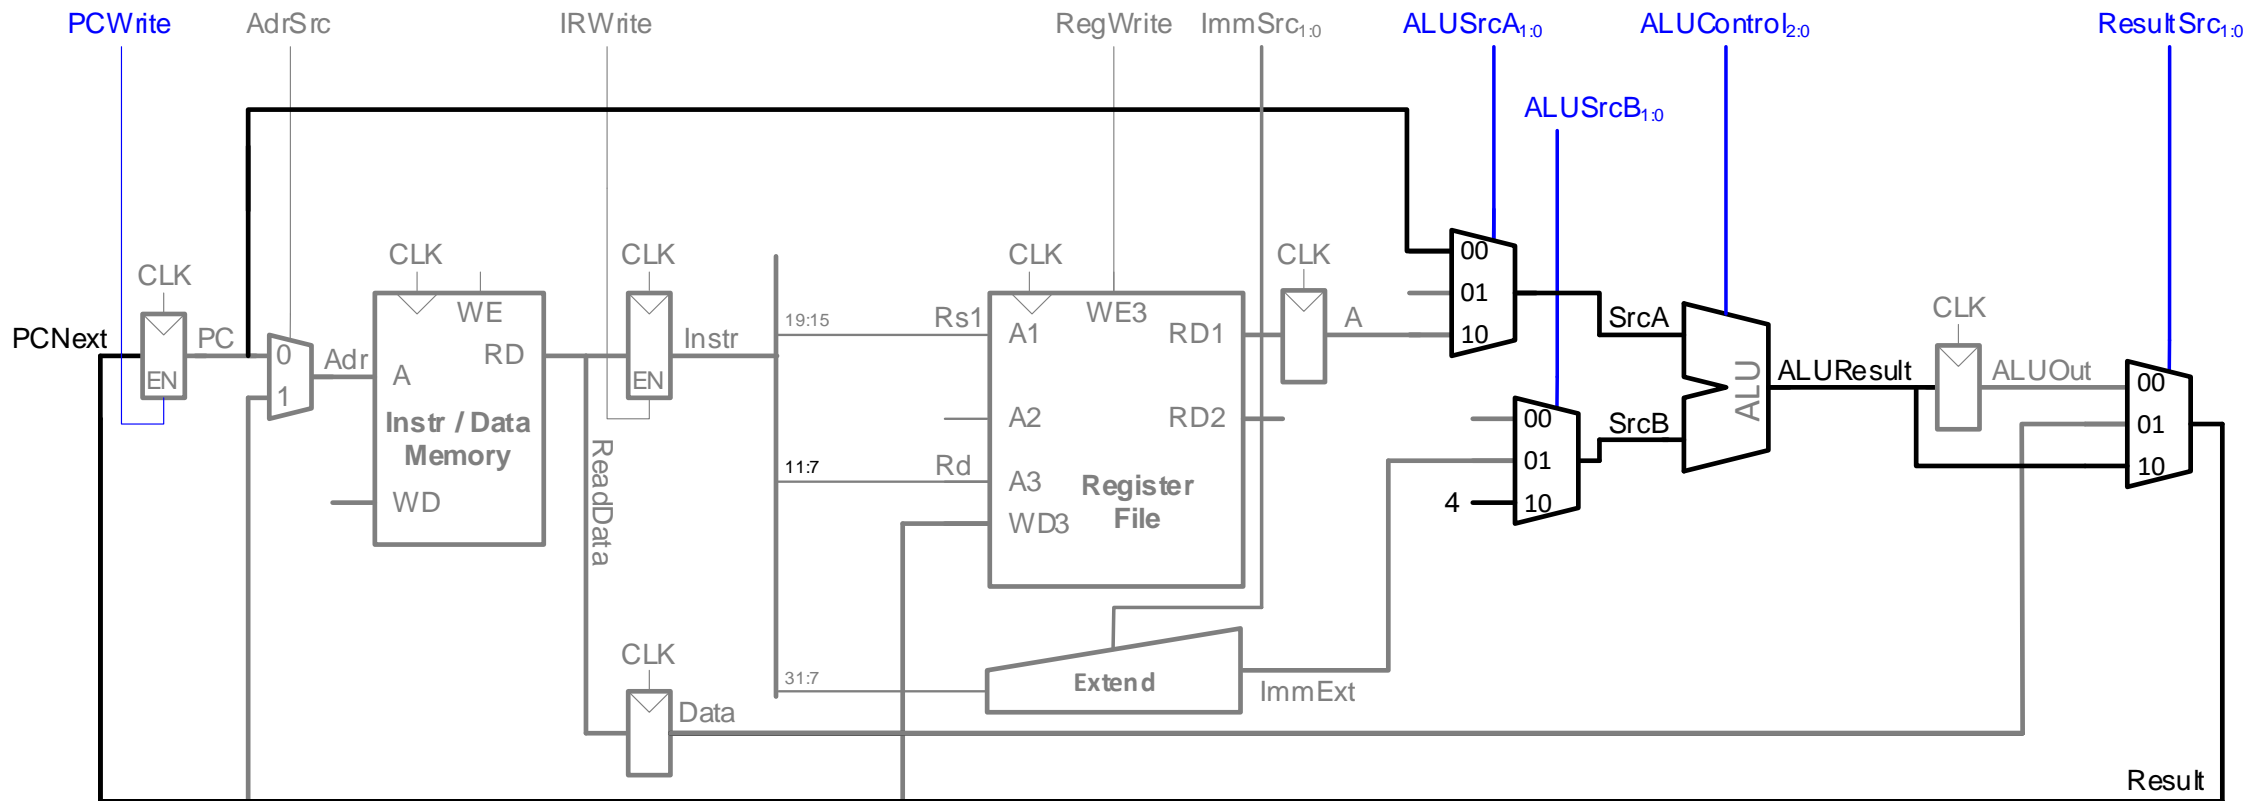
Multicycle Datapath: 1_W Write Register

STEP 5: Write data back to register file



Multicycle Datapath: Increment PC

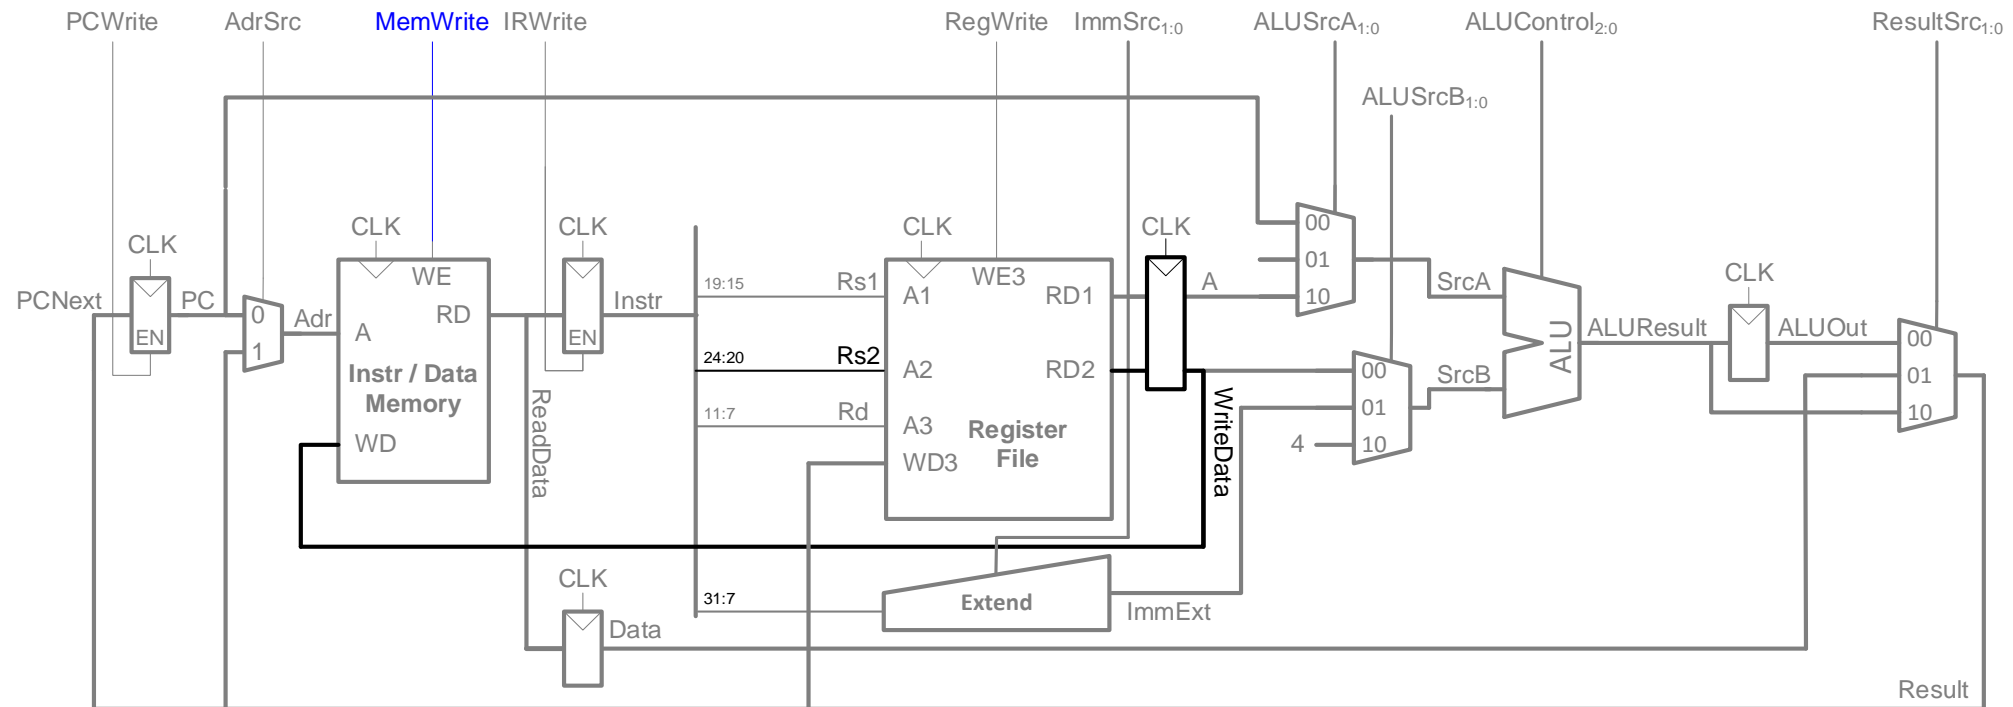
STEP 6: Increment PC: $PC = PC + 4$



Multicycle Datapath: Other Instructions

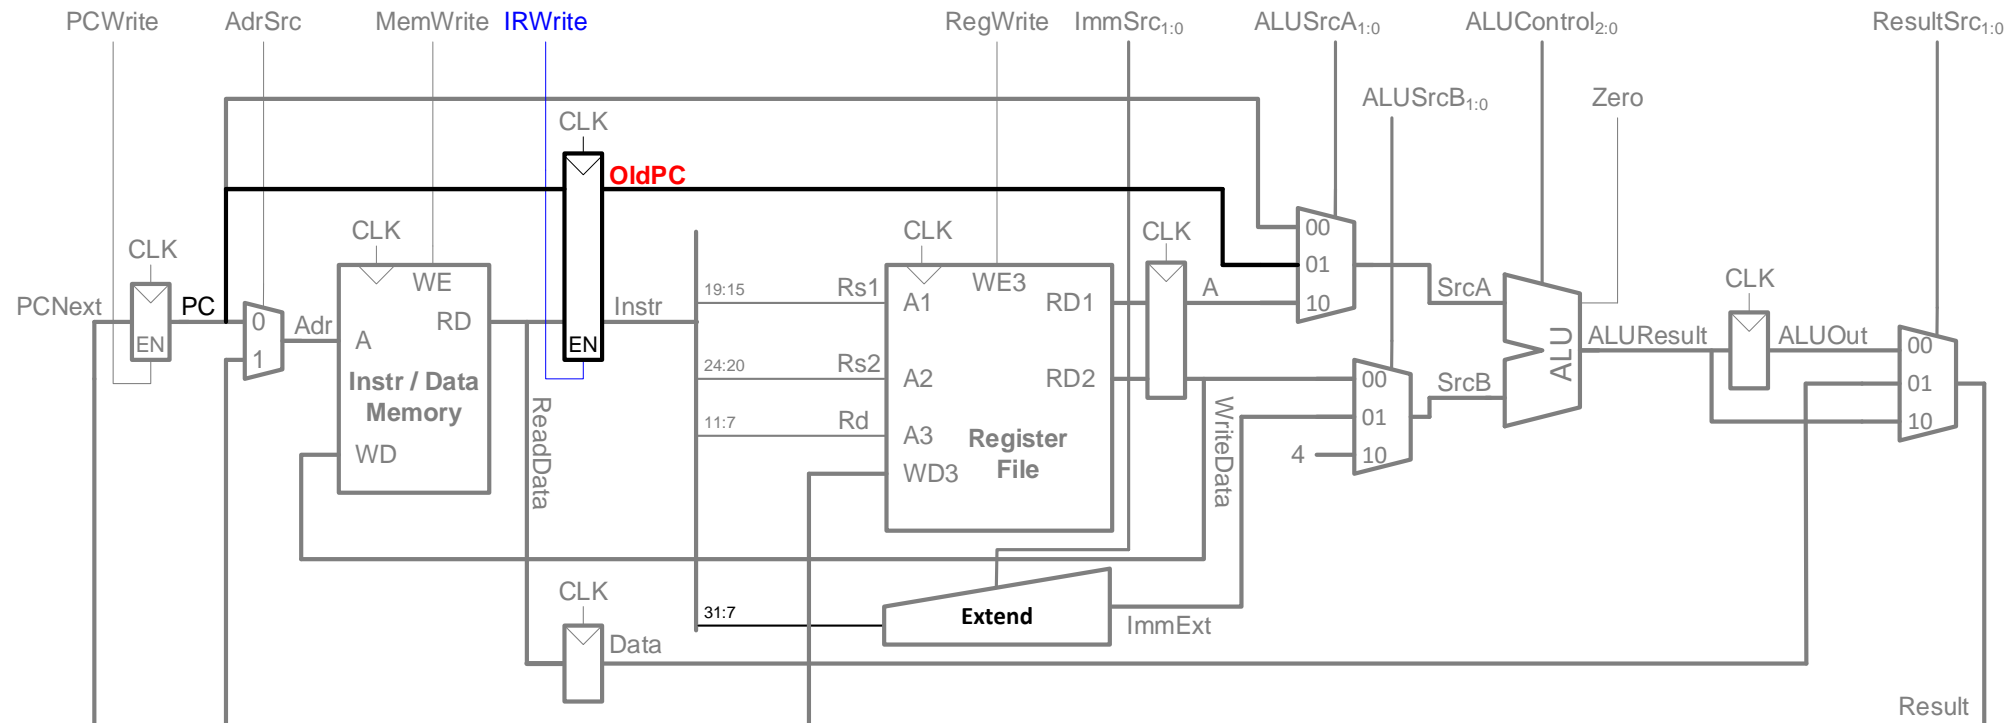
DDCA Ch7 - part 8: RISC-V Multicycle Processor - Other Instructions <https://www.youtube.com/watch?v=dnITBQQDmwU>

Write data from Register File (rs2) to memory



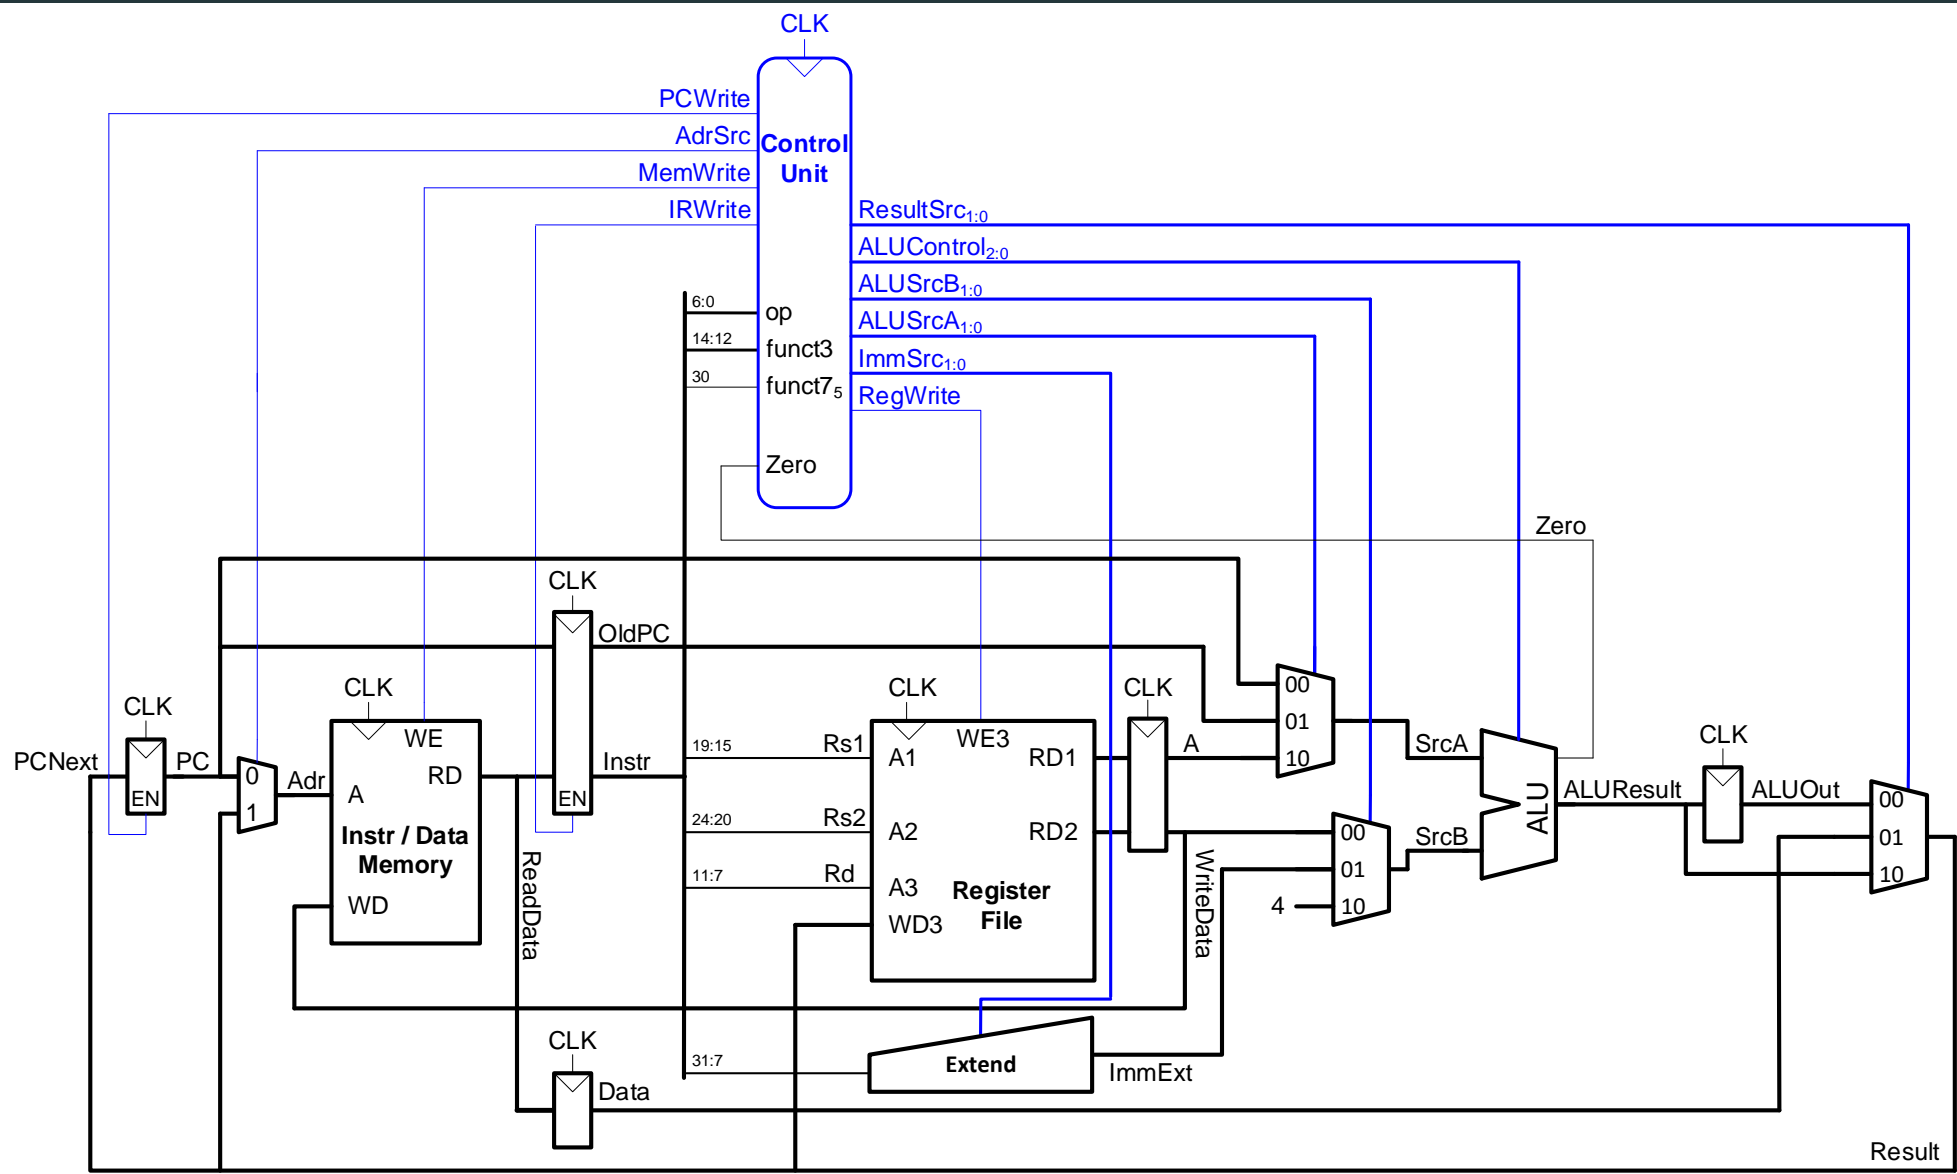
Multicycle Datapath: beq

- Calculate branch target address: $BTA = PC + imm$



PC is updated in Fetch stage, so need to save **old (current) PC**

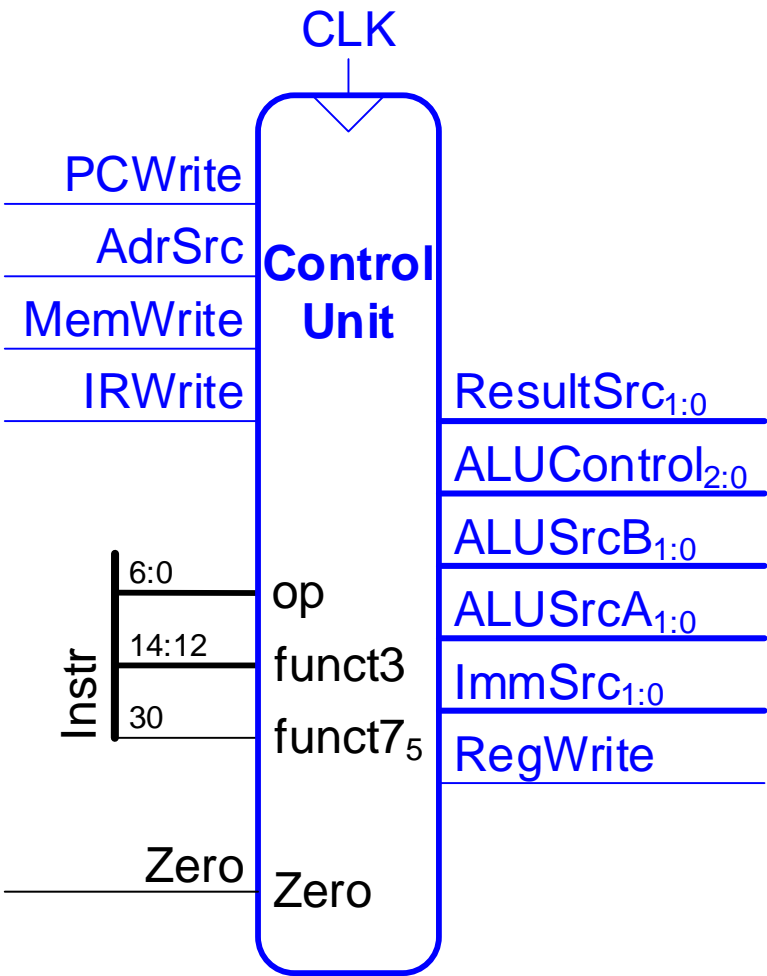
Multicycle RISC-V Processor



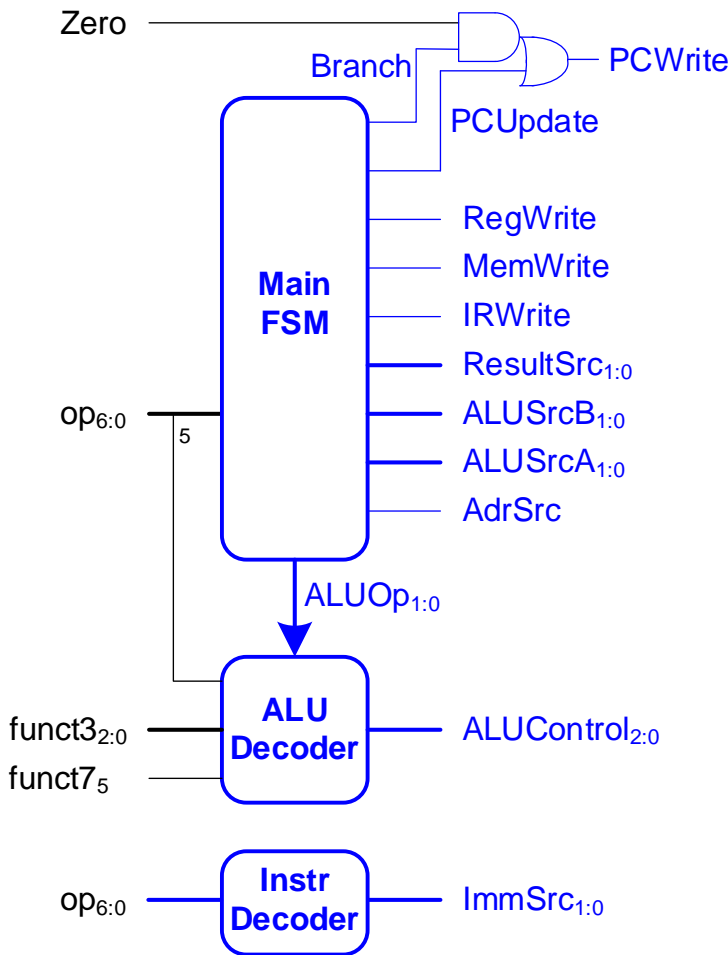
Multicycle Control

DDCA Ch7 - Part 9: RISC-V Multicycle Processor Control <https://www.youtube.com/watch?v=YUJhNTpunqI>

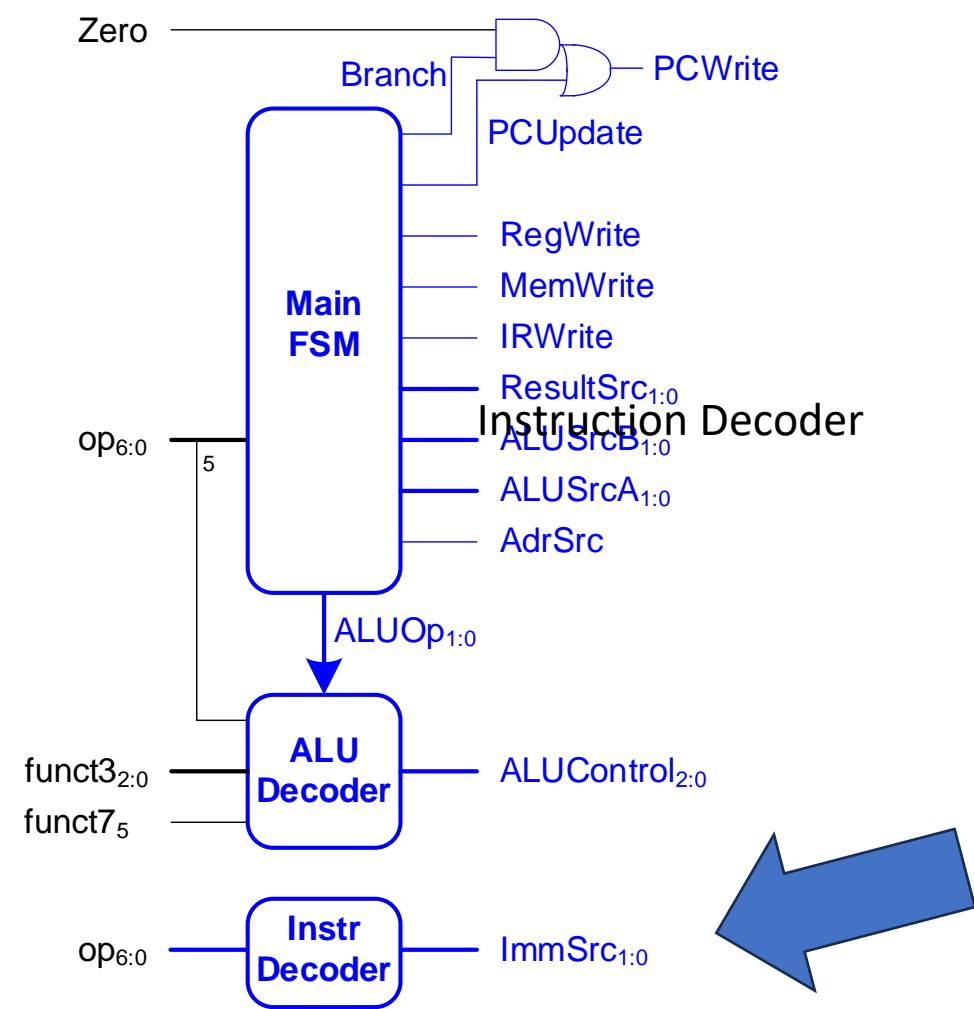
High-Level View



Low-Level View



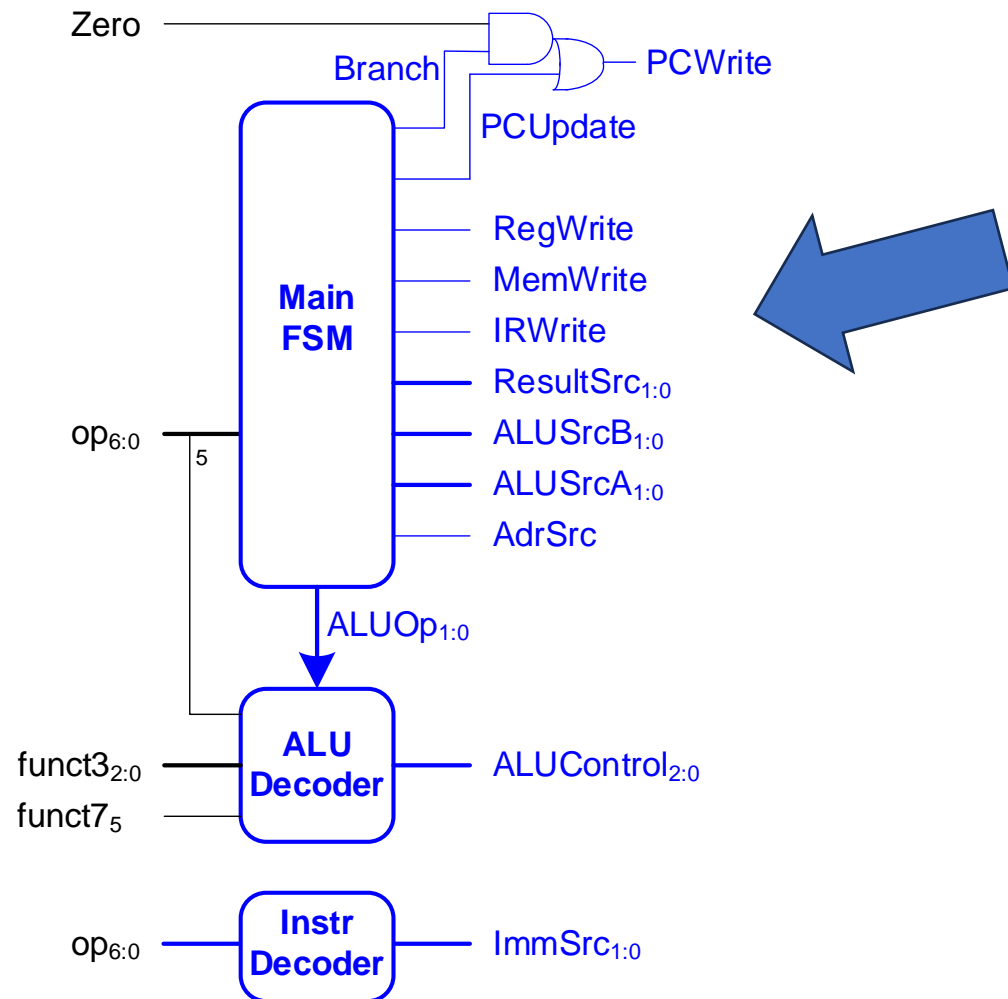
Multicycle Control: Instruction Decoder



Instruction Decoder

op	Instruction	ImmSrc
3	lw	00
35	sw	01
51	R-type	XX
99	beq	10

Multicycle Control: Main FSM

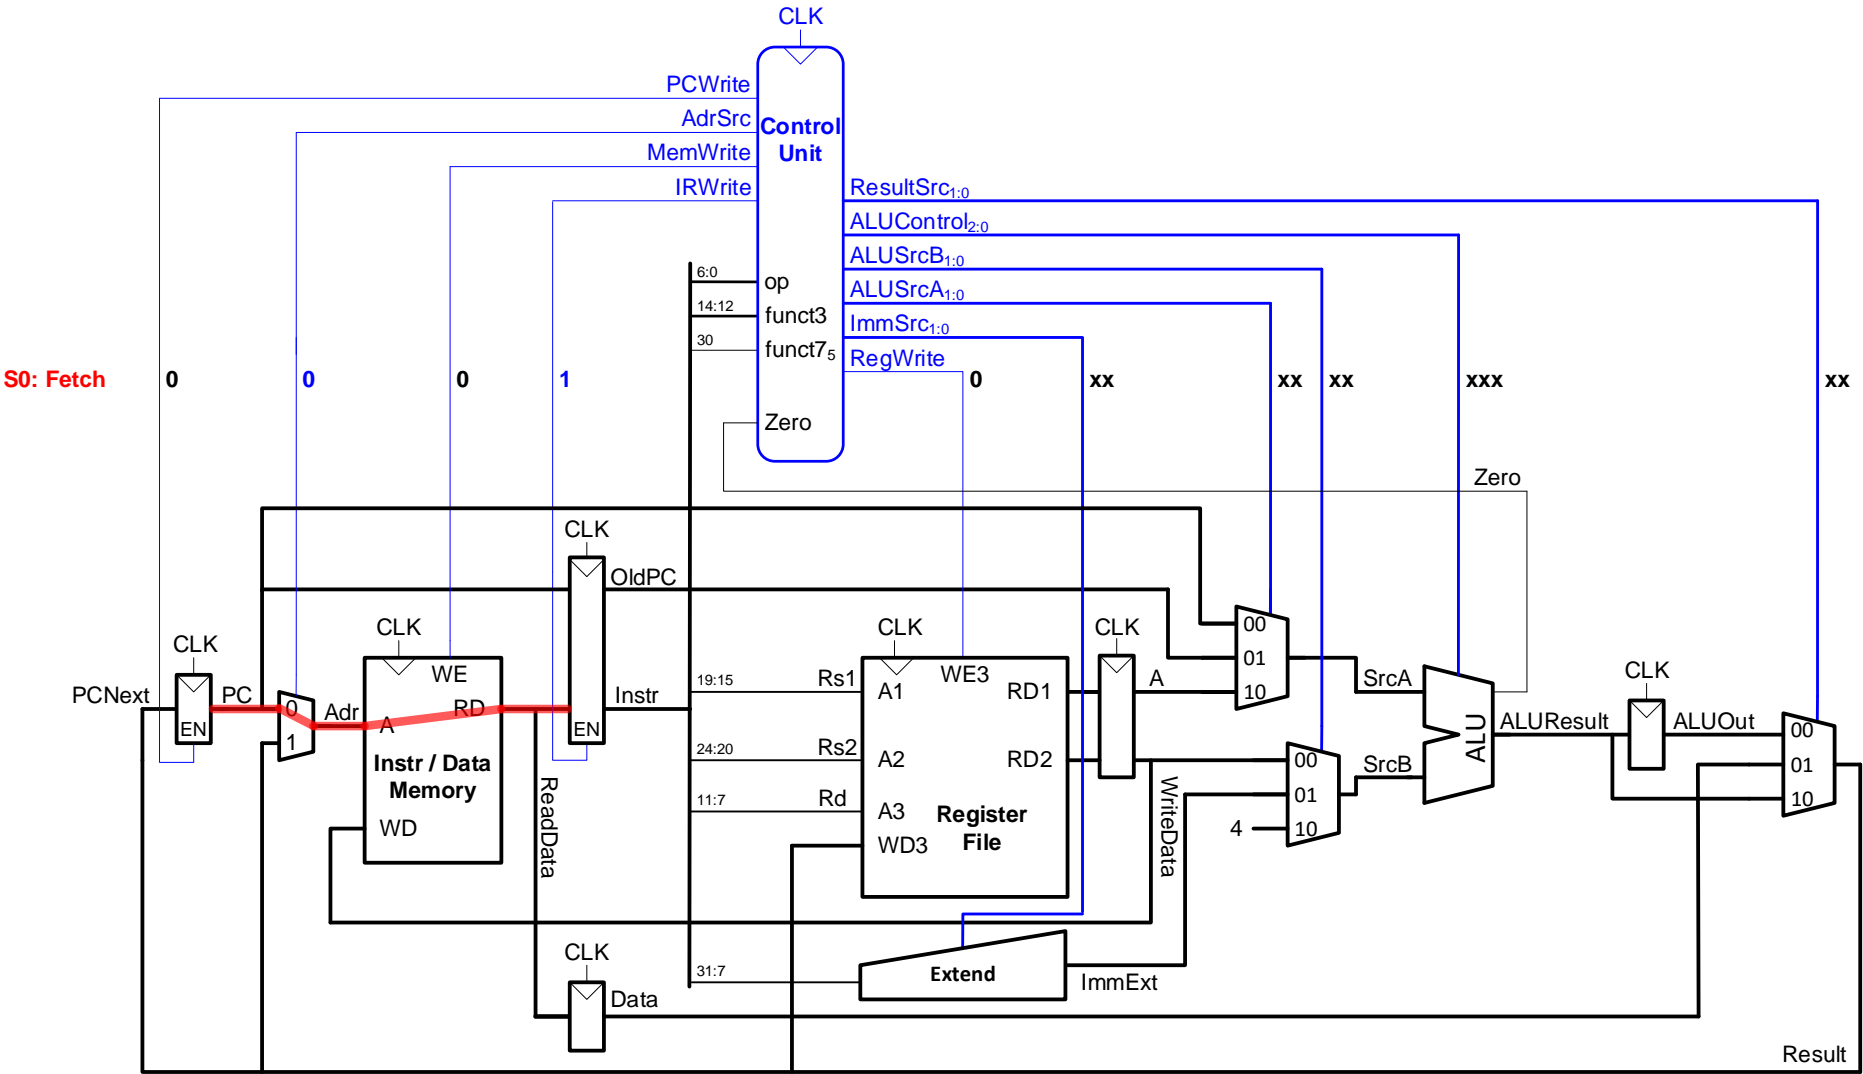



Main FSM

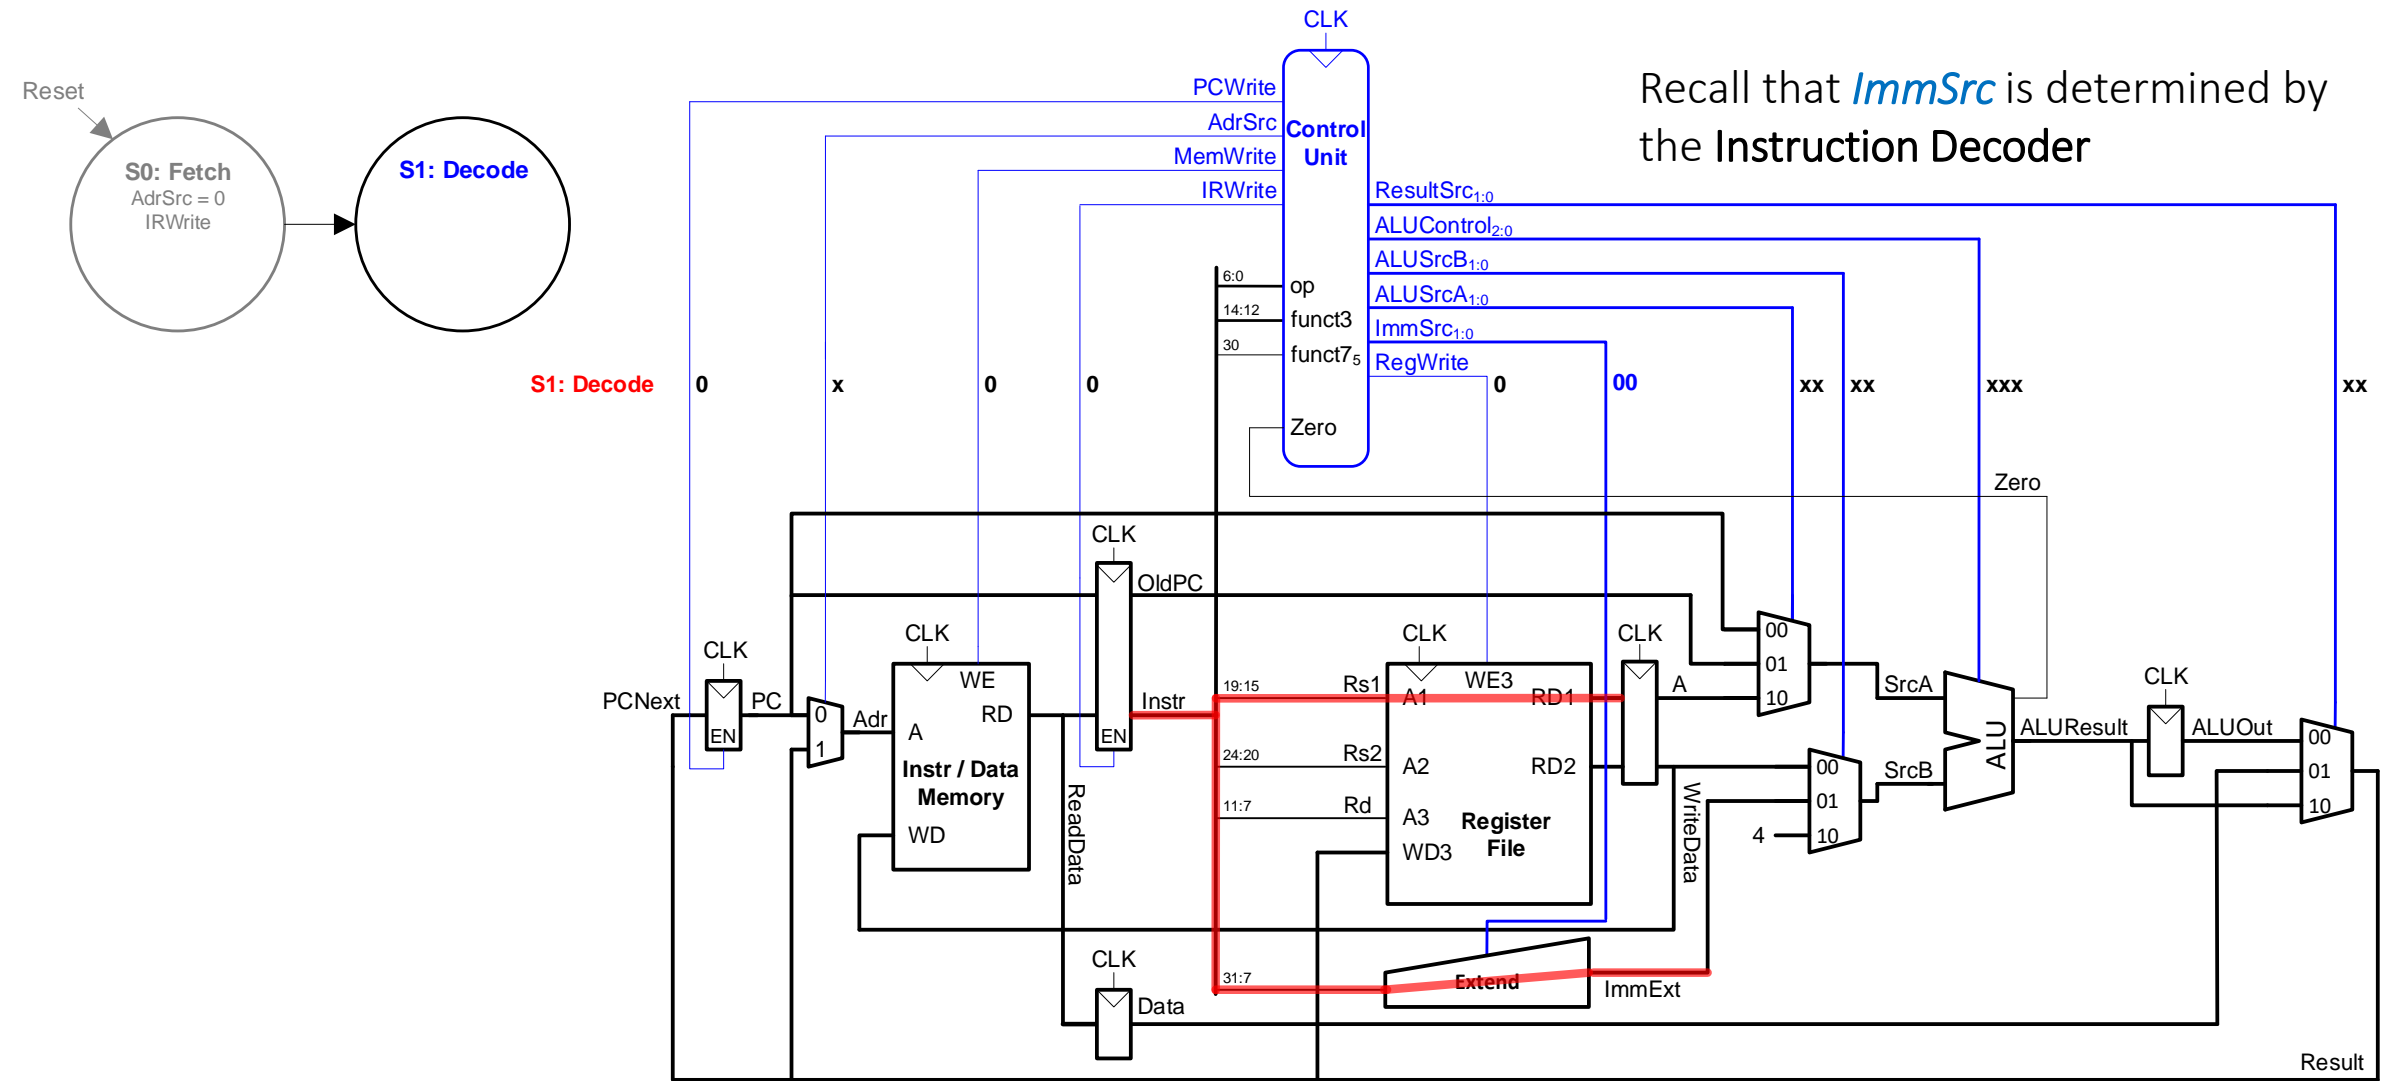
To declutter FSM:

- **Write enable signals** (RegWrite, MemWrite, IRWrite, PCUpdate, and Branch) are **0** if not listed in a state.
- **Other signals are don't care** if not listed in a state

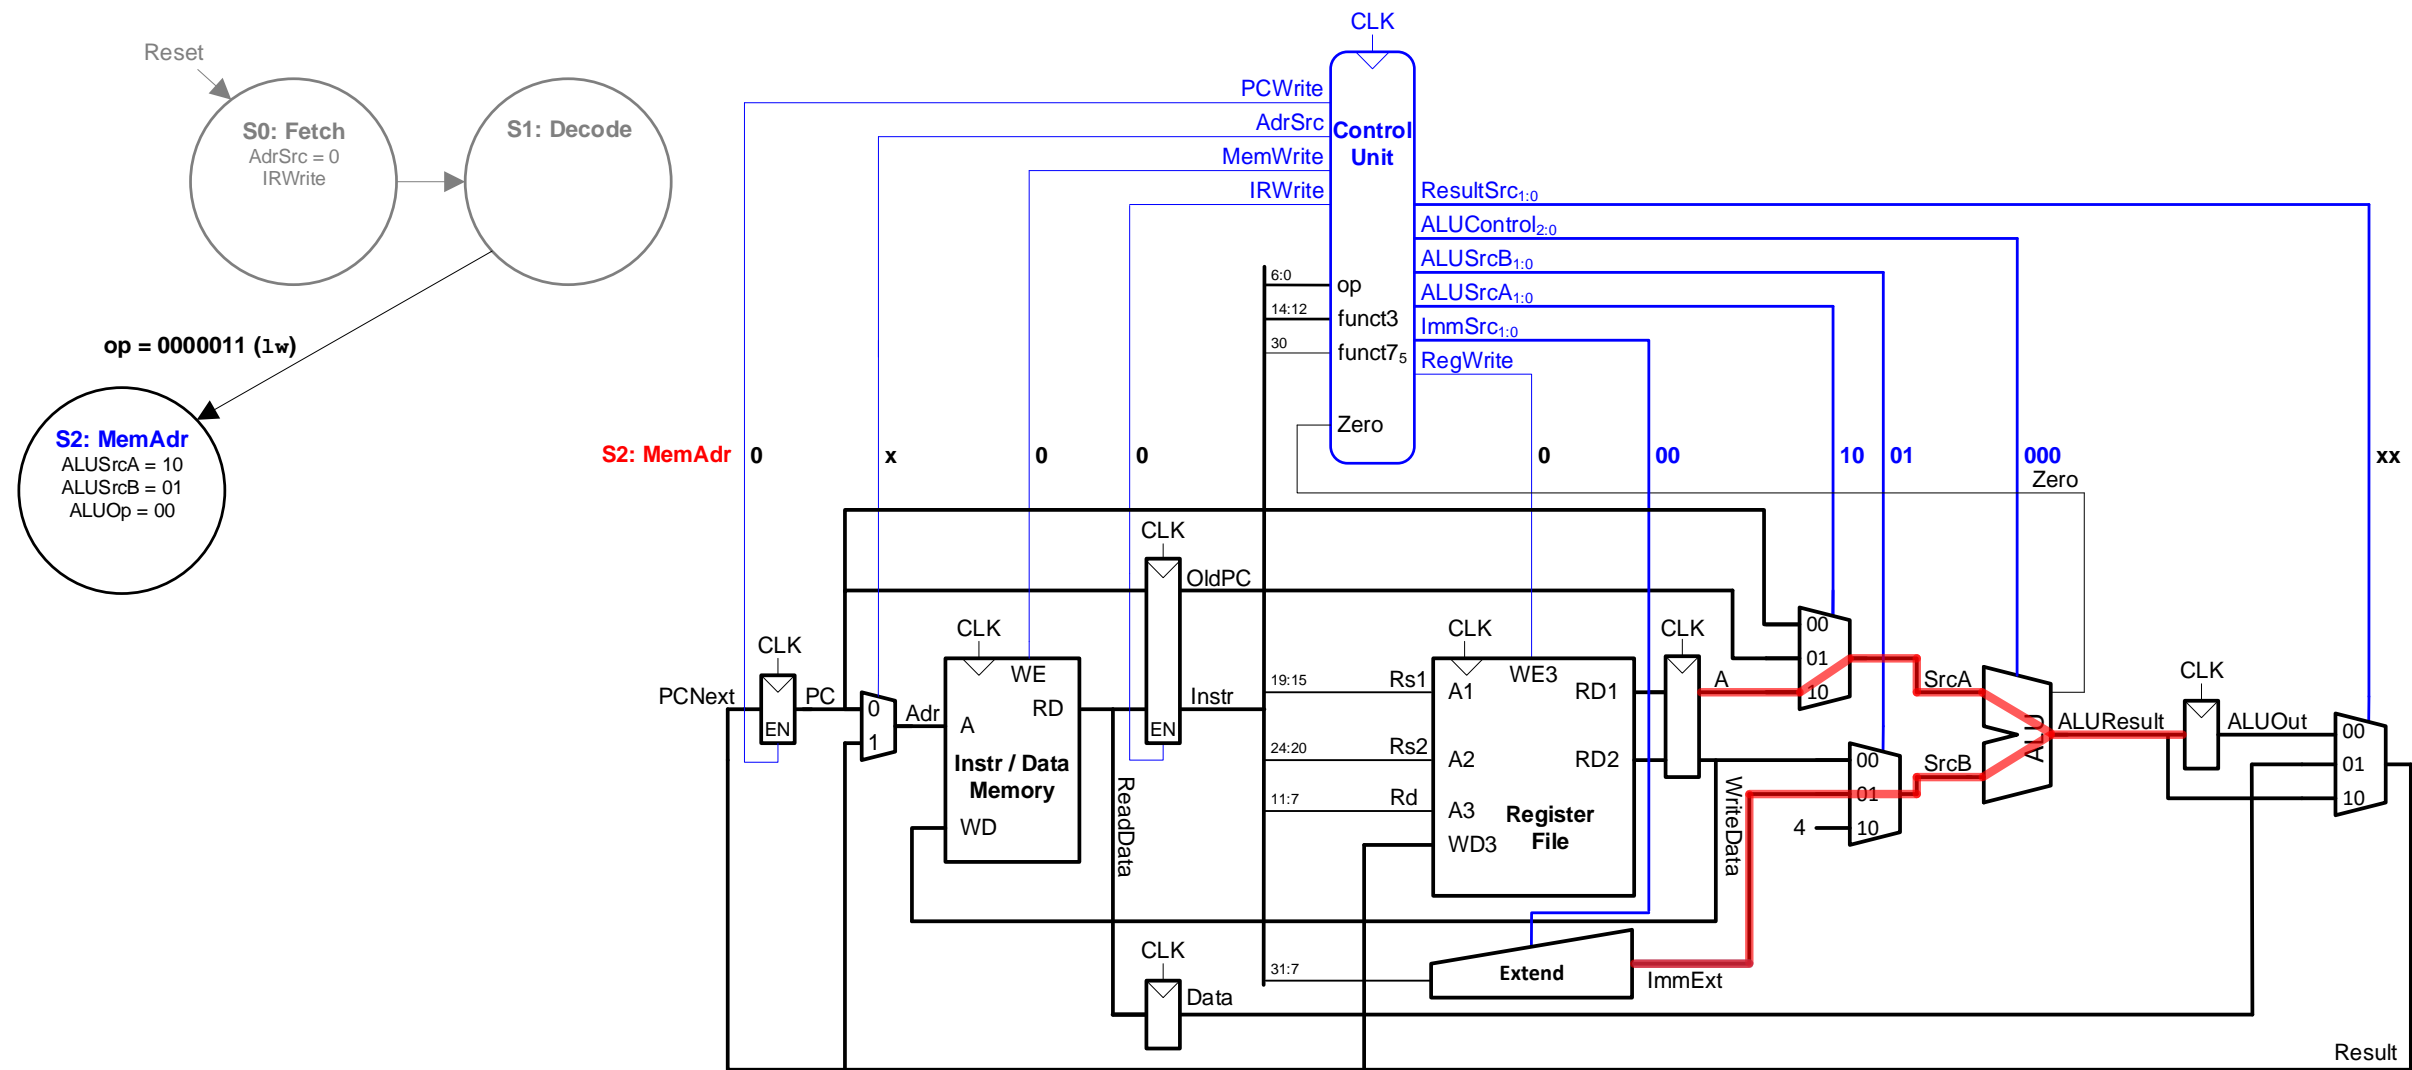
Main FSM: Fetch



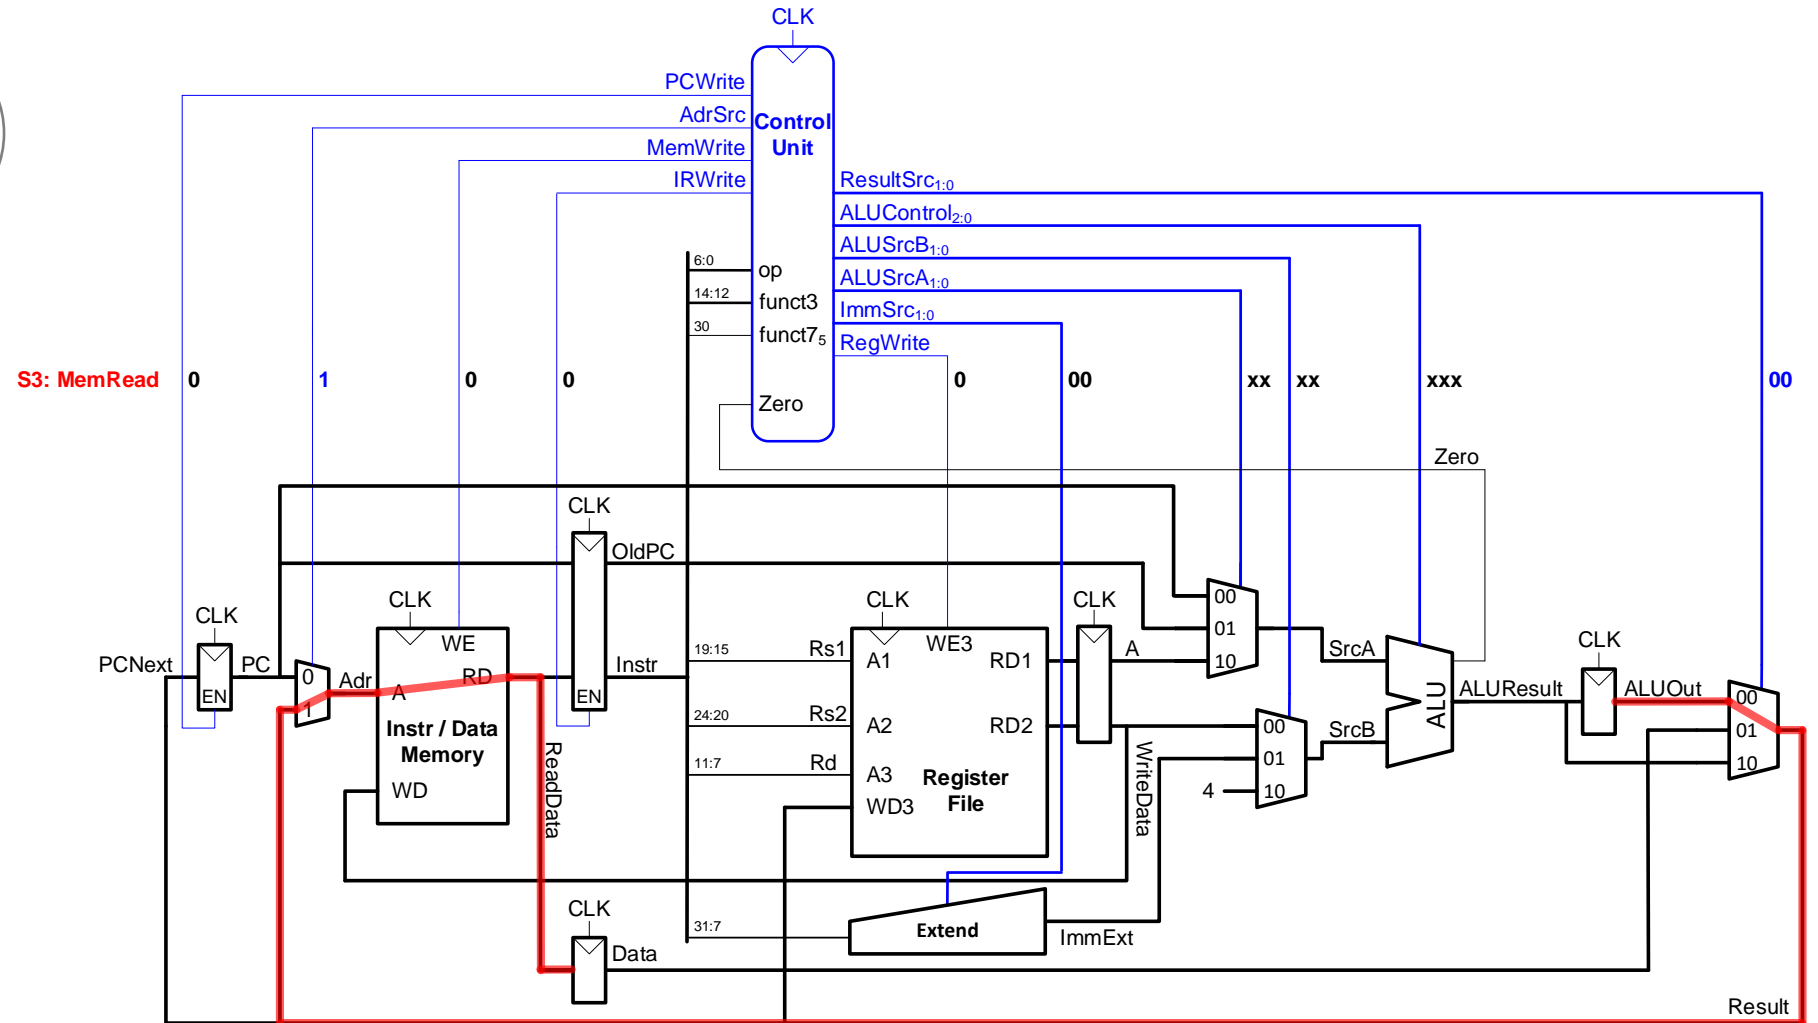
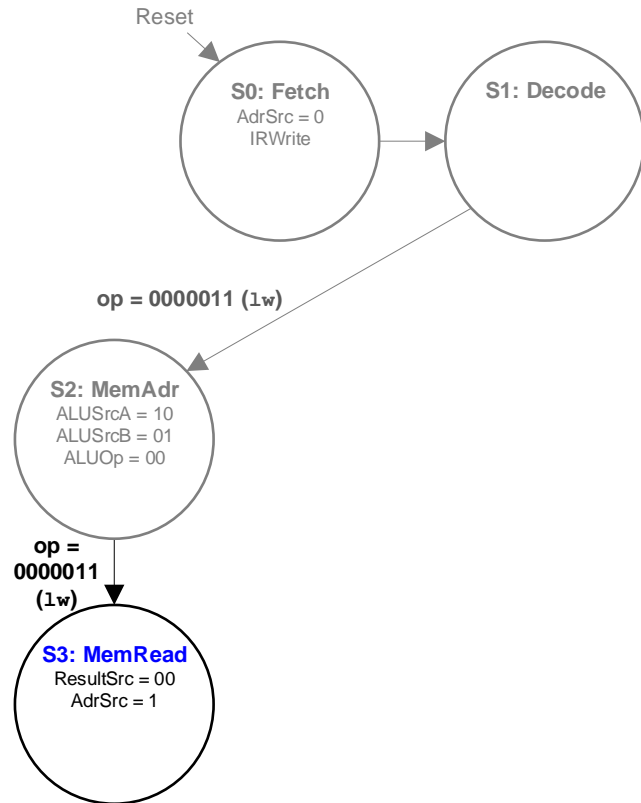
Main FSM: Decode



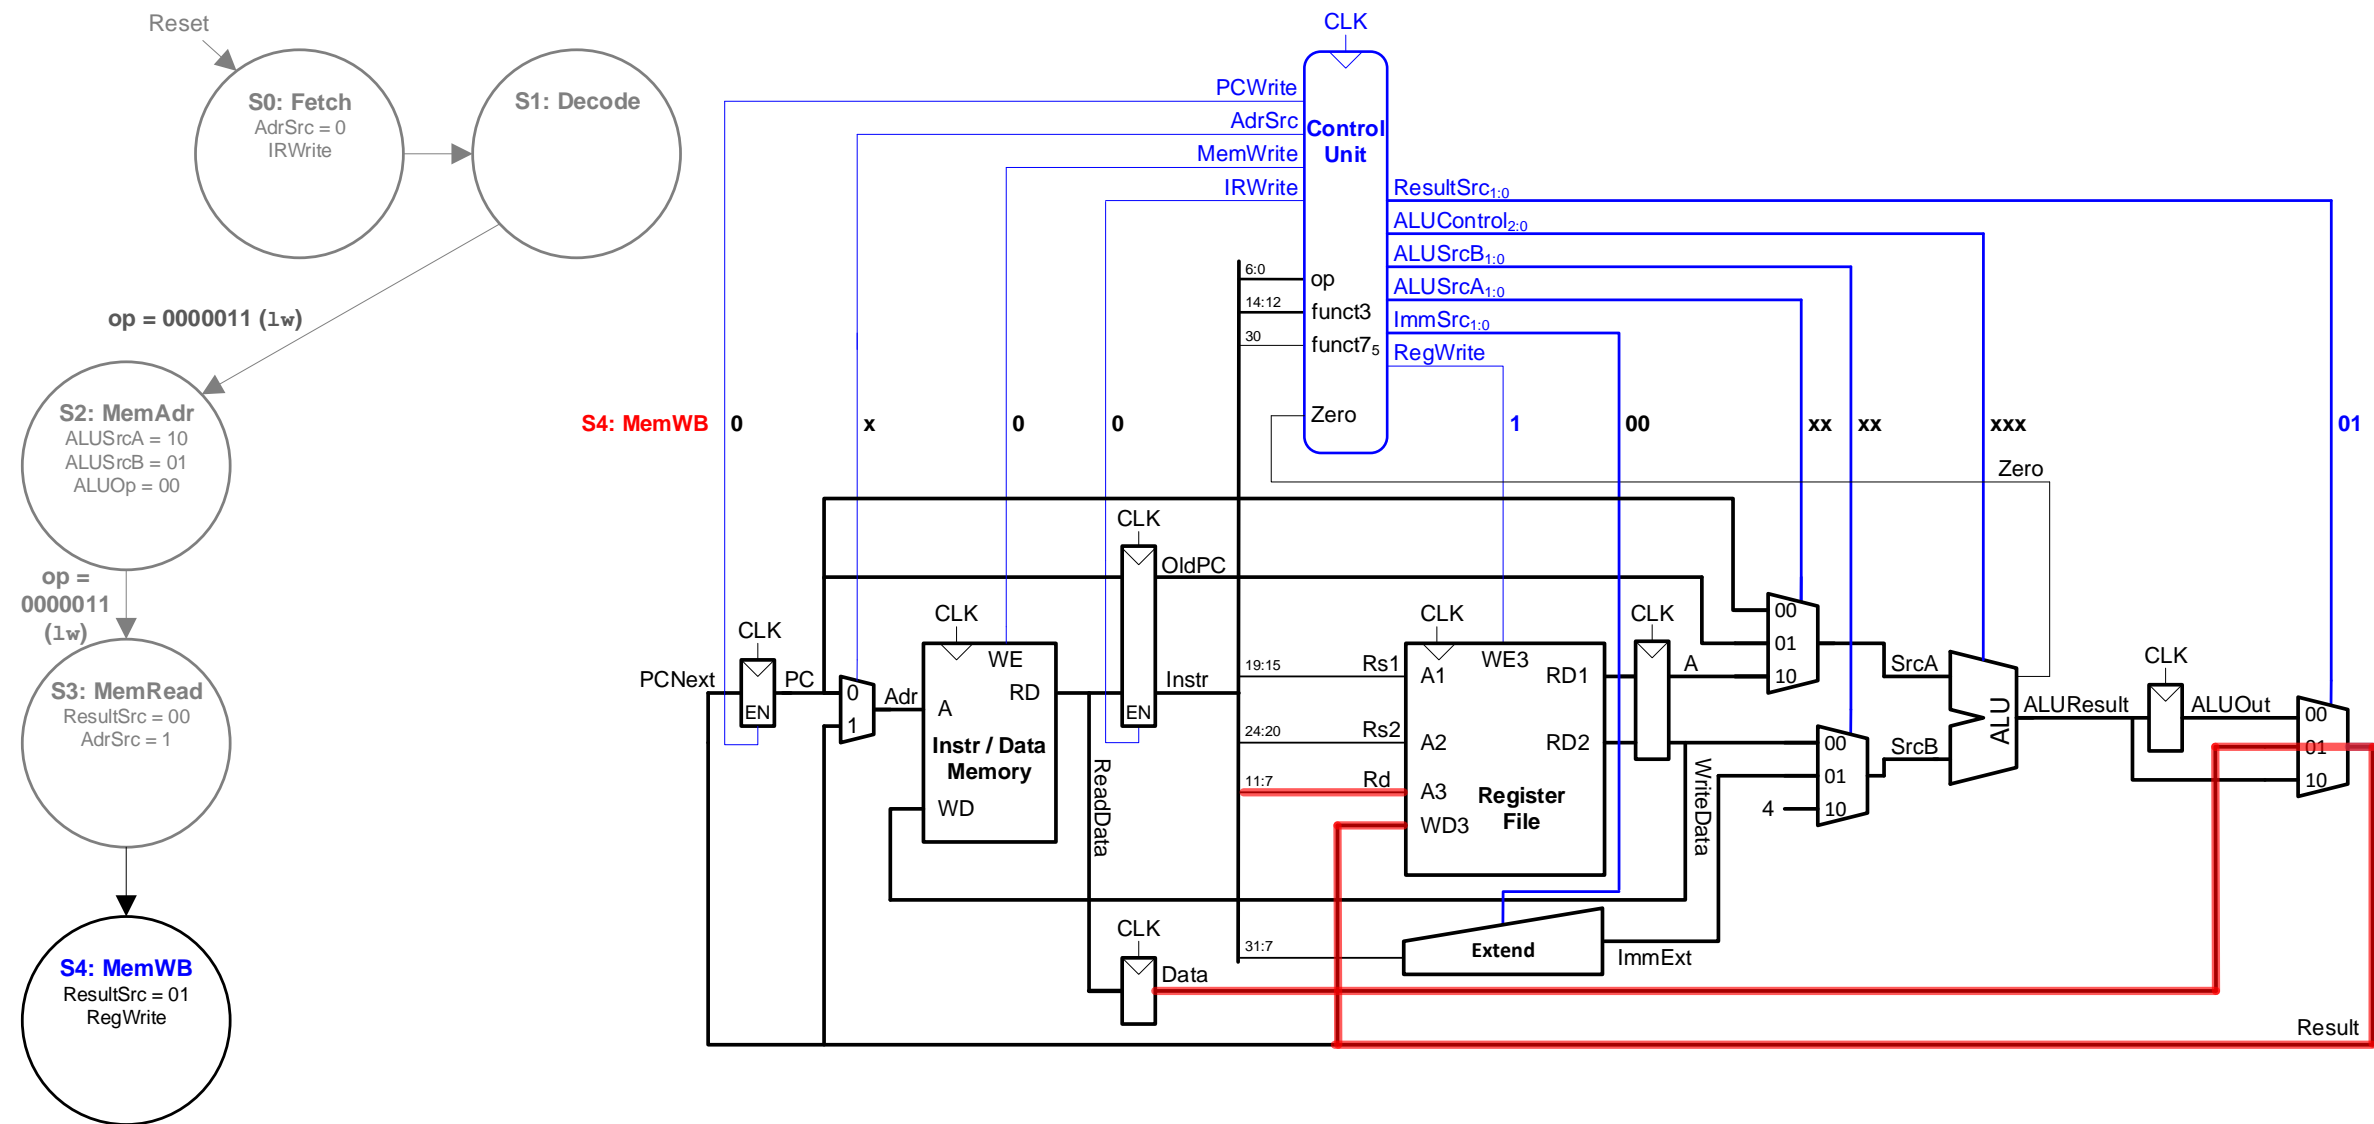
Main FSM: Address



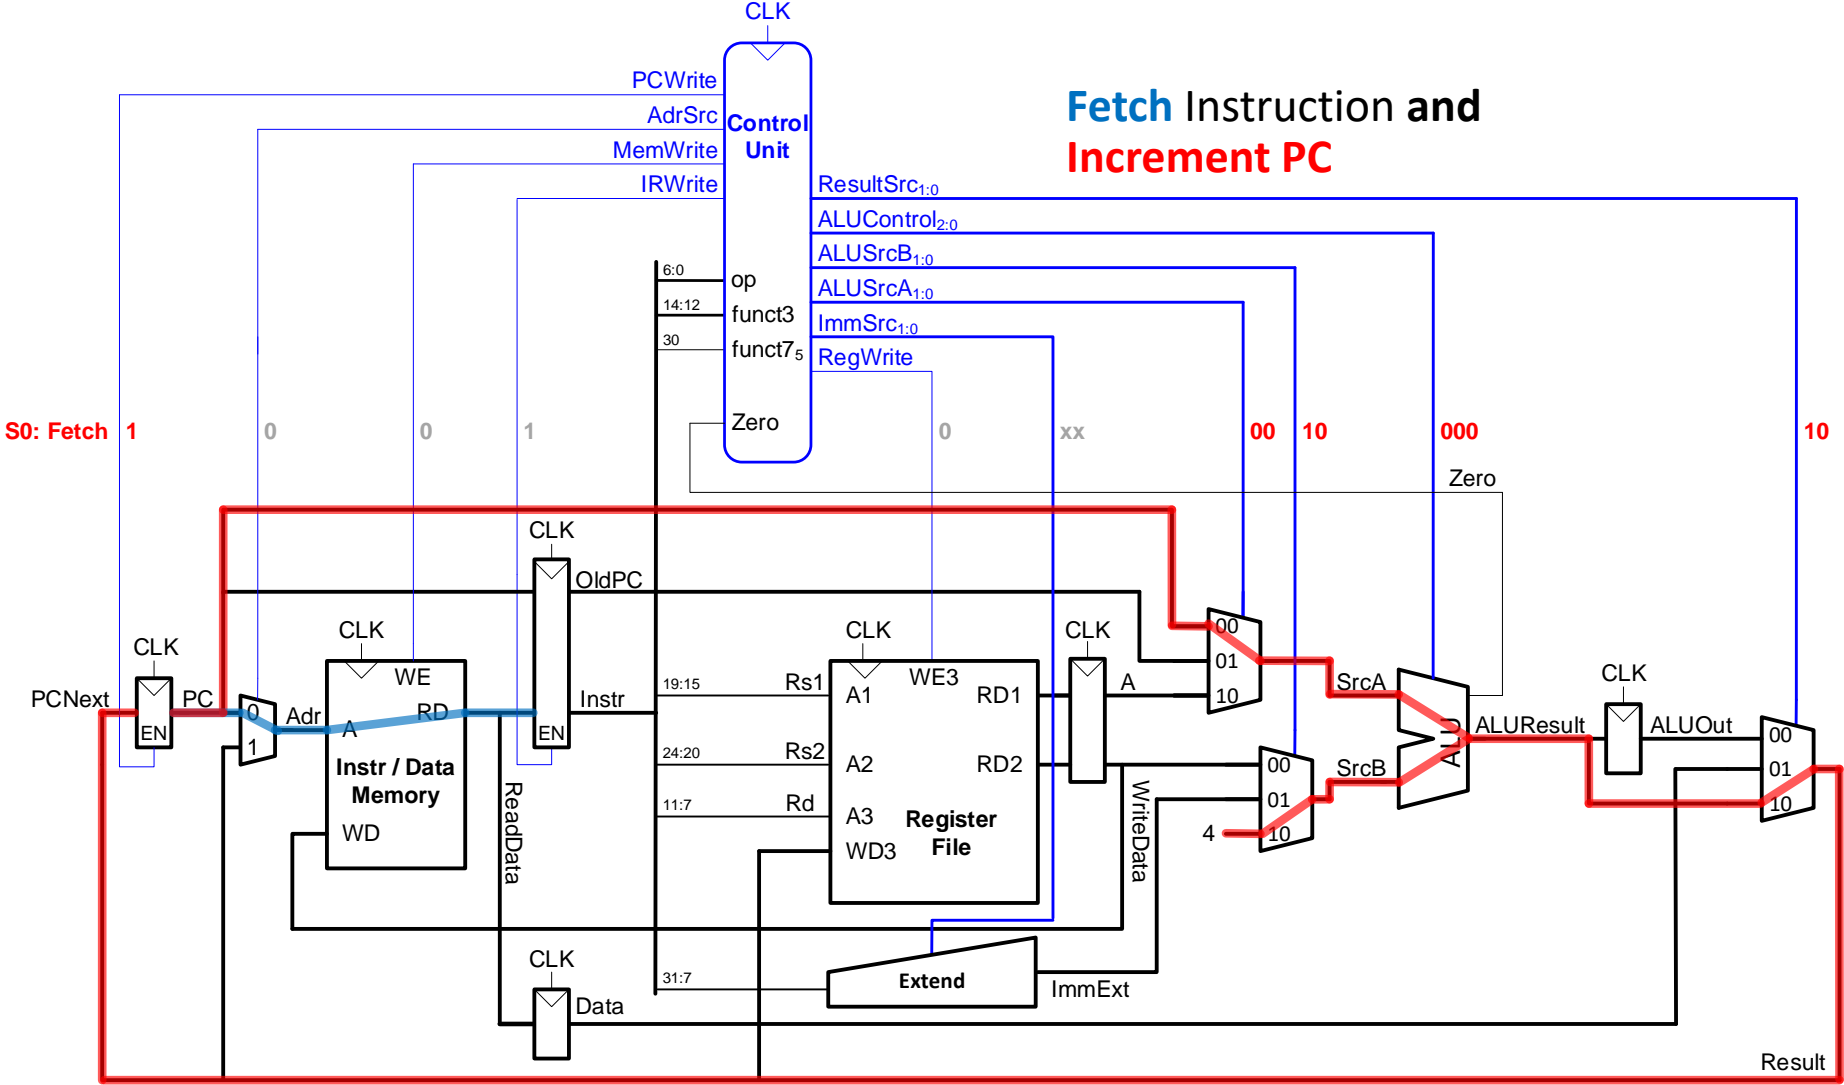
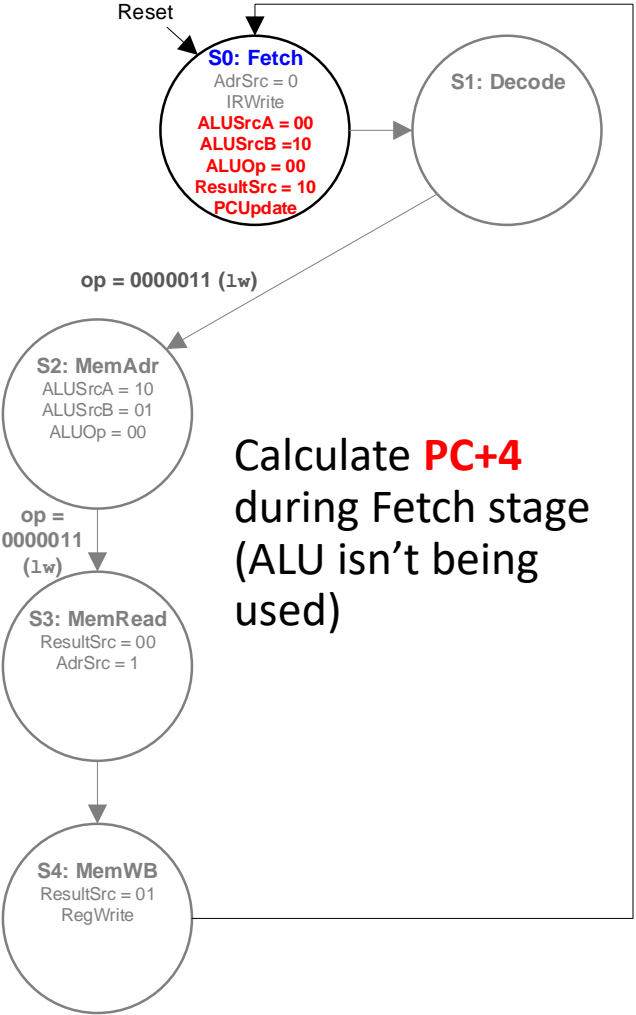
Main FSM: Read Memory



Main FSM: Write RF



Main FSM: Fetch Revisited



Multicycle Control: Other Instructions

DDCA Ch7 - Part 10: RISC-V Multicycle Processor Control: Other Instructions

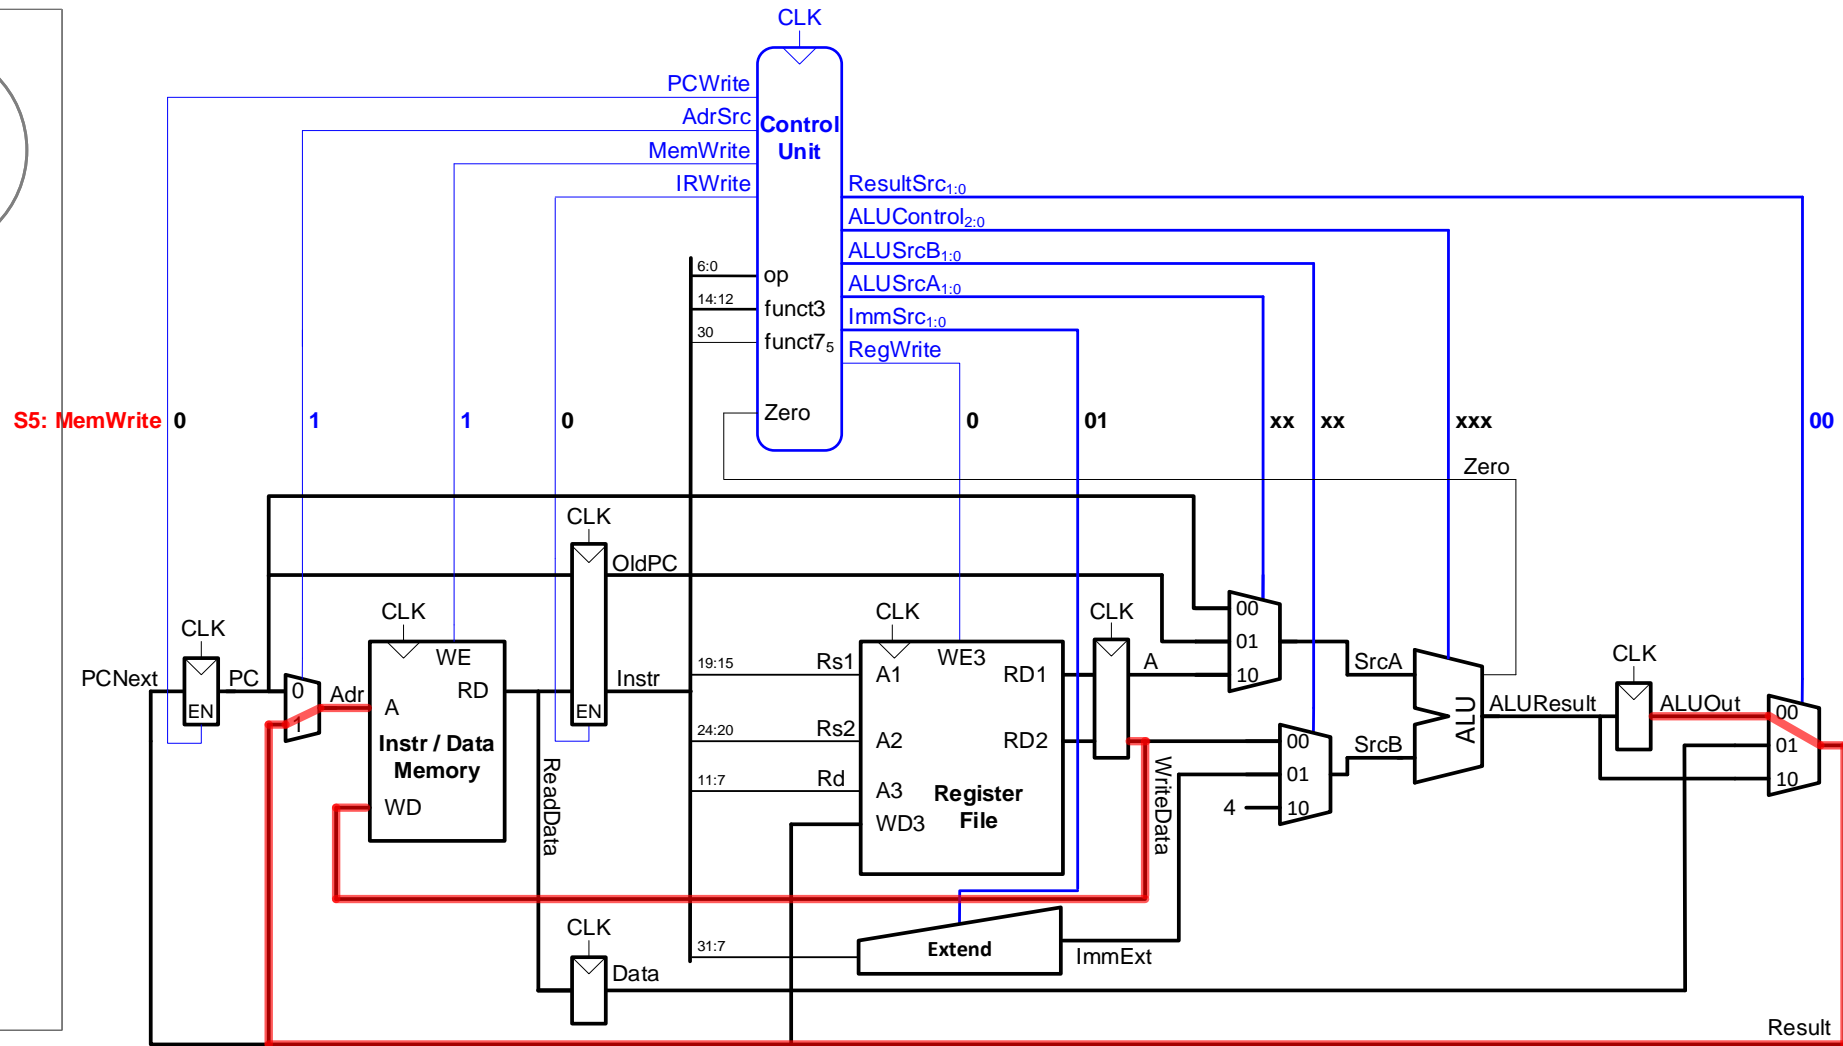
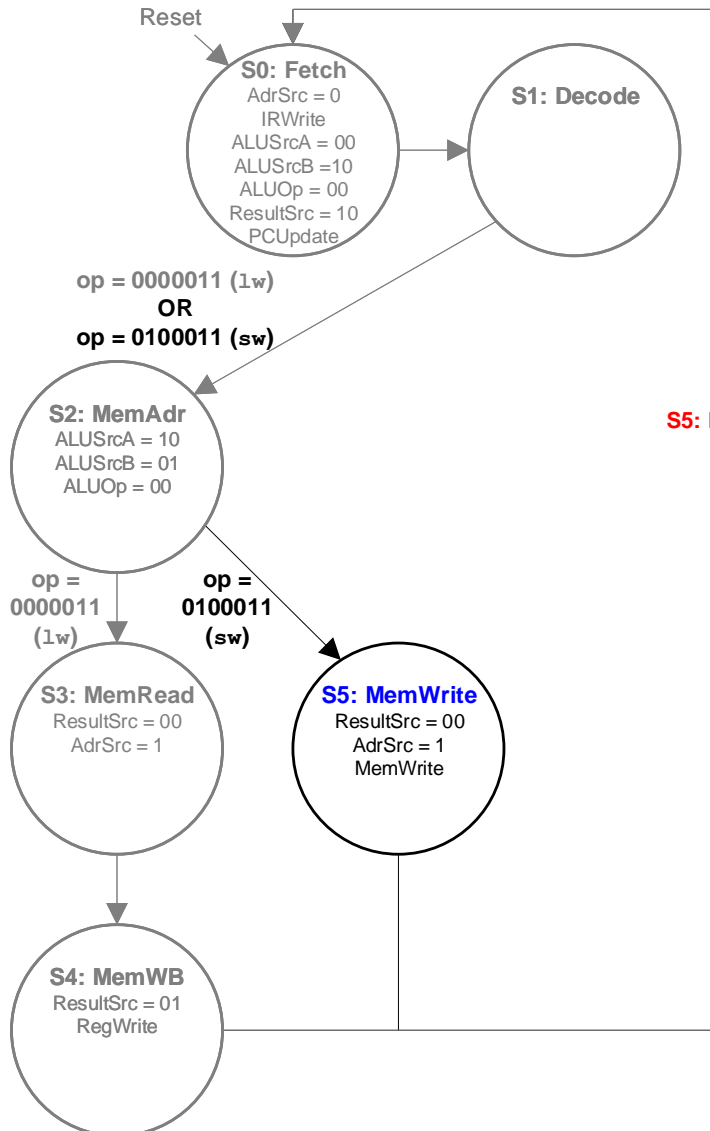
<https://www.youtube.com/watch?v=rSodrnsYXQ8>

24.04.2025

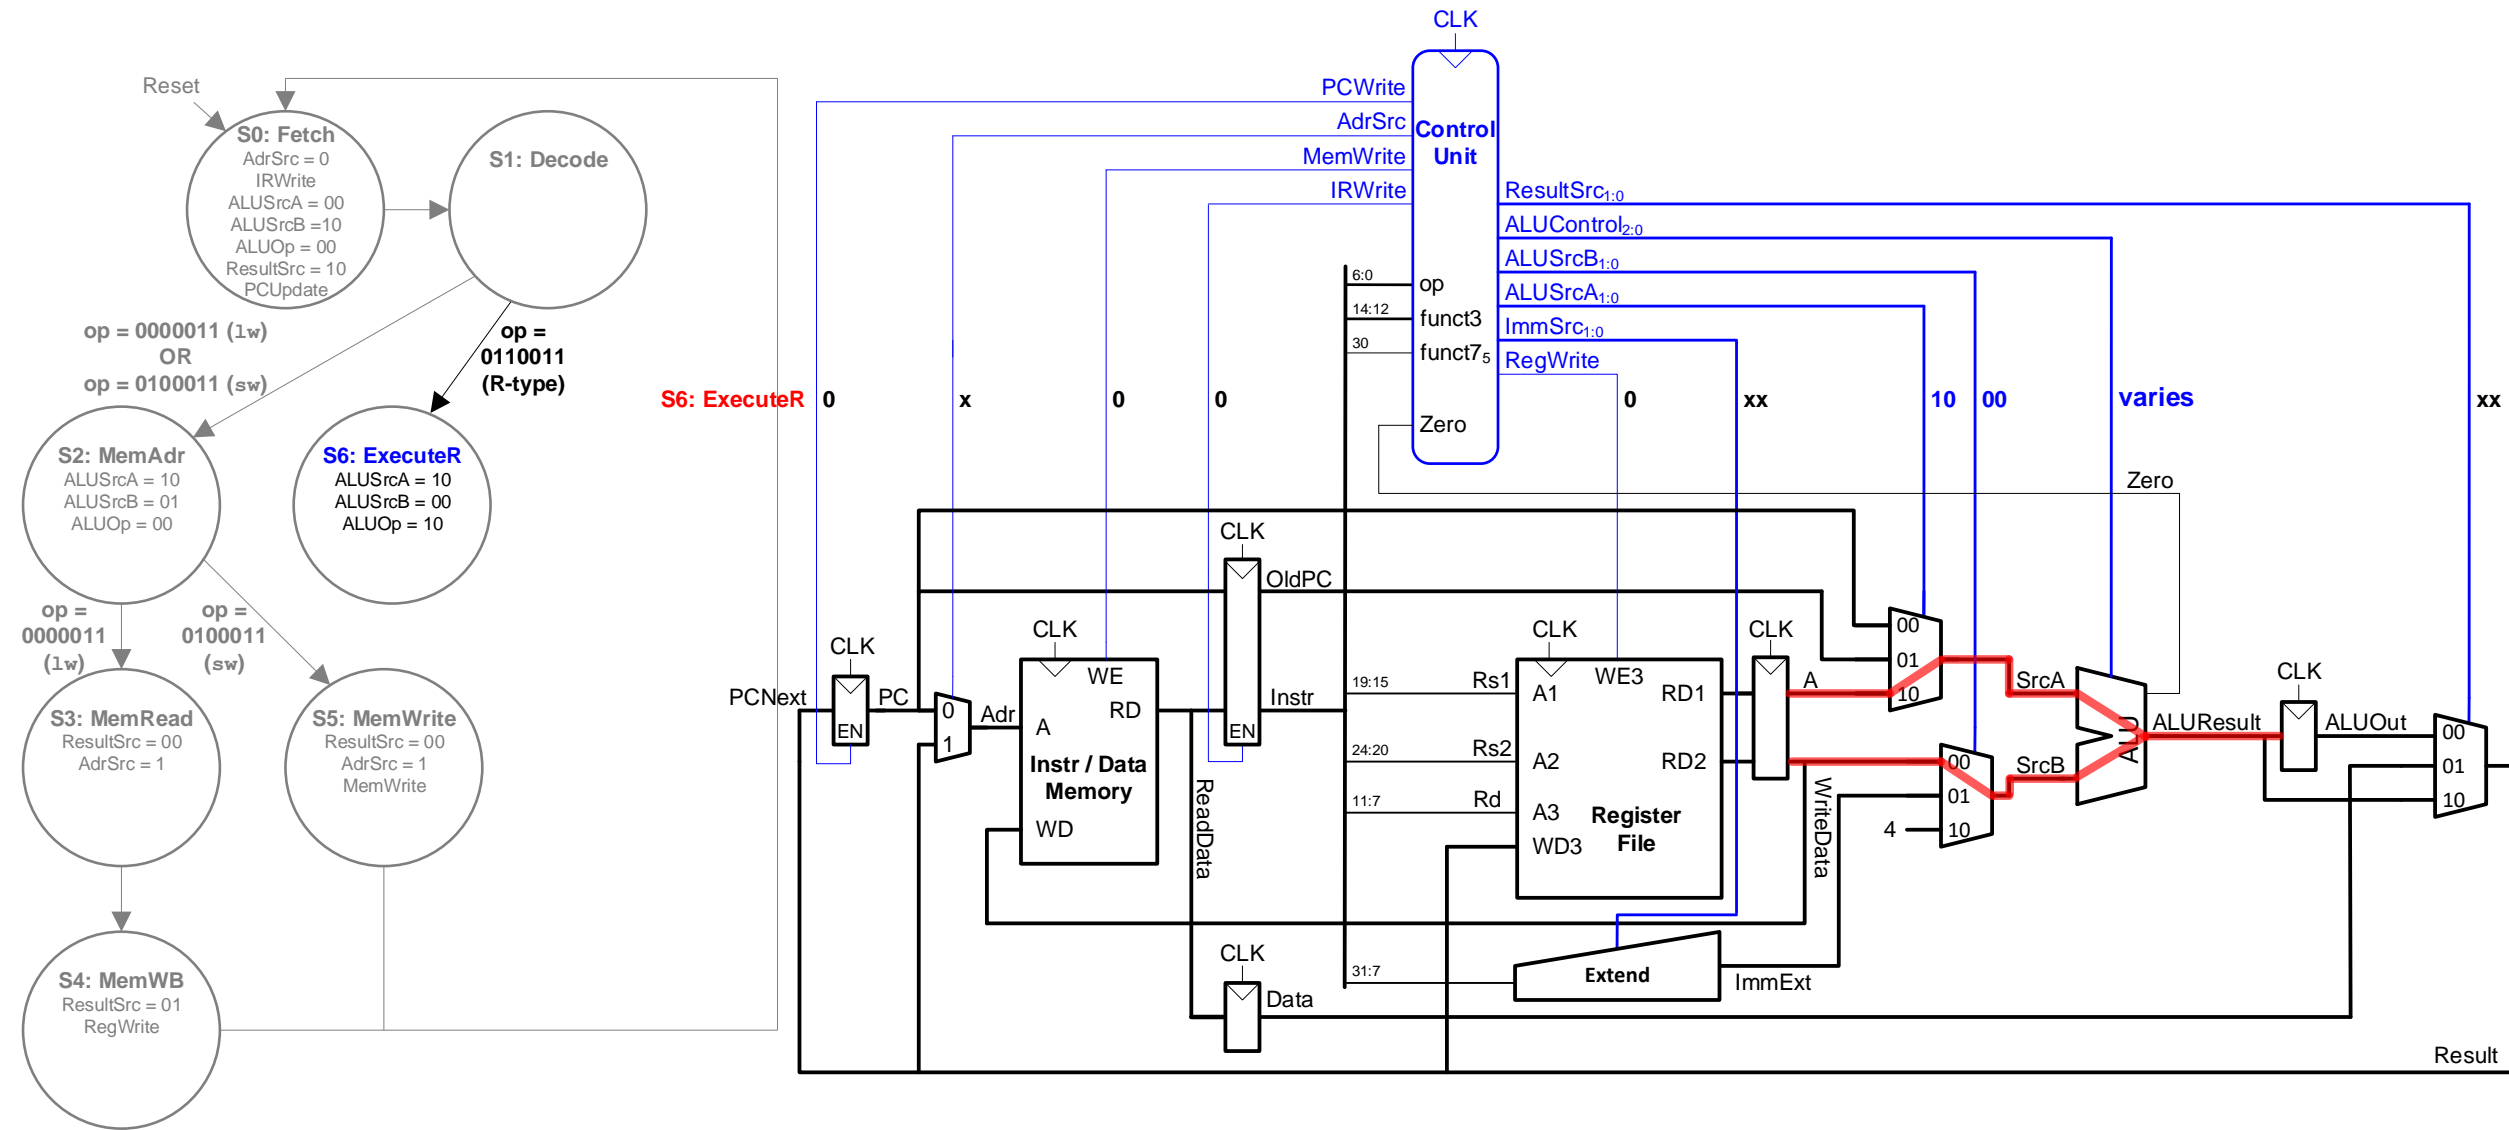
Computer Systems

75

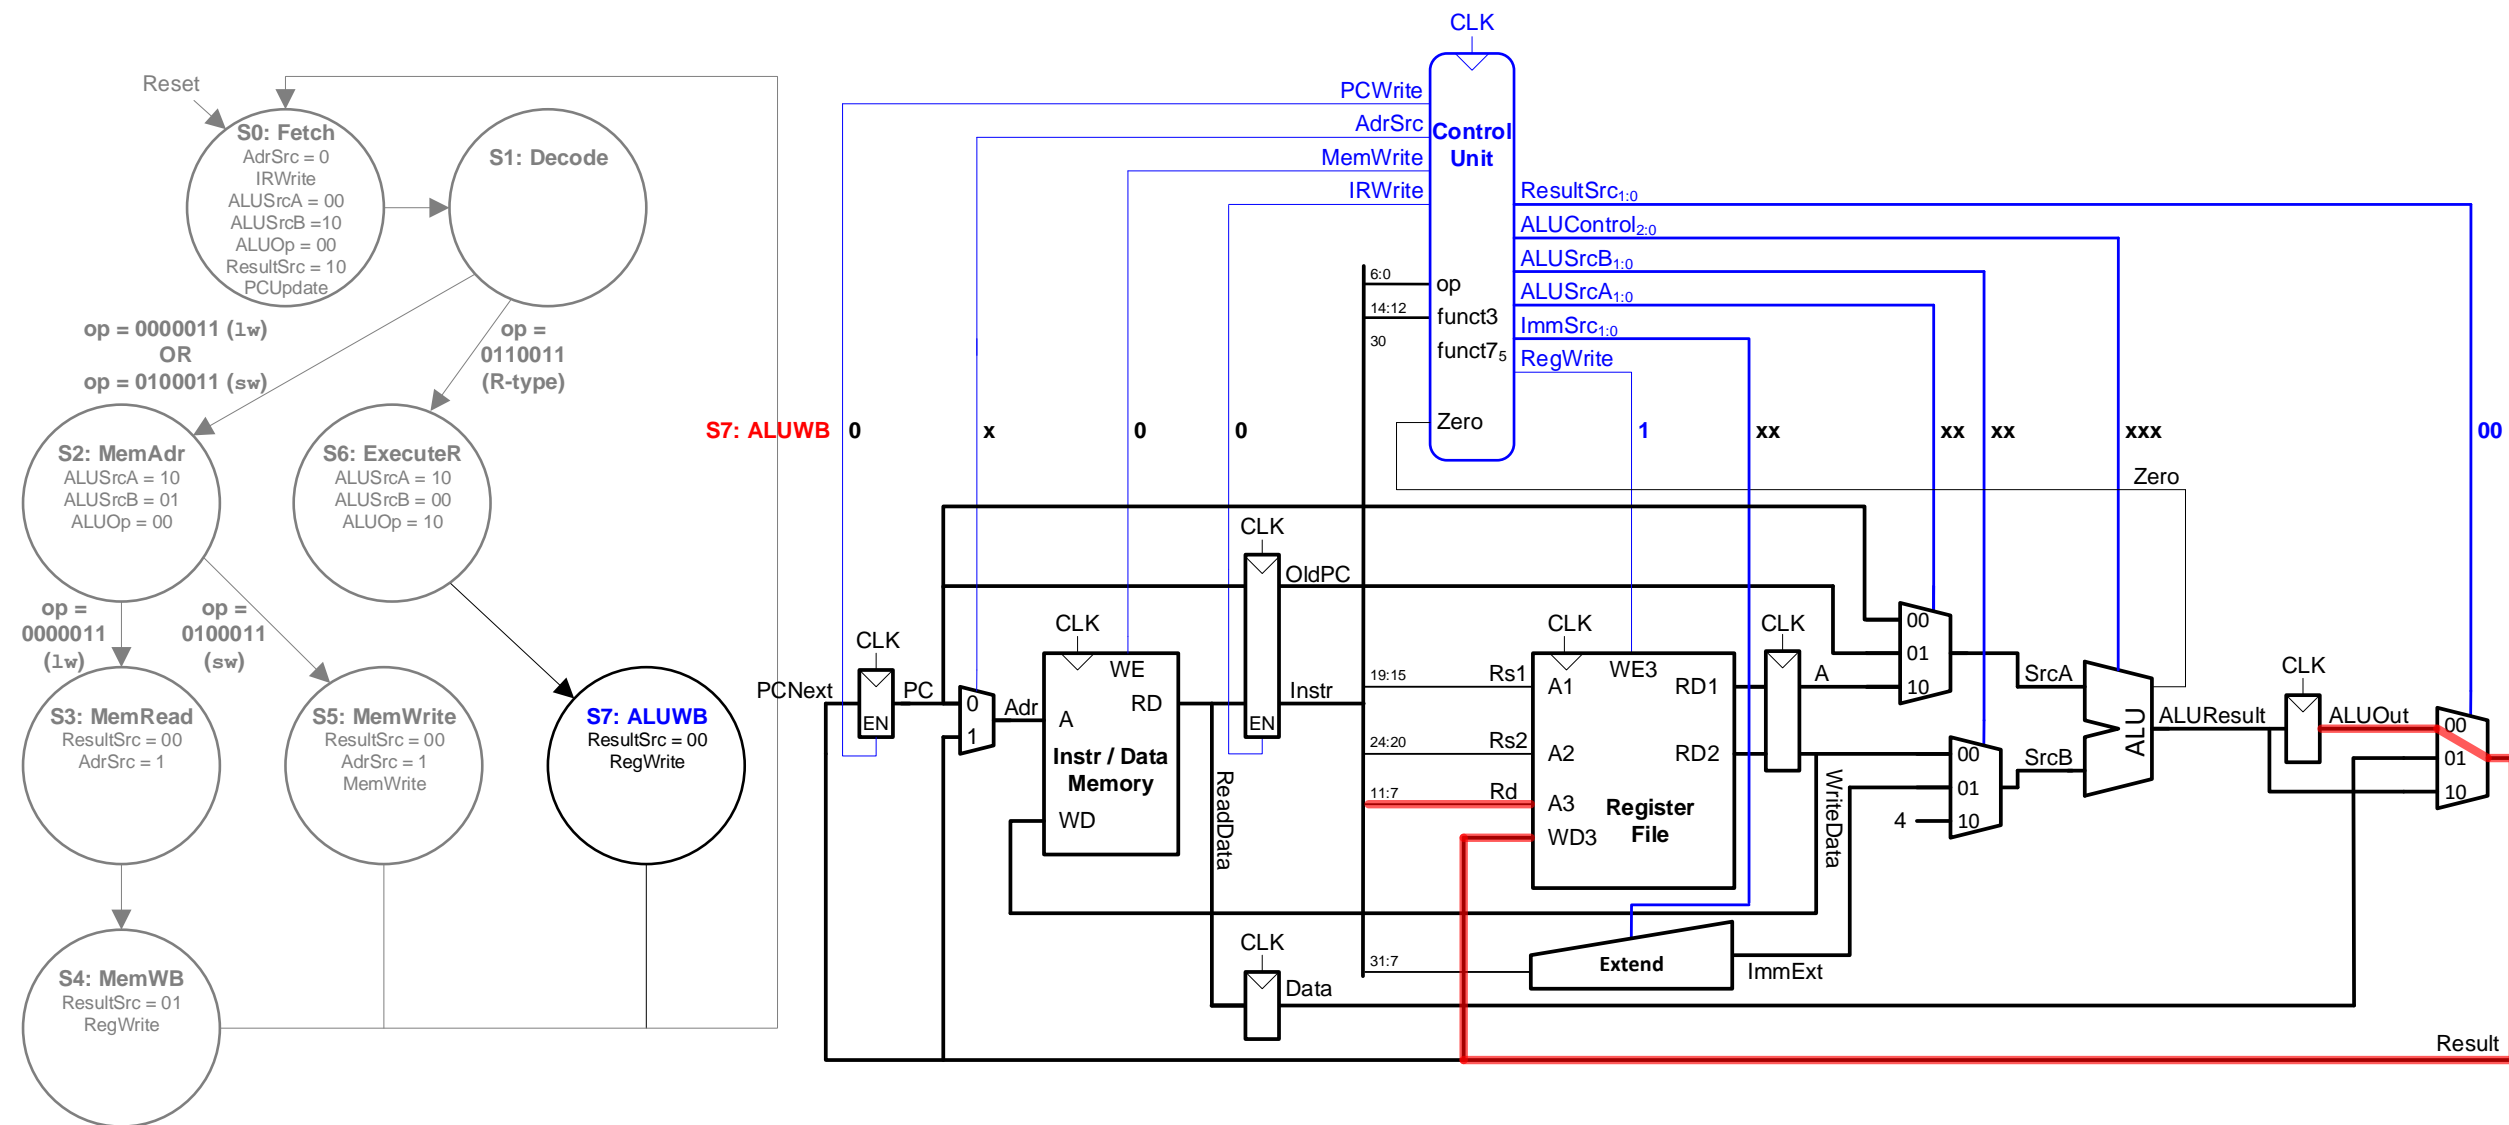
Main FSM: SW



Main FSM: R-Type Execute

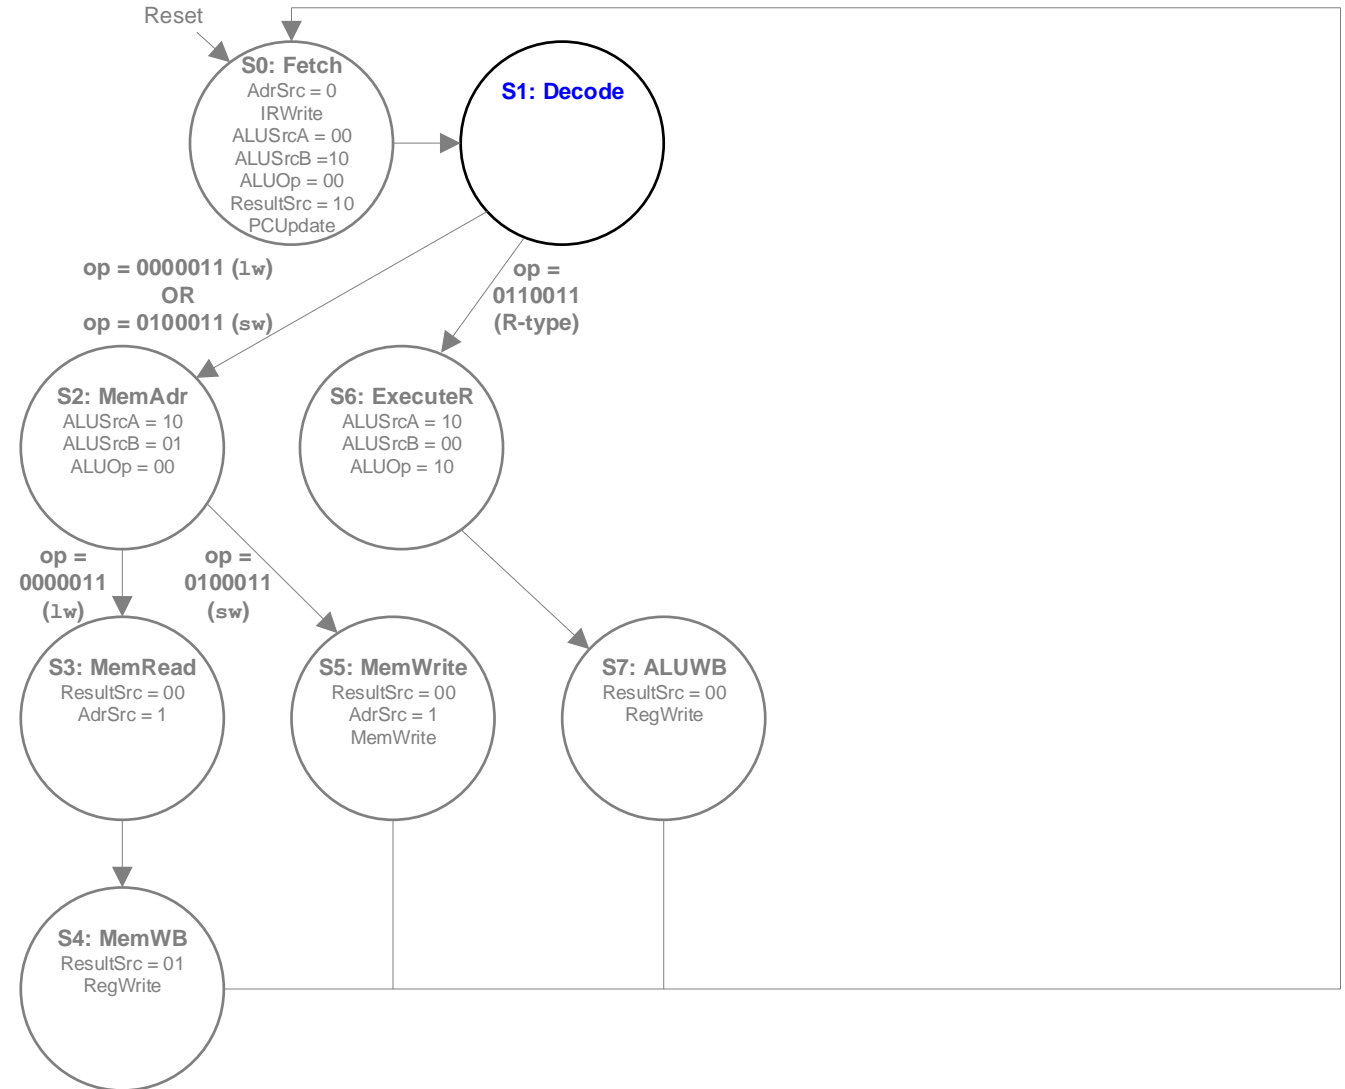


Main FSM: R-Type ALU Write Back



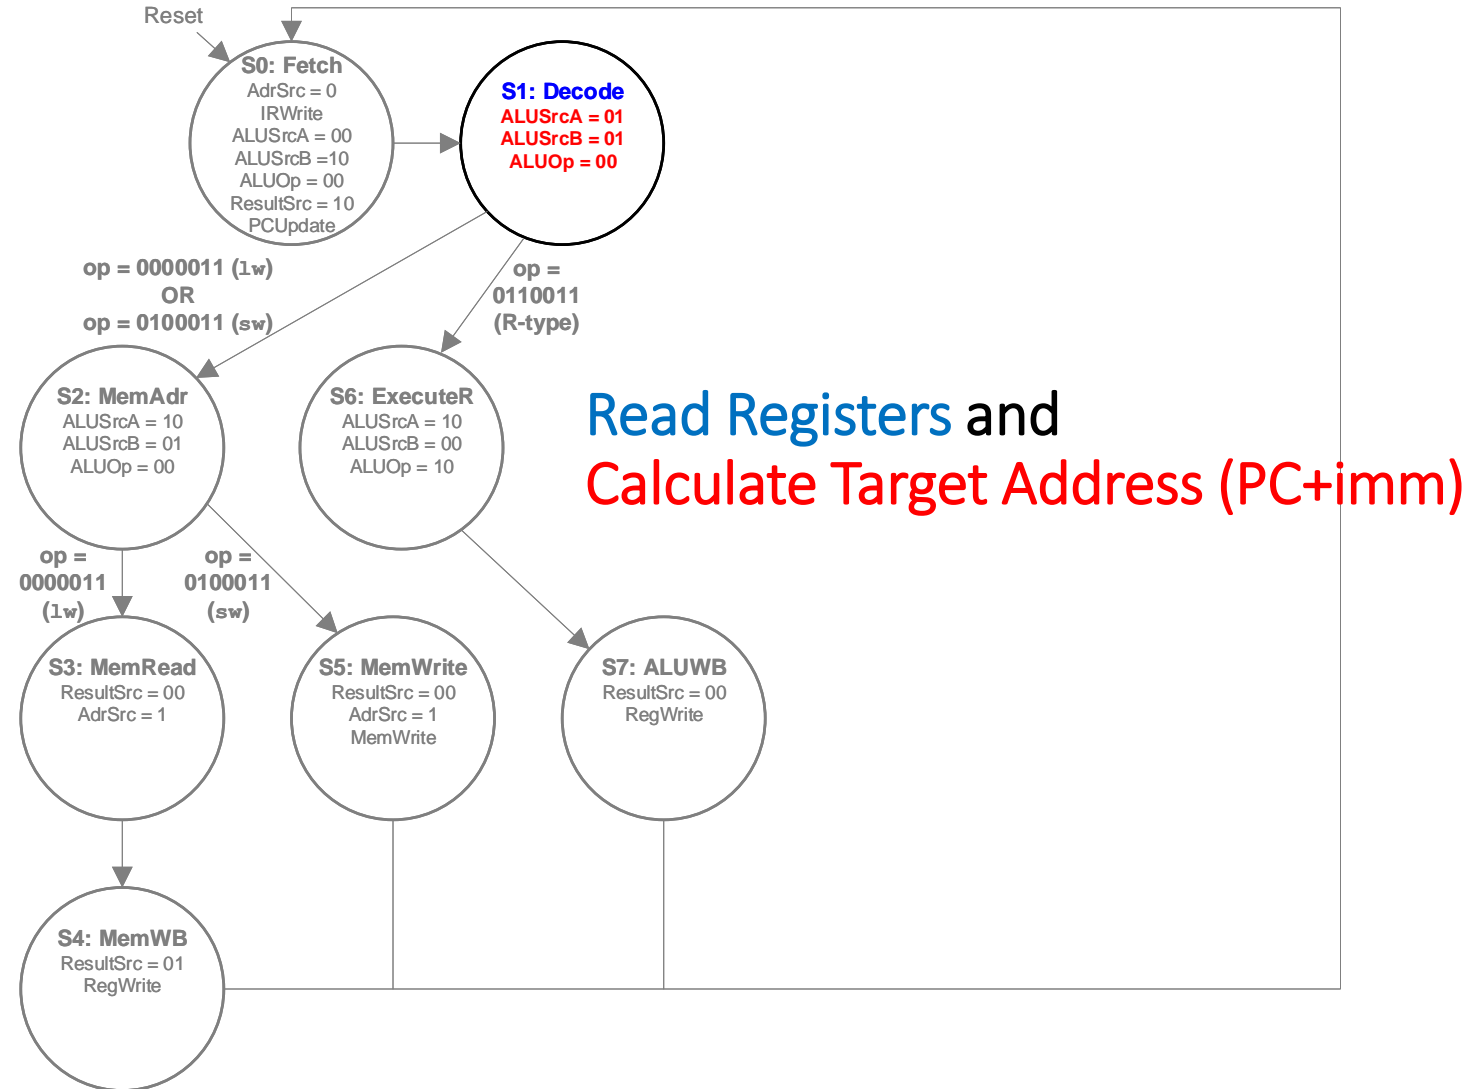
Main FSM: R-Type ALU Write Back

- **Need to calculate:**
 - Branch Target Address
 - **rs1 - rs2** (to see if equal)
- **ALU** isn't being used in Decode stage
 - Use it to calculate Target Address (PC + imm)

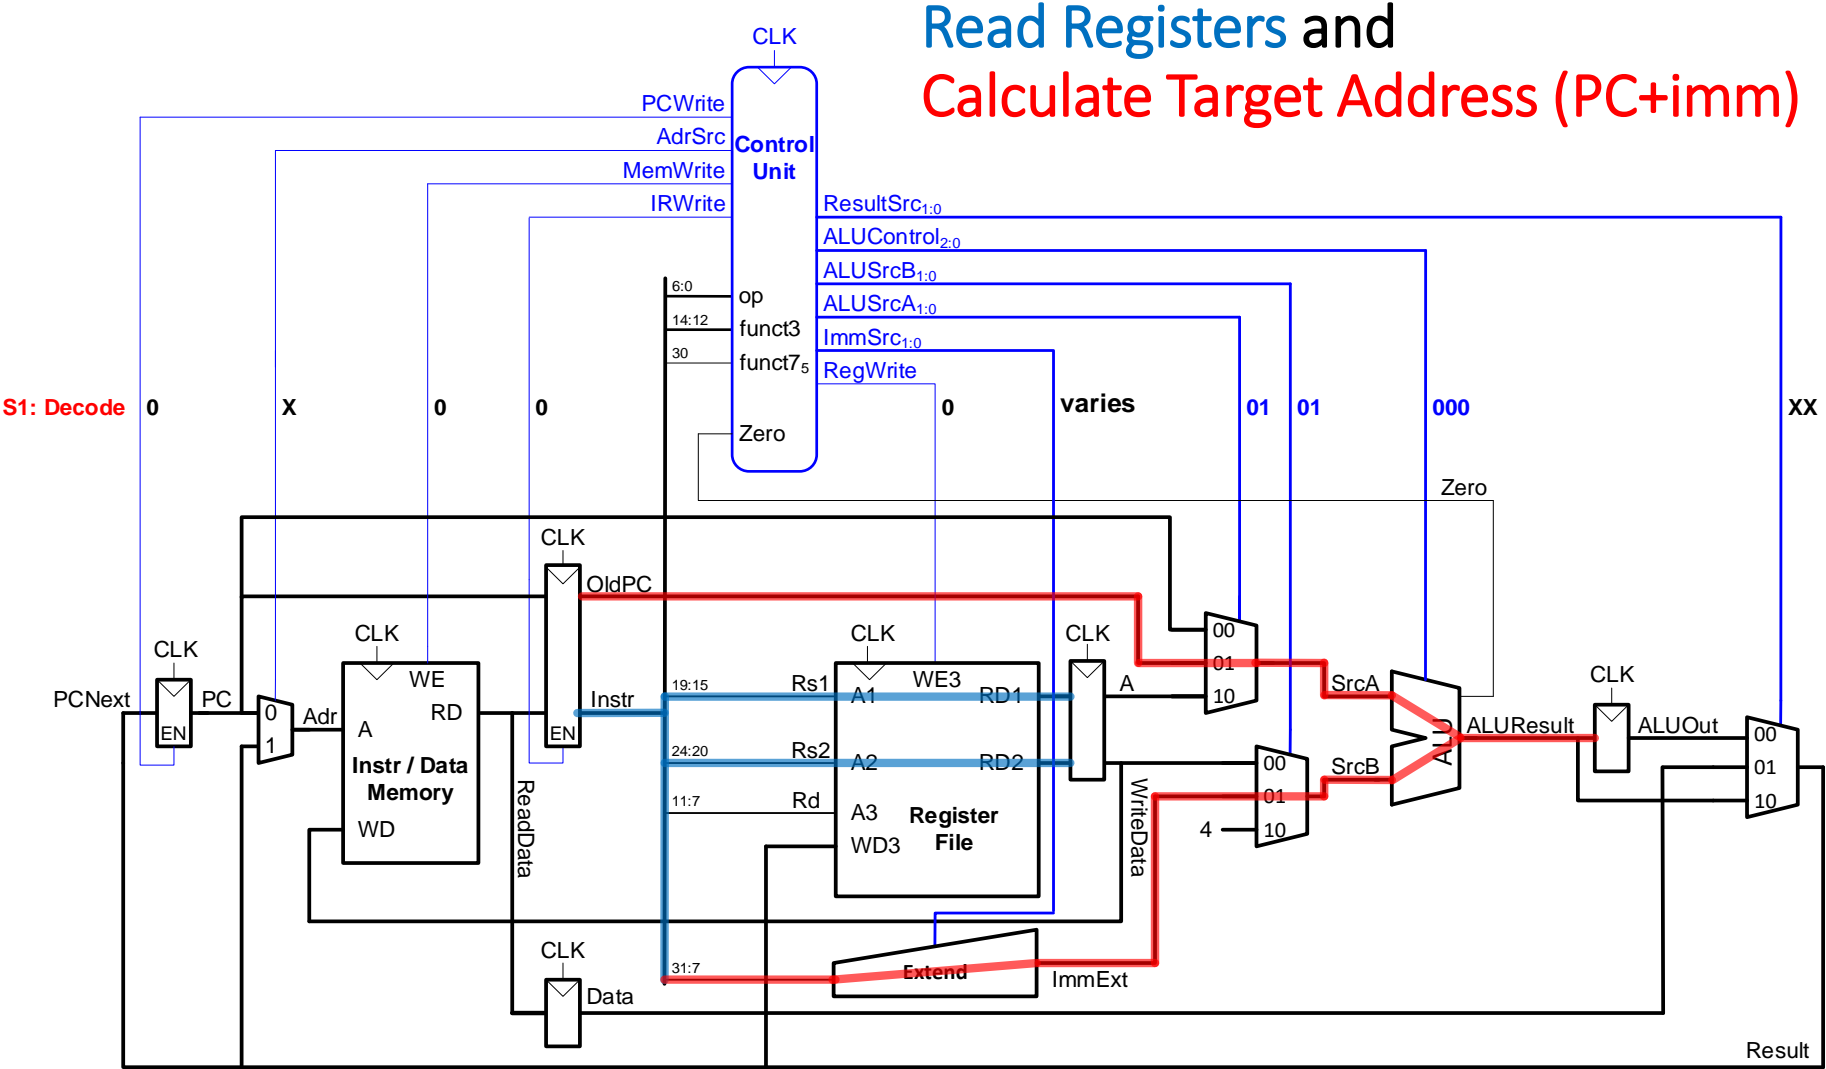
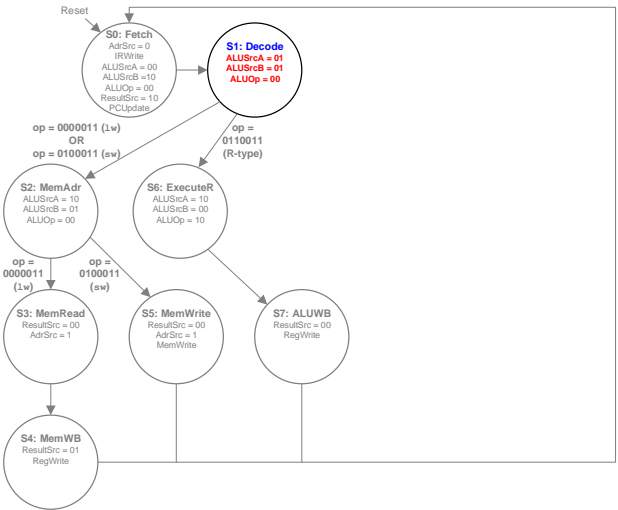
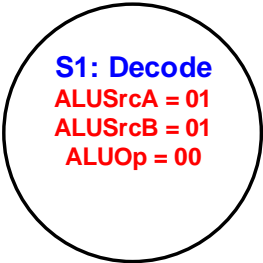


Main FSM: Decode Revisited

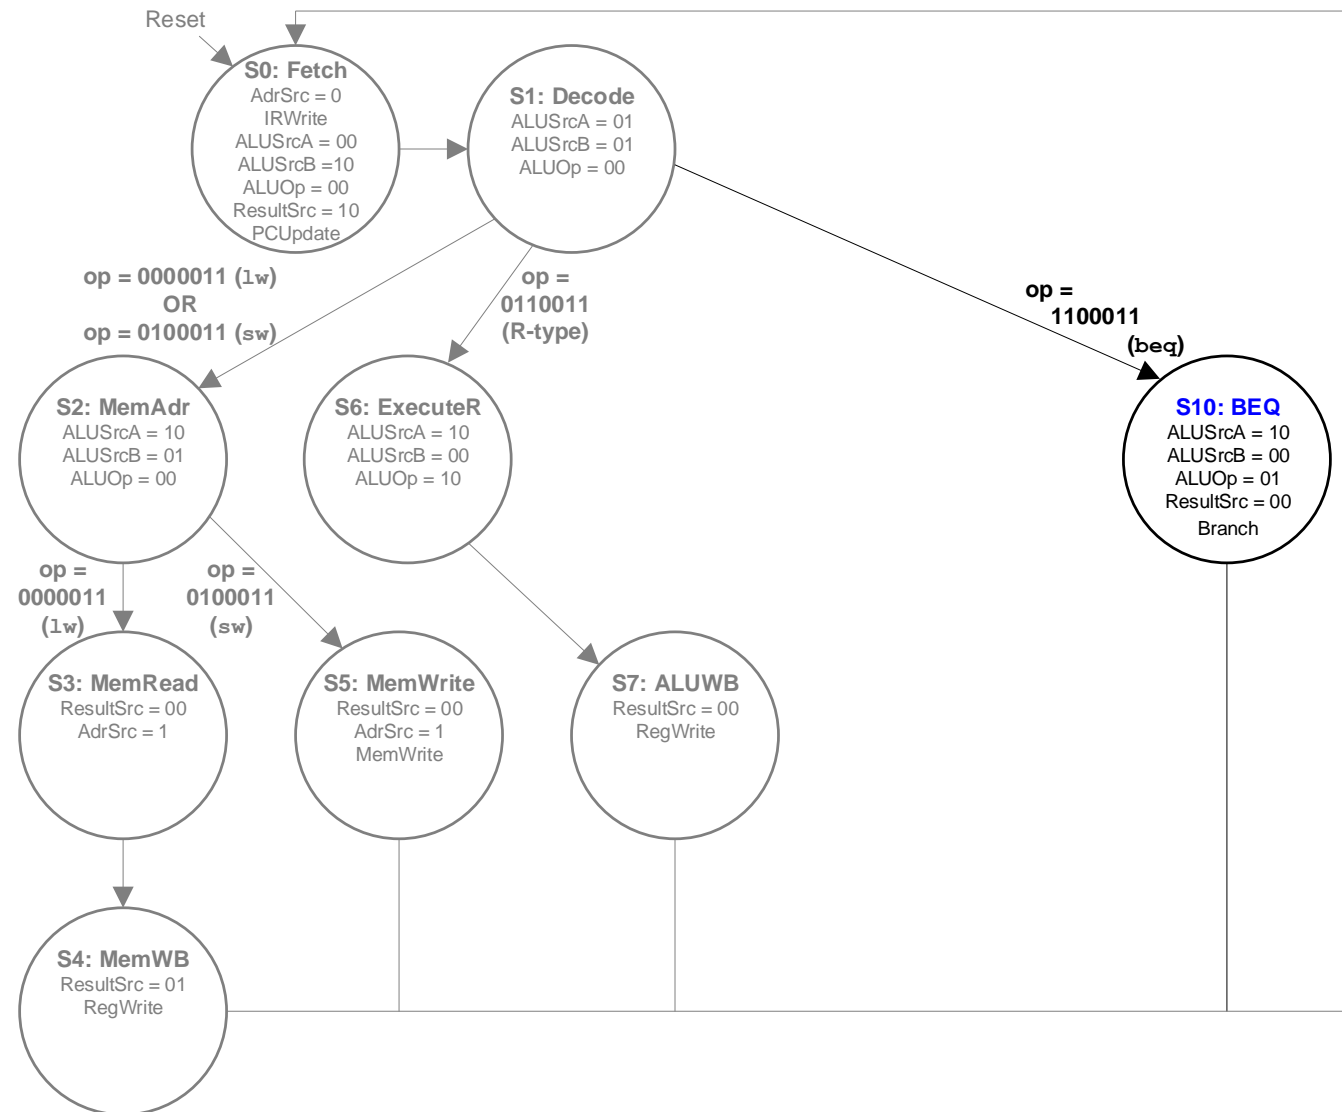
- **Need to calculate:**
 - Branch Target Address
 - **rs1 - rs2** (to see if equal)
- **ALU** isn't being used in Decode stage
 - Use it to calculate Target Address (PC + imm)



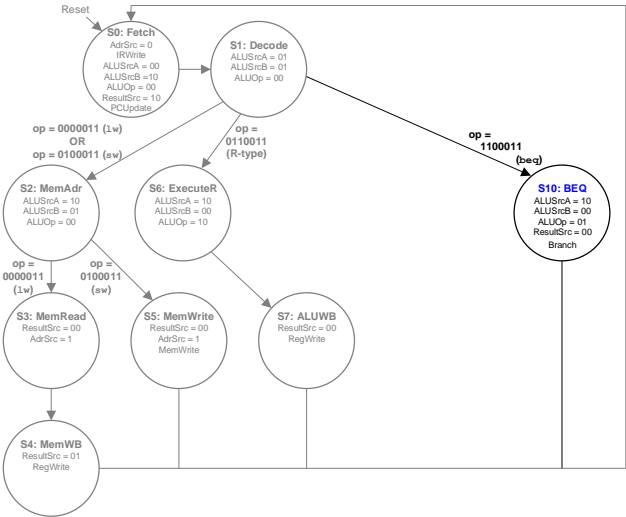

Main FSM: Decode (Target Address)



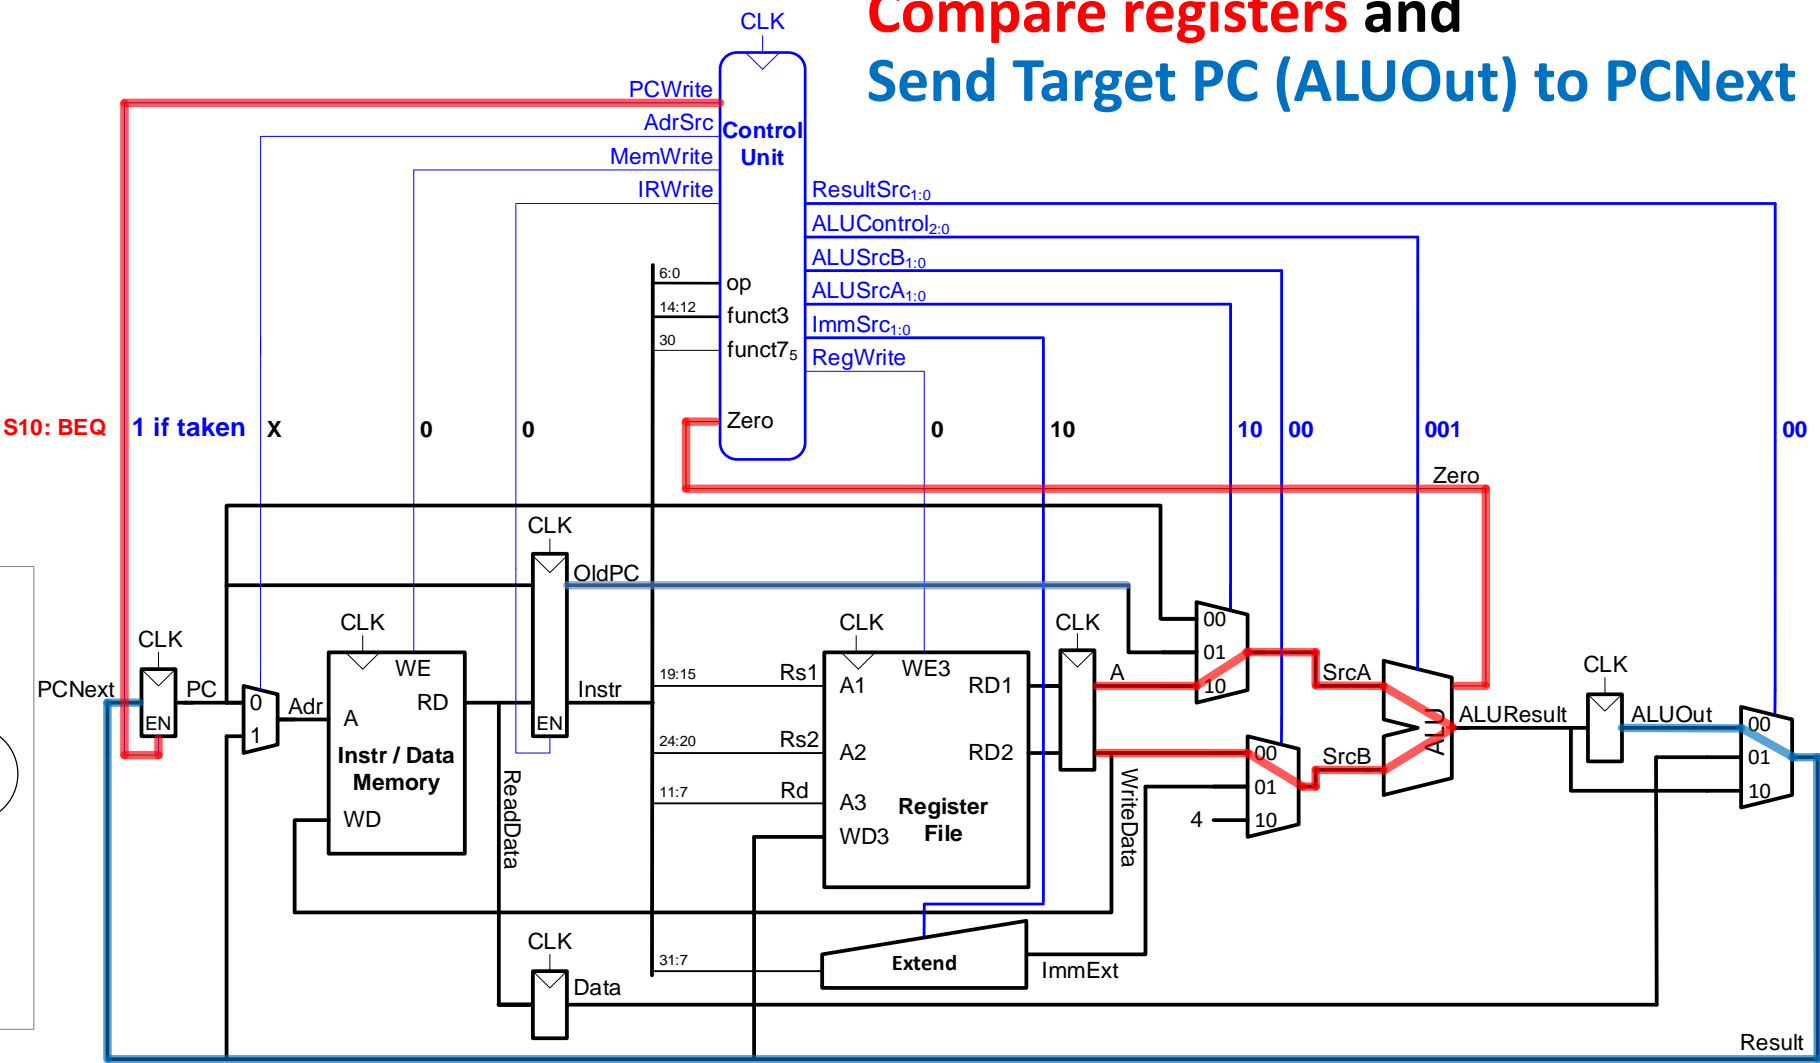
Main FSM: beq



Main FSM: beq Datapath



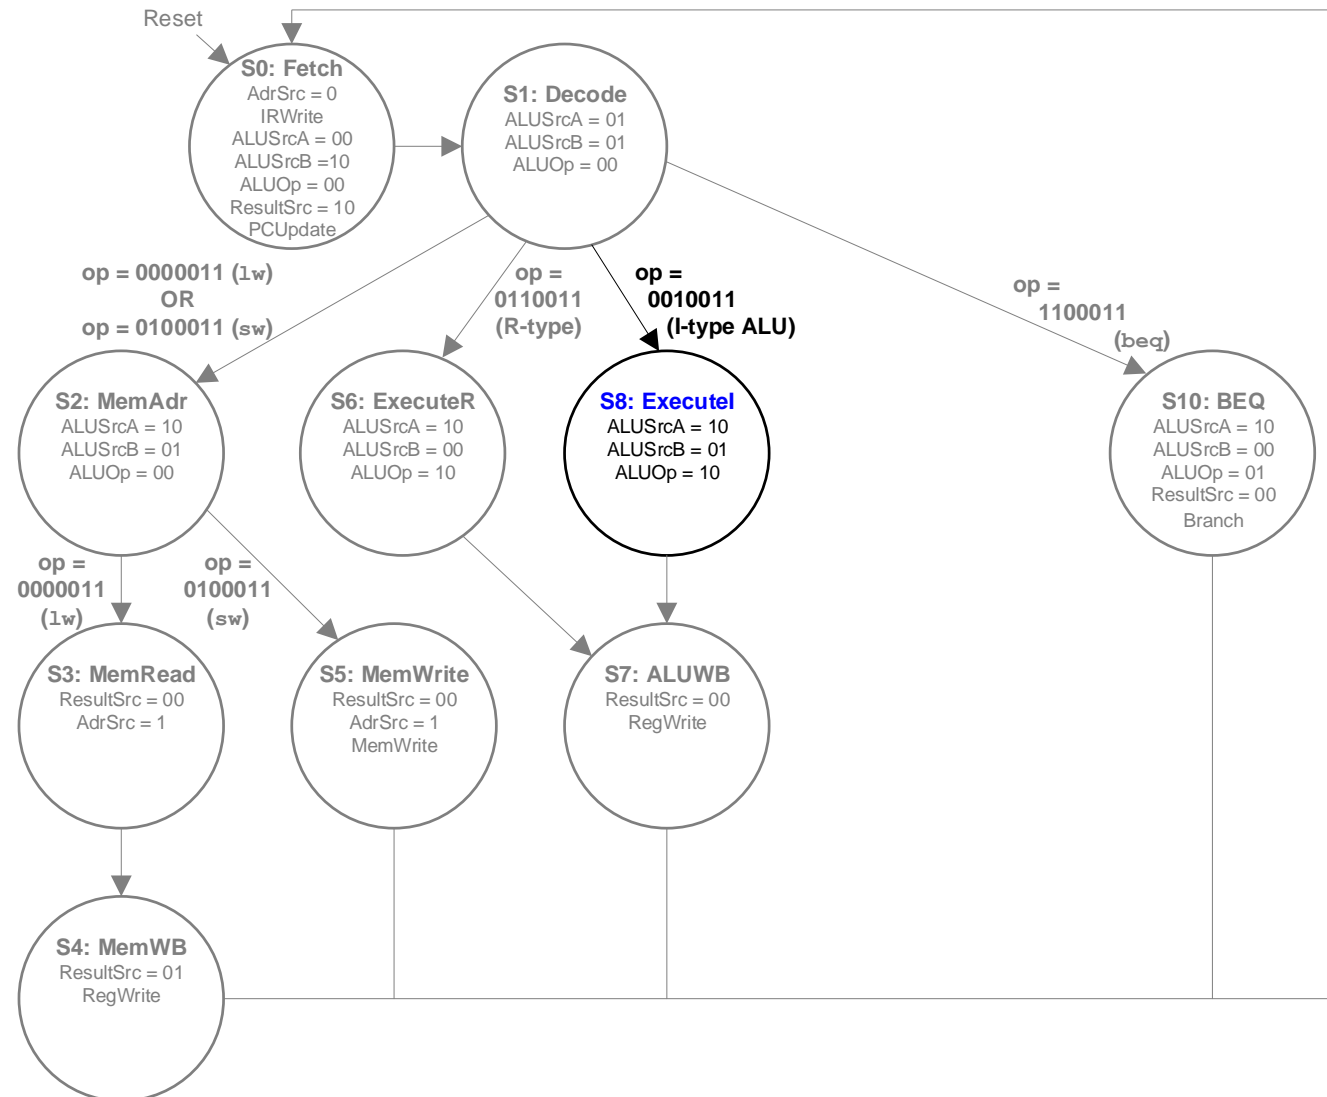
**Compare registers and
Send Target PC (ALUOut) to PCNext**



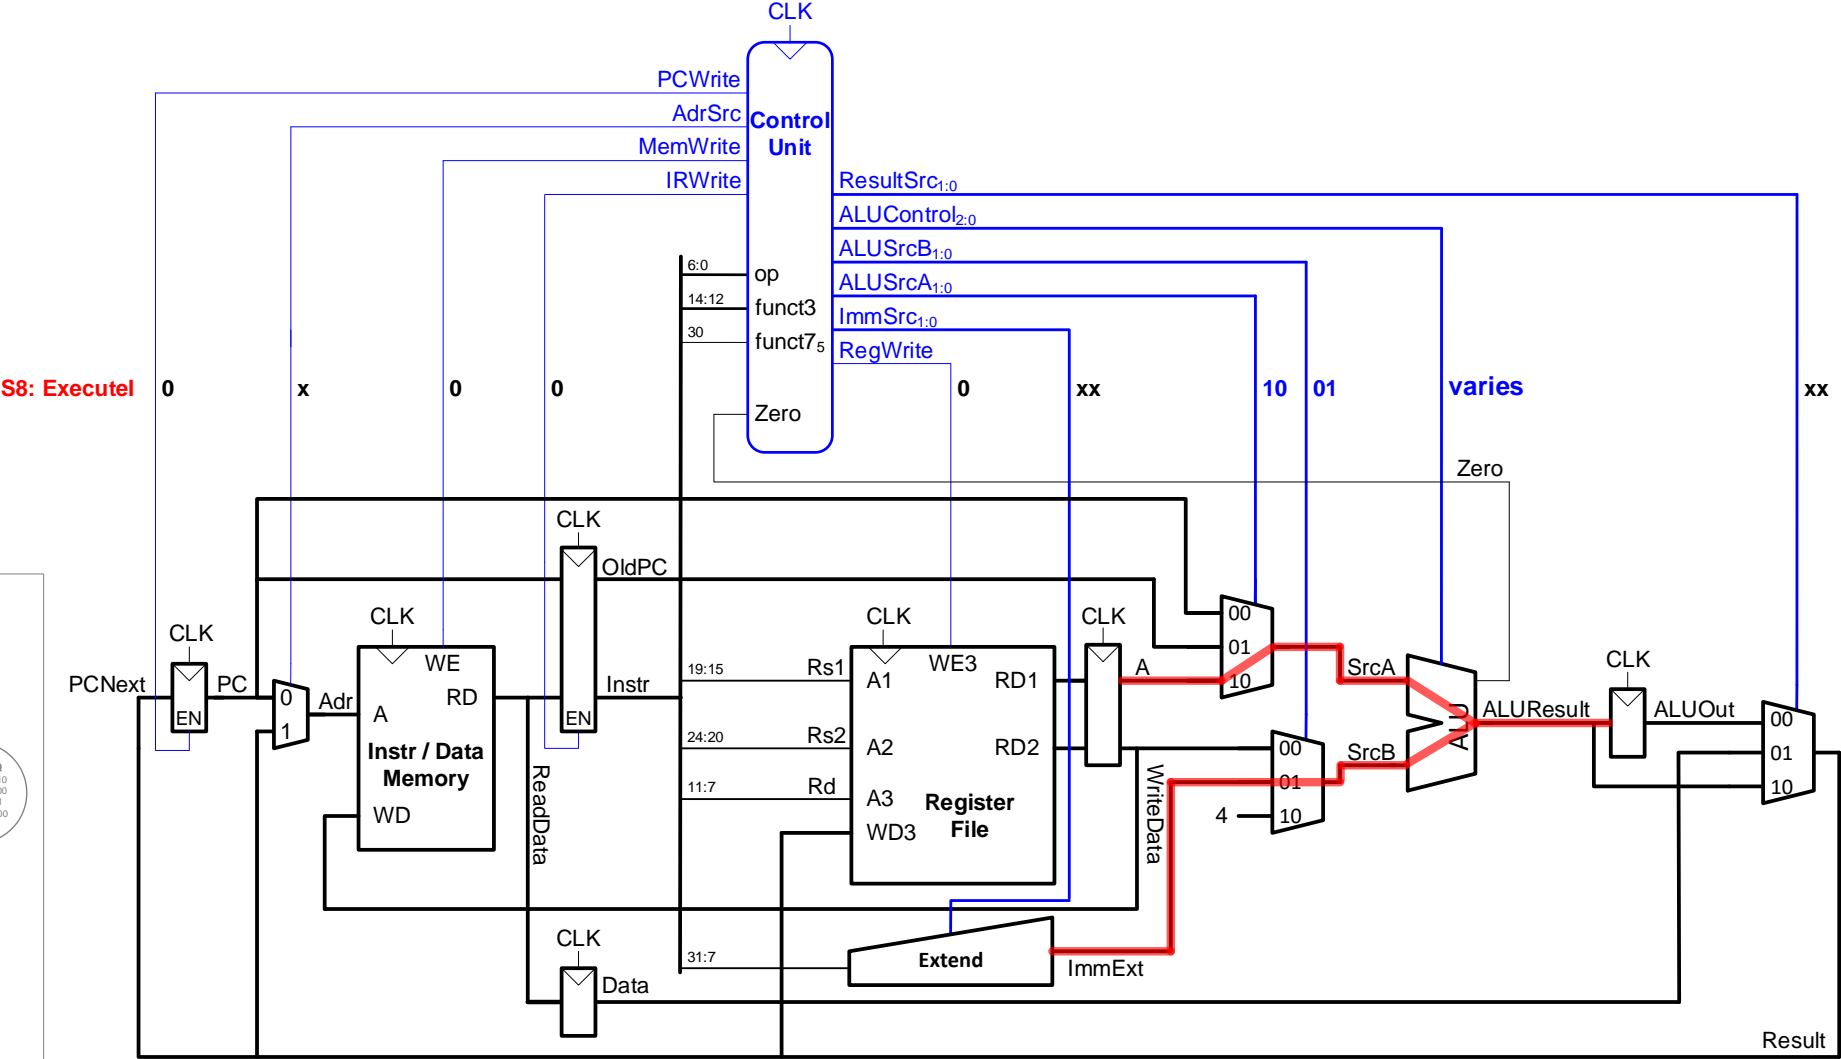
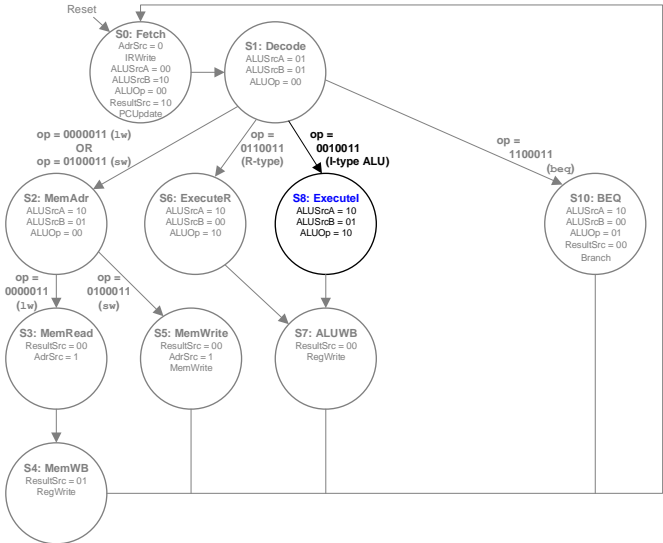

Extending the RISC-V Multicycle Processor

DDCA Ch7 - Part 11: Extending the RISC-V Multicycle Processor <https://www.youtube.com/watch?v=8EhVN192FRU>

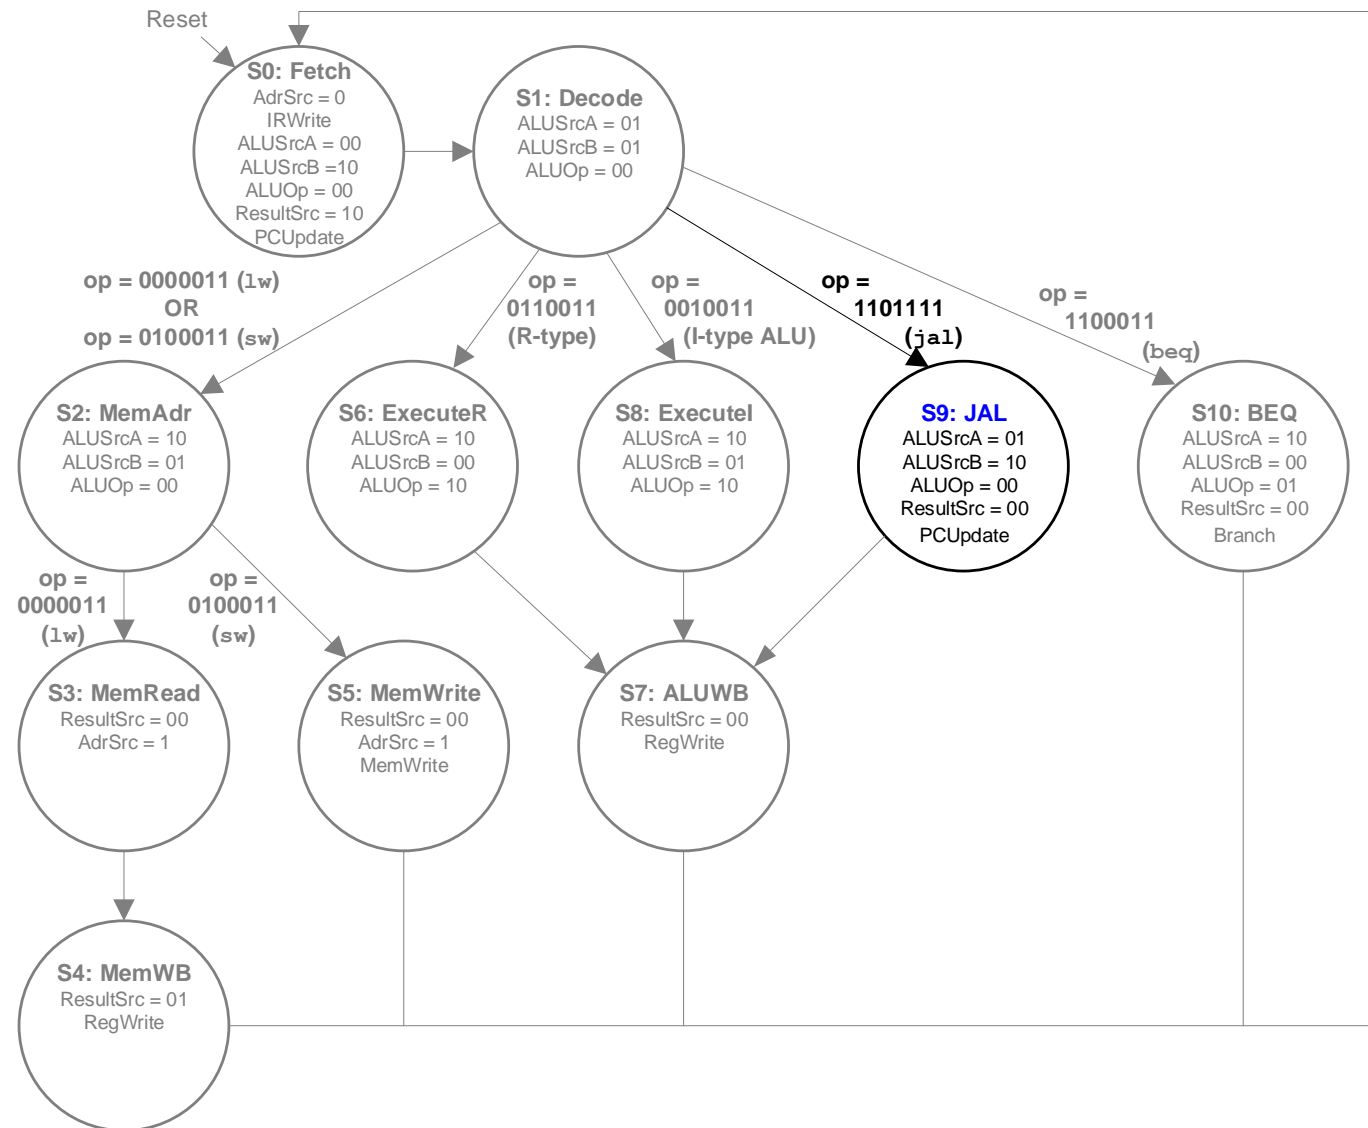
Main FSM: I-Type ALU Execute



Main FSM: I-Type ALU Exec. Datapath



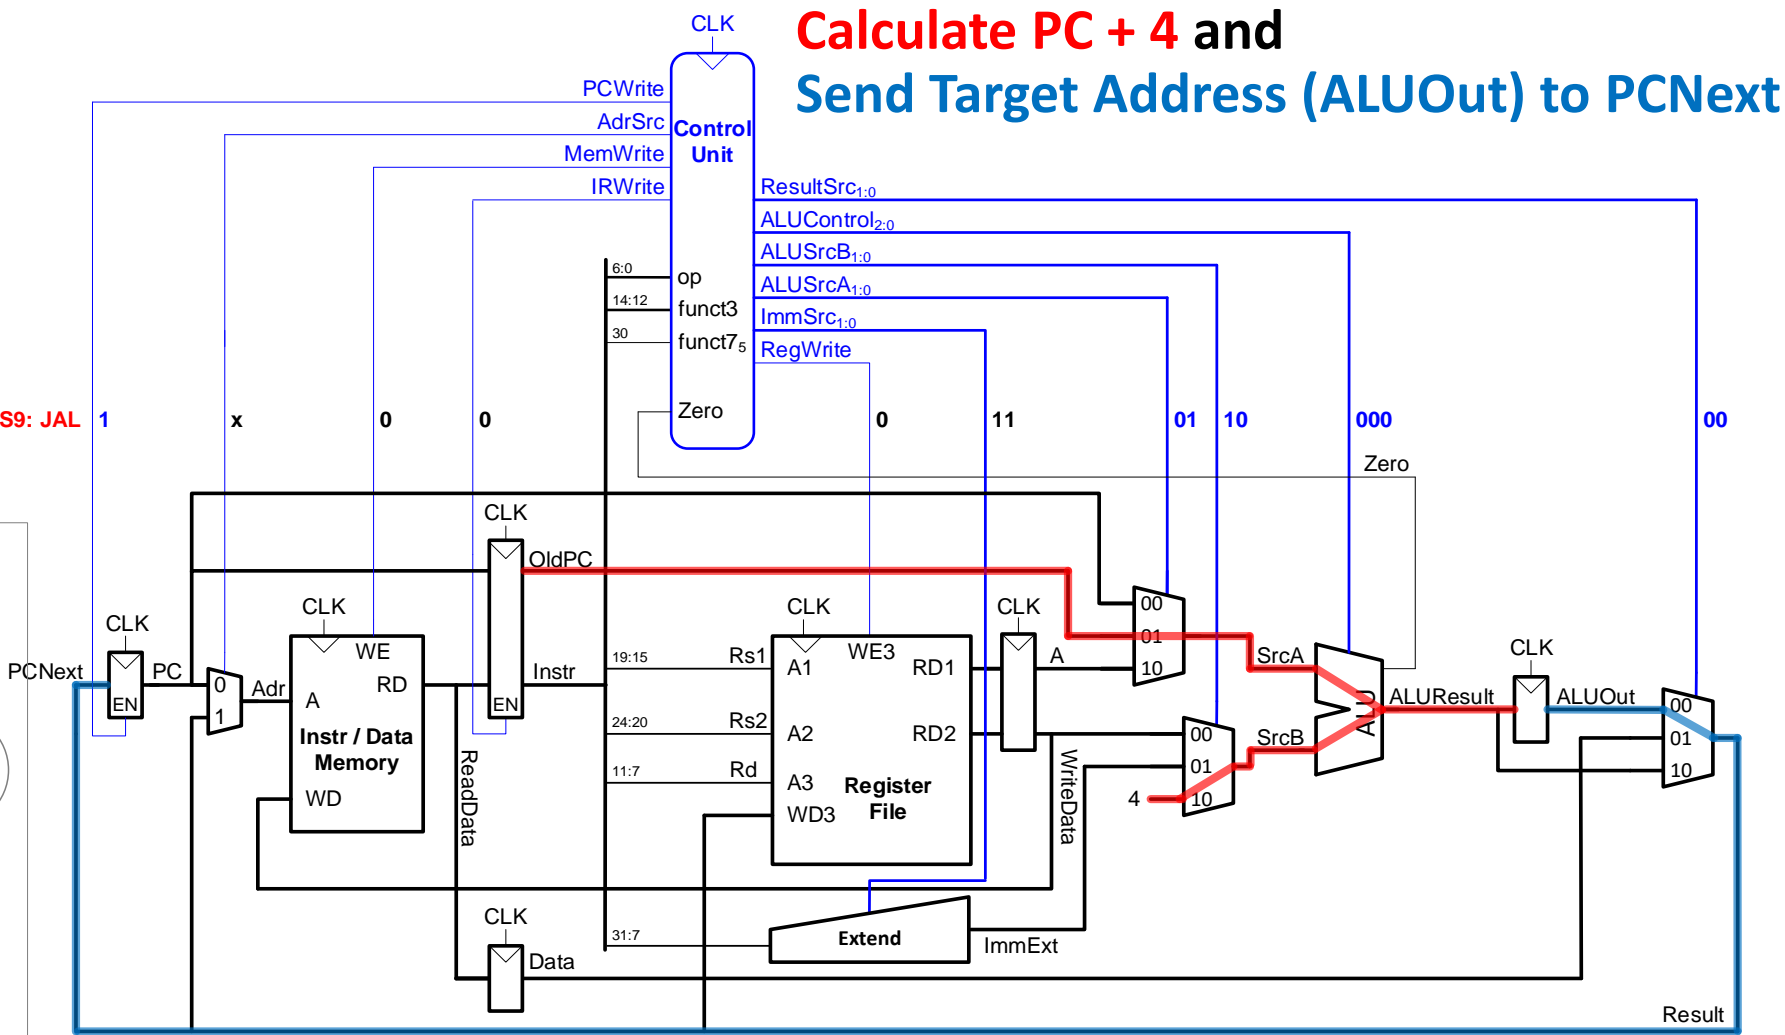
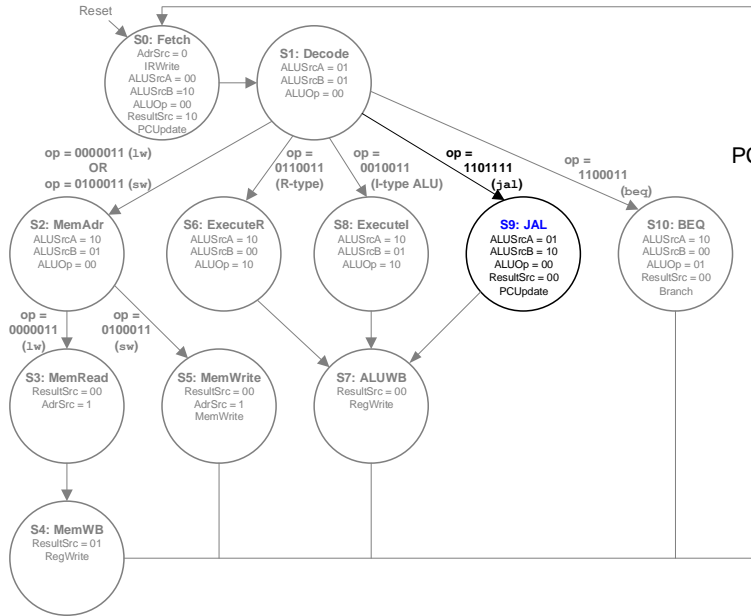
Main FSM: jal



Main FSM: jal Datapath

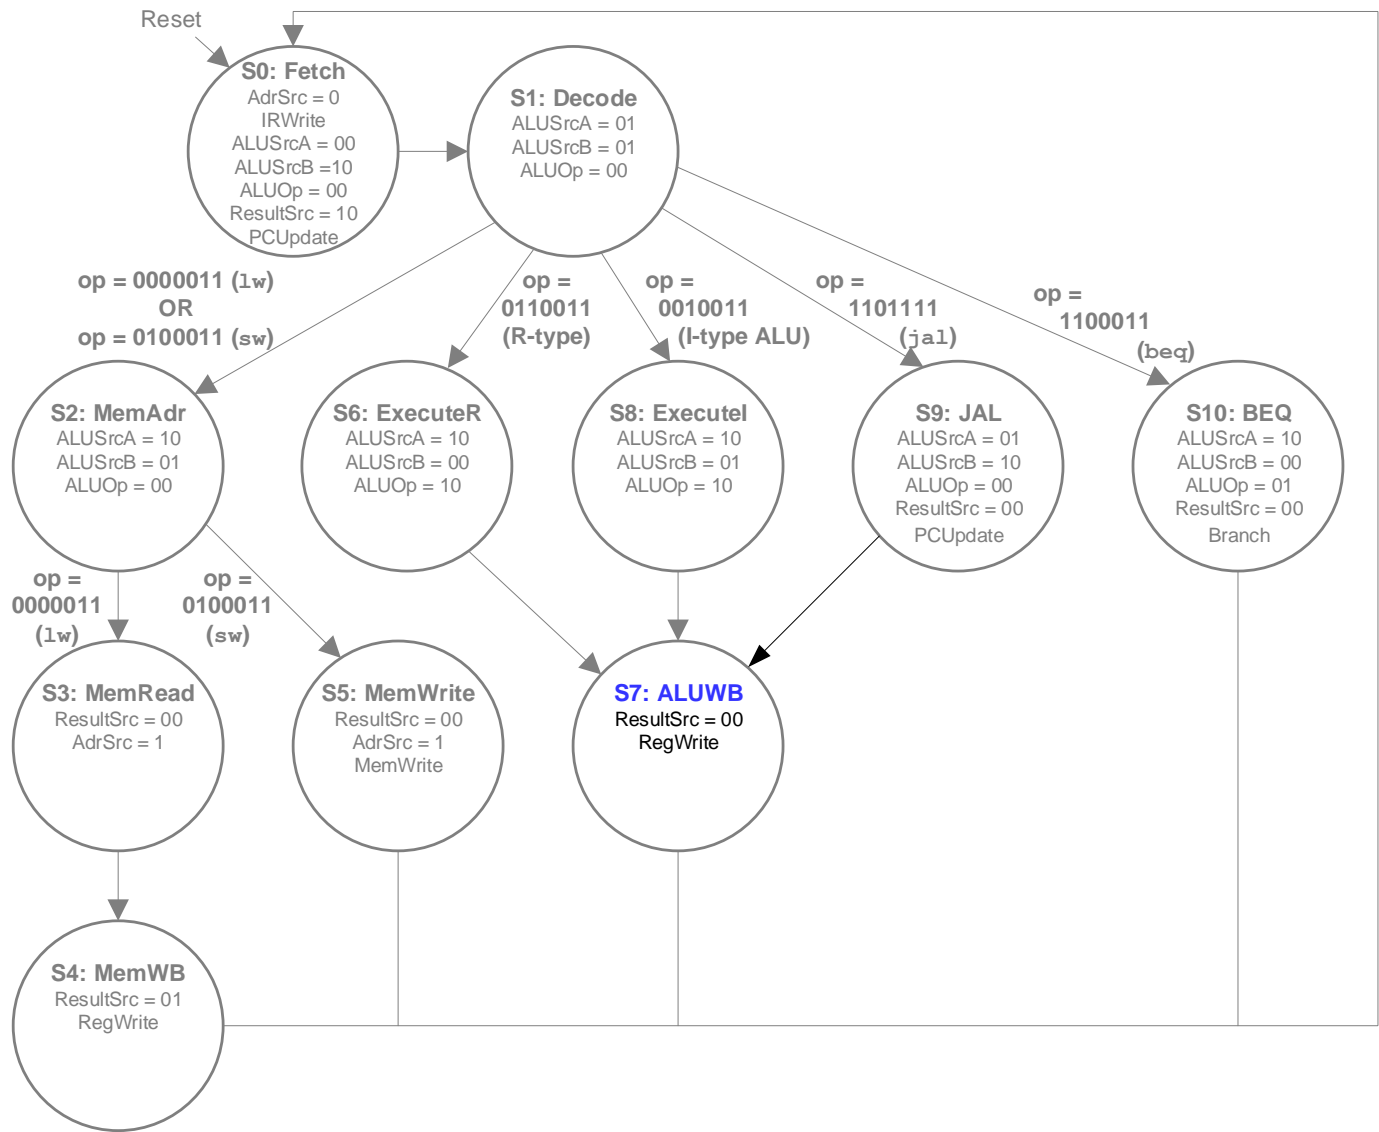
S9: JAL

ALUSrcA = 01
ALUSrcB = 10
ALUOp = 00
ResultSrc = 00
PCUpdate

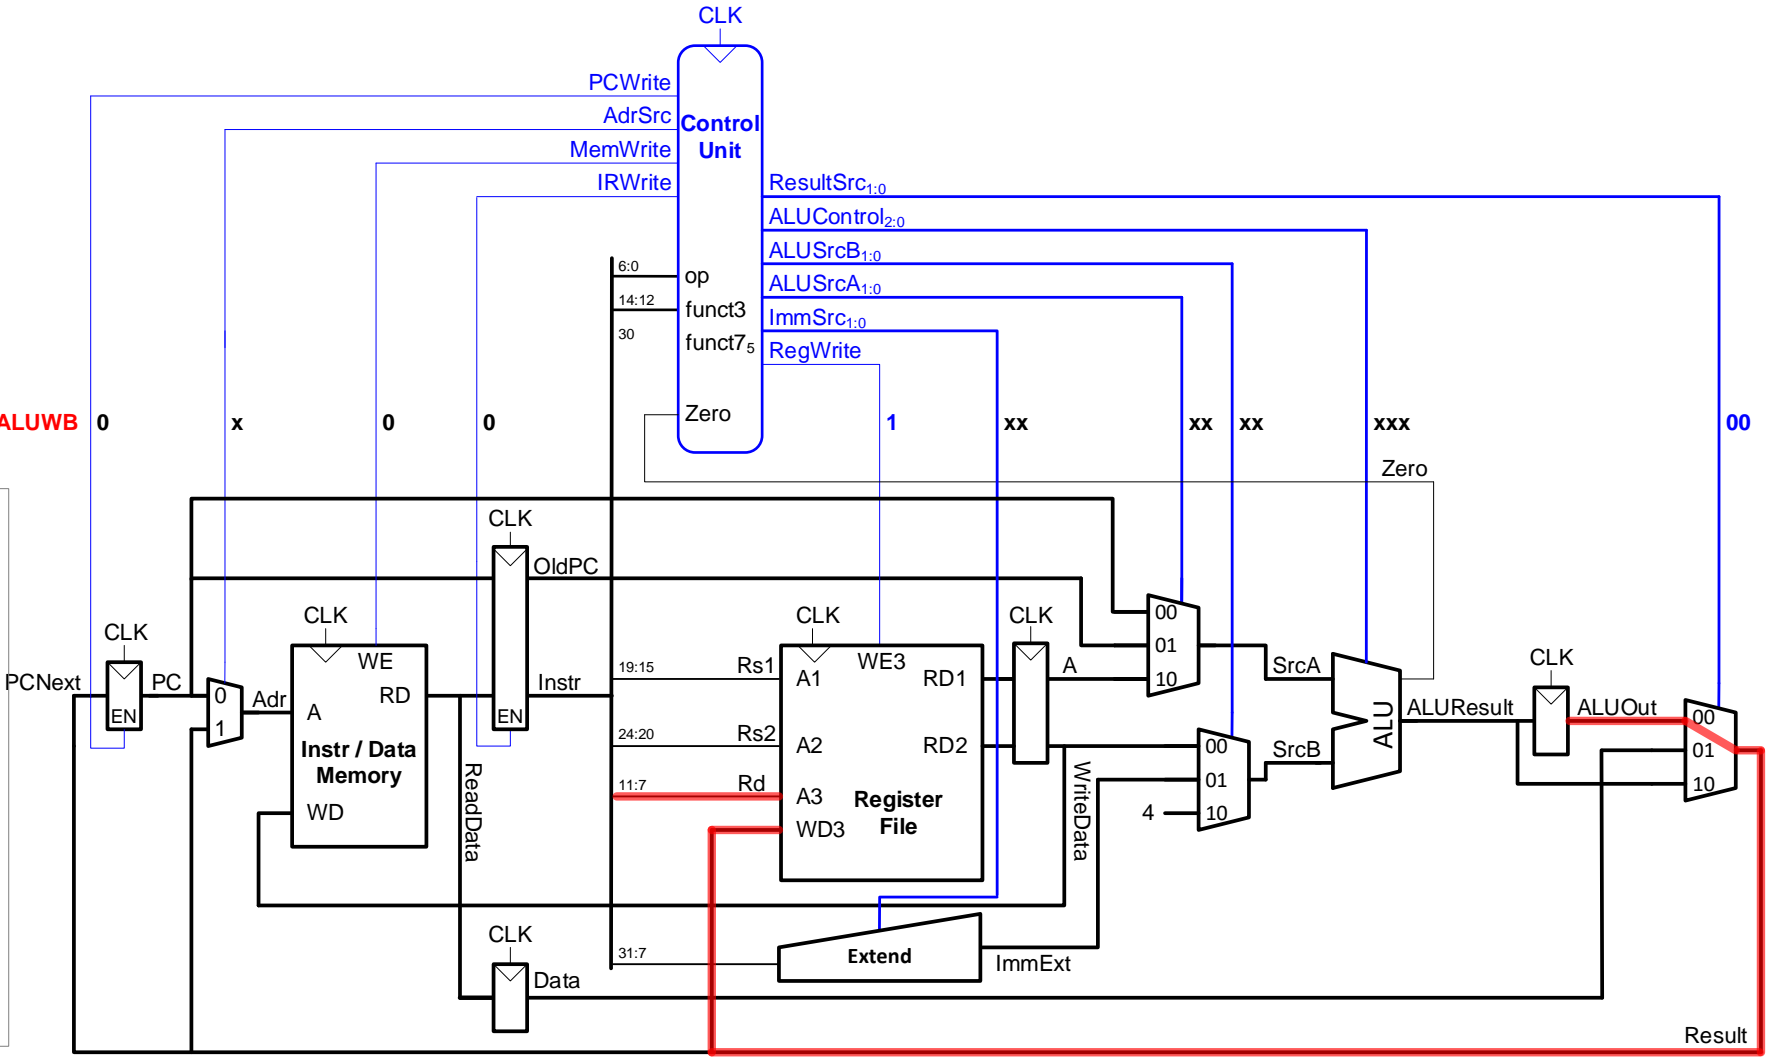
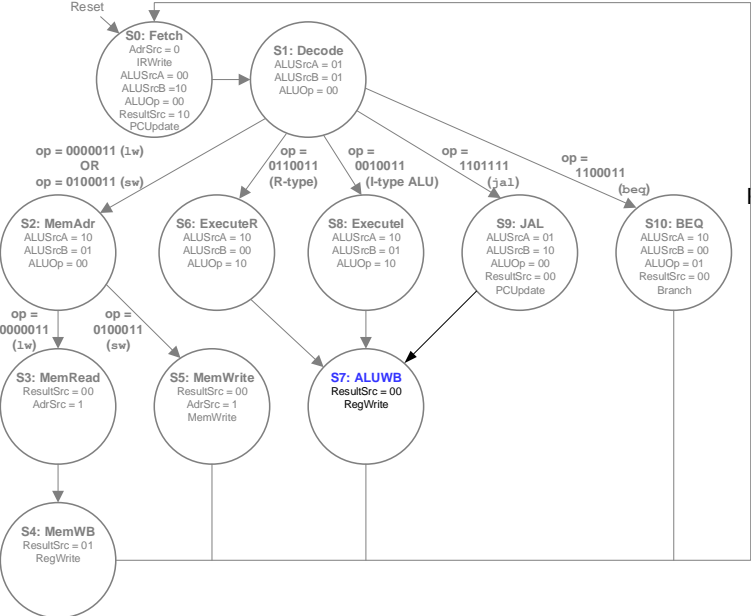


Main FSM: jal

PC + 4 is
written to rd
in S7: ALUWB



Main FSM: jal Datapath – WB same as before with PC+4



State

Fetch

Decode

MemAdr

MemRead

MemWB

MemWrite

ExecuteR

Executel

ALUWB

BEQ

JAL

Datapath μOp

$\text{Instr} \leftarrow \text{Mem}[\text{PC}]; \text{PC} \leftarrow \text{PC}+4$

$\text{ALUOut} \leftarrow \text{PCTarget}$

$\text{ALUOut} \leftarrow \text{rs1} + \text{imm}$

$\text{Data} \leftarrow \text{Mem}[\text{ALUOut}]$

$\text{rd} \leftarrow \text{Data}$

$\text{Mem}[\text{ALUOut}] \leftarrow \text{rd}$

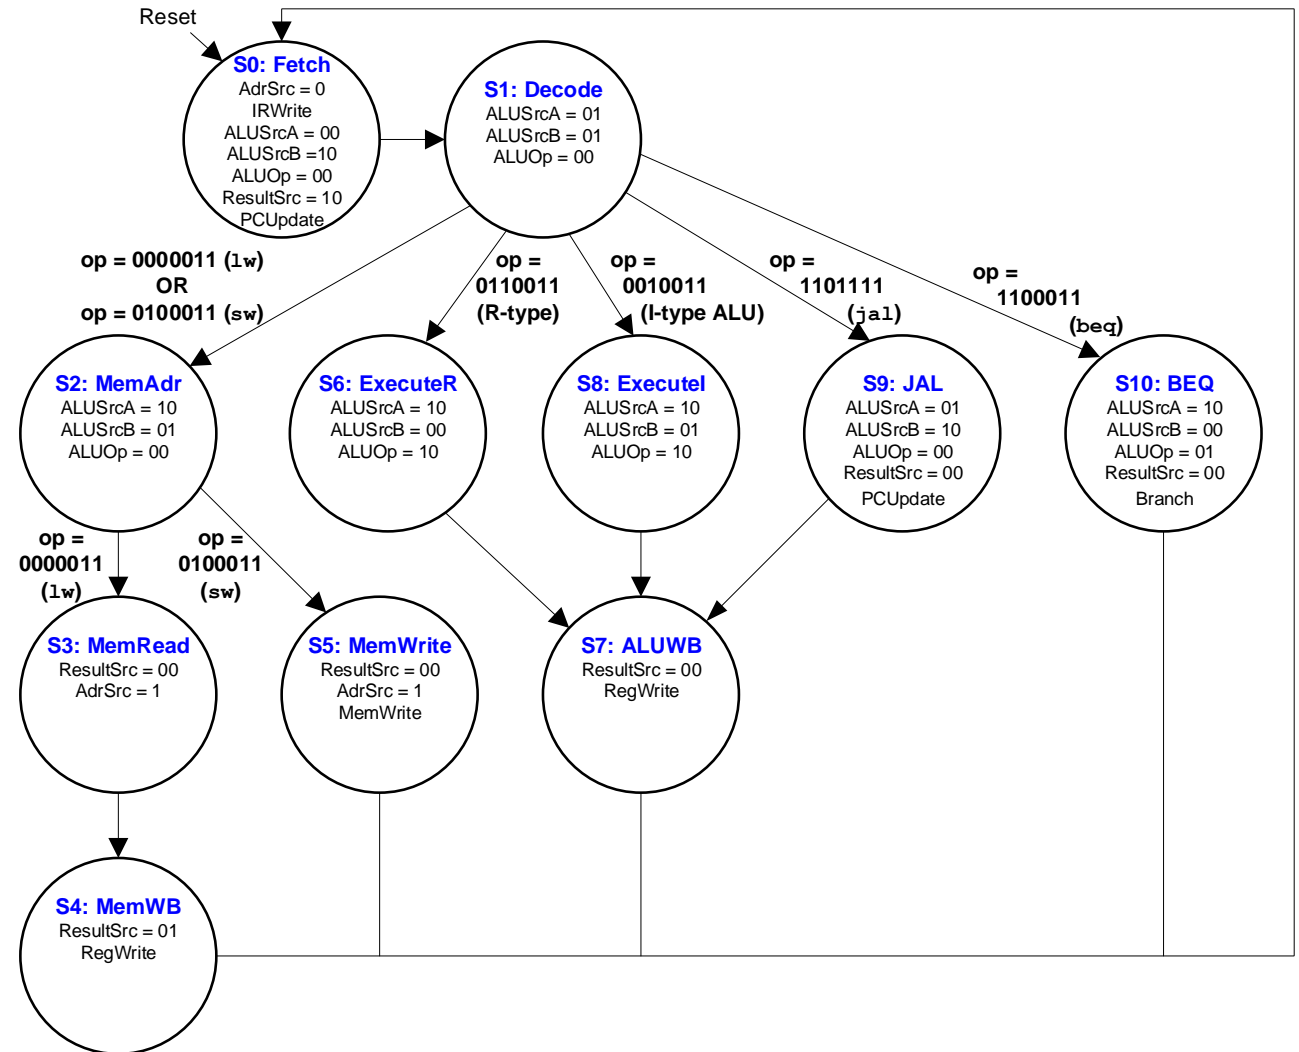
$\text{ALUOut} \leftarrow \text{rs1 op rs2}$

$\text{ALUOut} \leftarrow \text{rs1 op imm}$

$\text{rd} \leftarrow \text{ALUOut}$

$\text{ALUResult} = \text{rs1} - \text{rs2}; \text{if Zero}, \text{PC} \leftarrow \text{ALUOut}$

$\text{PC} \leftarrow \text{ALUOut}; \text{ALUOut} \leftarrow \text{PC}+4$



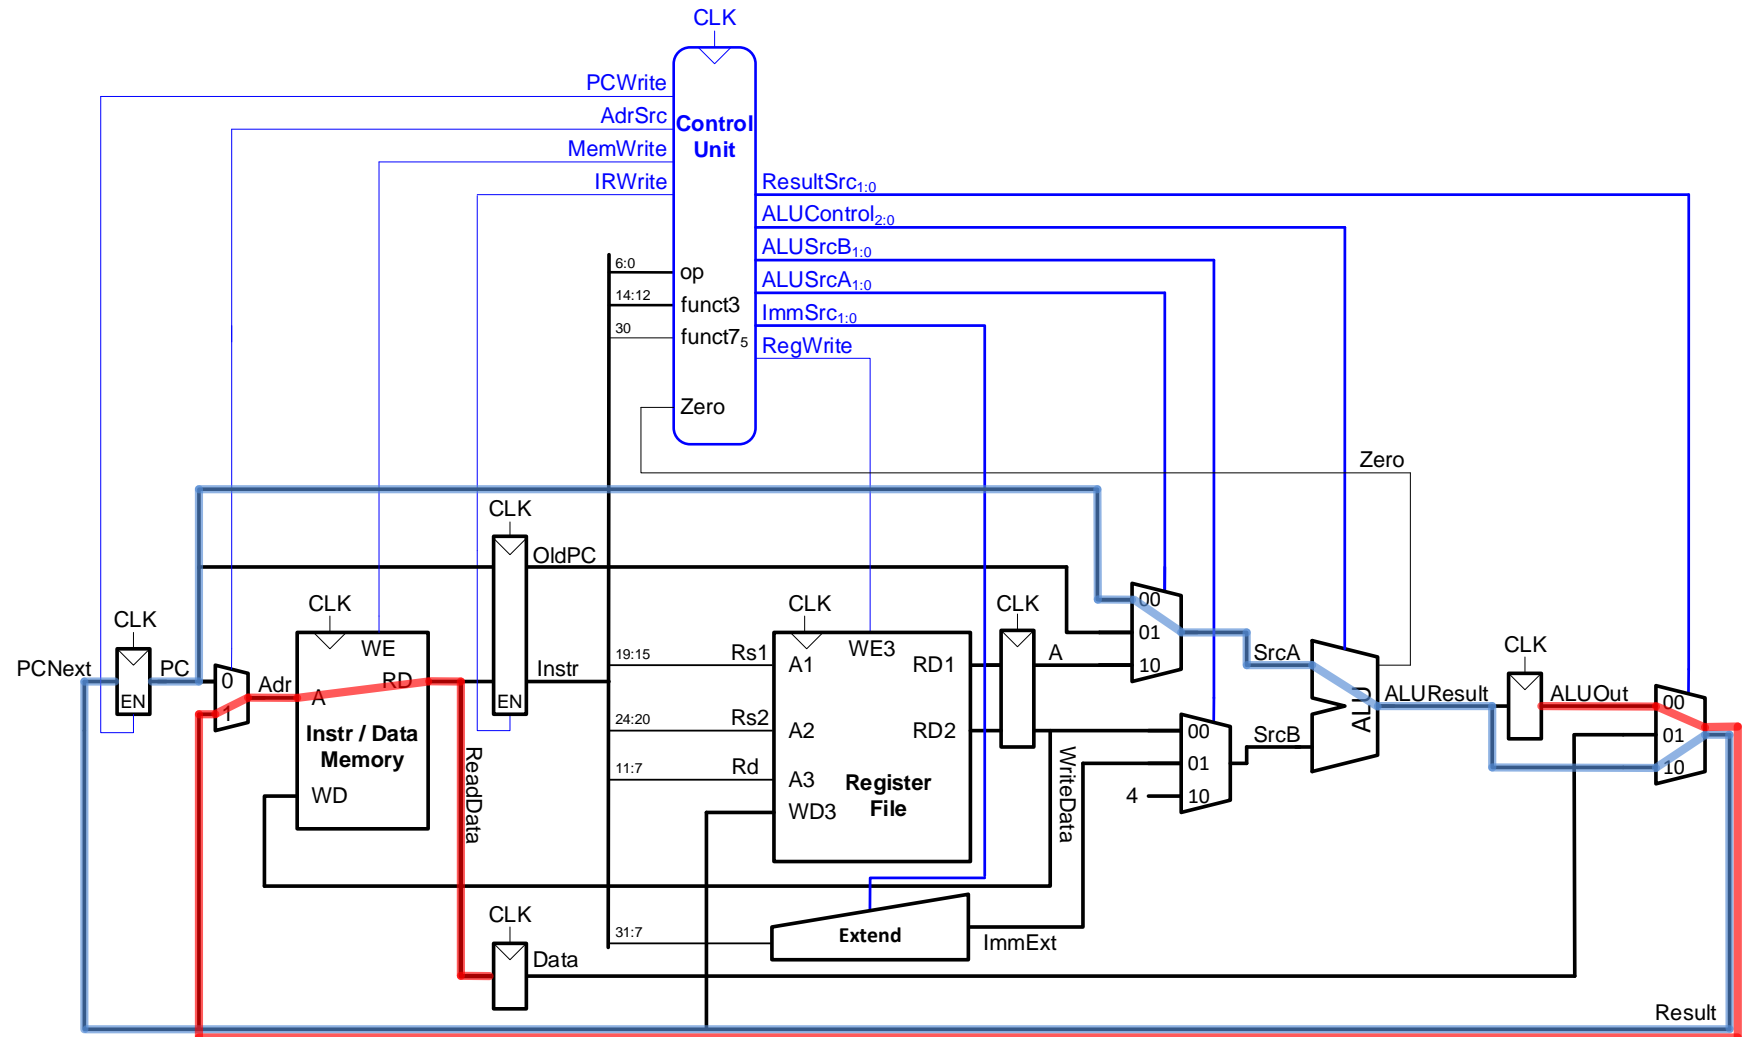
Multicycle Performance

- Instructions take different number of cycles:
 - 3 cycles: beq
 - 4 cycles: R-type, addi, sw, jal
 - 5 cycles: lw
- CPI is weighted average
- SPECINT2000 benchmark:
 - 25% loads
 - 10% stores
 - 13% branches
 - 52% R-type
- **Average CPI** = $(0.13)(3) + (0.52 + 0.10)(4) + (0.25)(5) = 4.12$

Multicycle Critical Path

Potential Critical Paths:

- Calculate PC + 4 or
- Read Memory



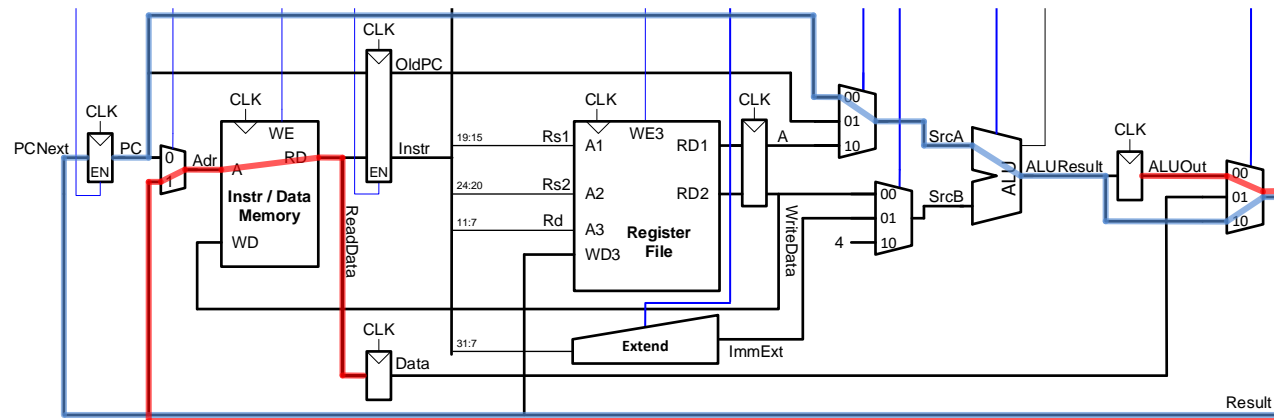
Multicycle Processor Performance

Multicycle critical path:

- **Assumptions:**

- Registry File is faster than memory
- Writing memory is faster than reading memory

$$T_{c_multi} = t_{pcq_PC} + t_{dec} + 2t_{mux} + \max(t_{ALU}, t_{mem}) + t_{setup}$$



Multicycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	t_{pcq_PC}	40
Register setup	t_{setup}	50
Multiplexer	t_{mux}	30
AND-OR gate	t_{AND-OR}	20
ALU	t_{ALU}	120
Decoder (Control Unit)	t_{dec}	25
Extend unit	t_{dec}	35
Memory read	t_{mem}	200
Register file read	t_{RFread}	100
Register file setup	$t_{RFsetup}$	60

$$\begin{aligned}T_{c_multi} &= t_{pcq_PC} + t_{dec} + 2t_{mux} + \max(t_{ALU}, t_{mem}) + t_{setup} \\&= (40 + 25 + 2 * 30 + 200 + 50) ps = \mathbf{375 ps}\end{aligned}$$

- For a program with **100 billion** instructions executing on a **multicycle** RISC-V processor
 - **CPI** = 4.12 cycles/instruction
 - **Clock cycle time:** $T_{c_multi} = 375 \text{ ps}$
- **Execution Time** = (# instructions) \times CPI $\times T_c$
 - = $(100 \times 10^9)(4.12)(375 \times 10^{-12})$
 - = **155 seconds**
- This is **slower** than the single-cycle processor (75 sec.)

Parallelism

DDCA Ch3 - Part 16: Parallelism <https://www.youtube.com/watch?v=xX2Crru3xCg>

Two types of parallelism:

- **Spatial parallelism**
 - duplicate hardware performs multiple tasks at once
- **Temporal parallelism**
 - task is broken into multiple stages
 - also called pipelining
 - for example, an assembly line

- **Token**: Group of inputs processed to produce group of outputs
- **Latency**: Time for one token to pass from start to end
- **Throughput**: Number of tokens produced per unit time

Parallelism increases throughput

Parallelism Example 1/3

- Ben Bitdiddle bakes cookies to celebrate traffic light controller installation
 - **5 minutes** to roll cookies
 - **15 minutes** to bake
- What is the latency and throughput without parallelism?
 - Latency (when is the first cookie finished?)
 - Throughput (how many cookies can Ben finish in an hour?)

Latency = 5 + 15 = 20 minutes = $\frac{1}{3}$ hour

Throughput = 1 tray / $\frac{1}{3}$ hour = 3 trays/hour

What is the latency and throughput if Ben uses parallelism?

- **Spatial parallelism:**

Ben asks Allysa P. Hacker to help, using her own oven

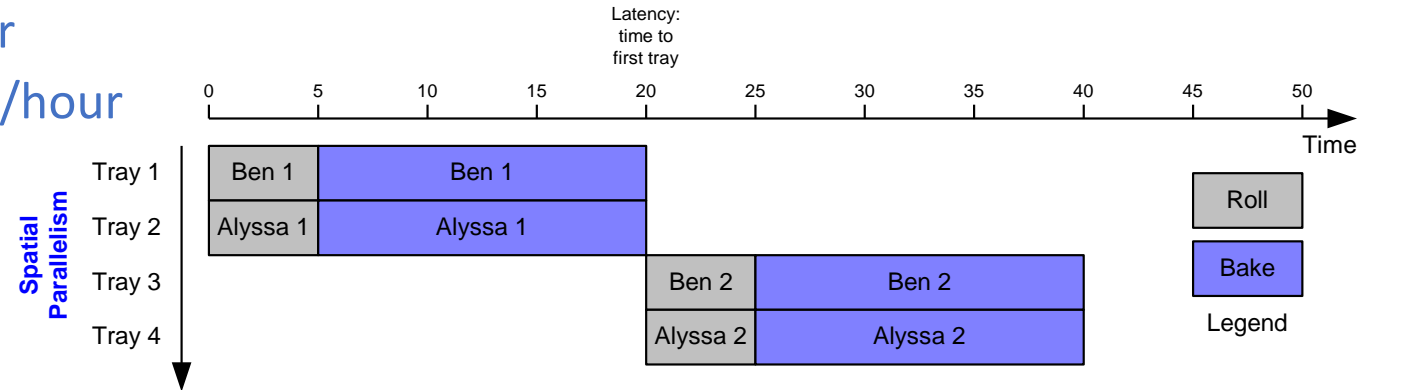
- **Temporal parallelism:**

- two stages: rolling and baking
- He uses two trays
- While first batch is baking, he rolls the second batch, etc.

Spatial vs. Temporal Parallelism

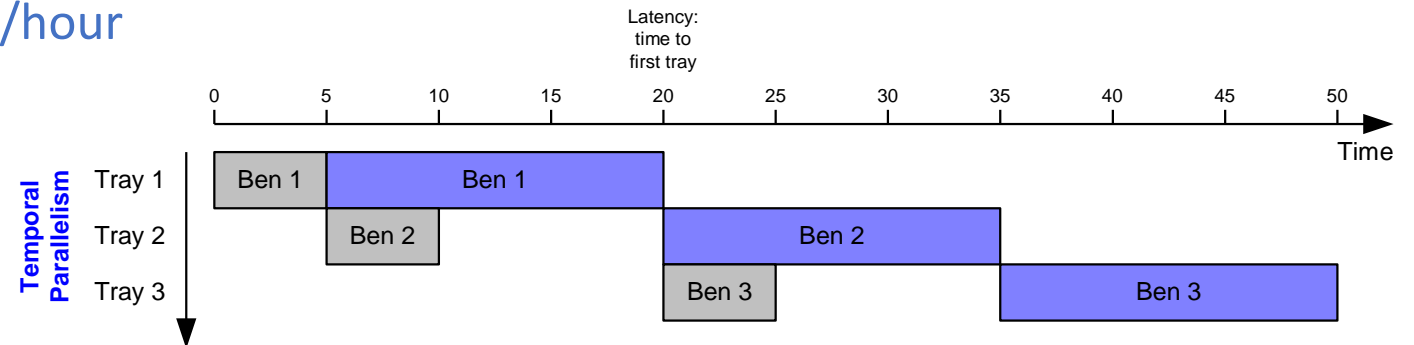
- Spatial Parallelism

- Latency = 5 + 15 = 20 minutes = $\frac{1}{3}$ hour
- Throughput = 2 trays/ $\frac{1}{3}$ hour = 6 trays/hour



- Temporal Parallelism

- Latency = 5 + 15 = 20 minutes = $\frac{1}{3}$ hour
- Throughput = 1 trays/ $\frac{1}{4}$ hour = 4 trays/hour



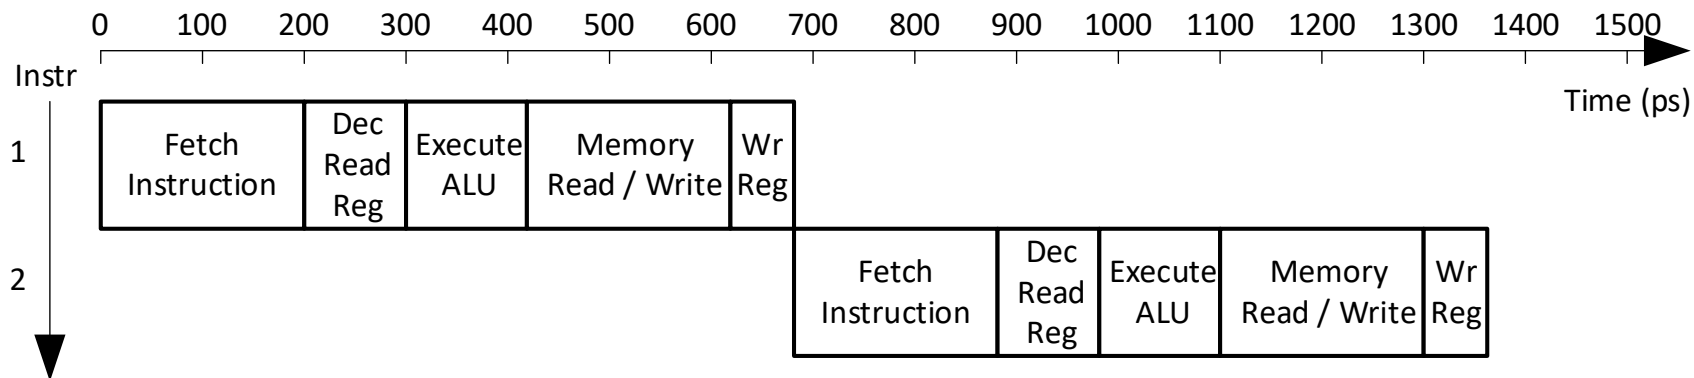
Pipelined RISC-V Processor

DDCA Ch7 - Part 13: Pipelined Processor <https://www.youtube.com/watch?v=UZdURUwQMmk>

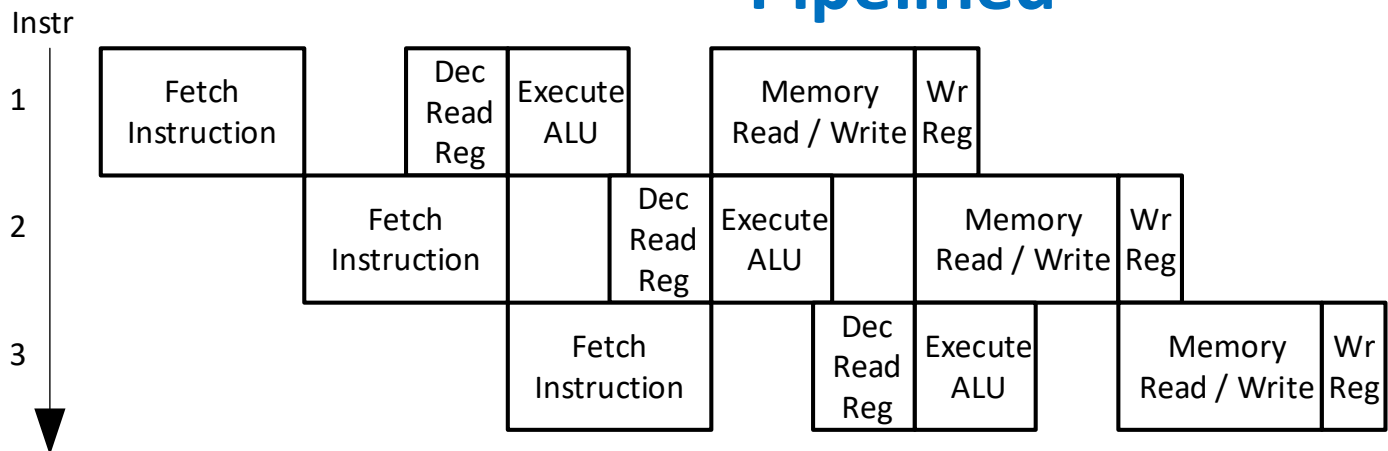
- **Temporal parallelism**
- Divide single-cycle processor into **5 stages**:
 - Fetch
 - Decode
 - Execute
 - Memory
 - Writeback
- Add **pipeline registers** between stages

Single-Cycle vs. Pipelined Processor

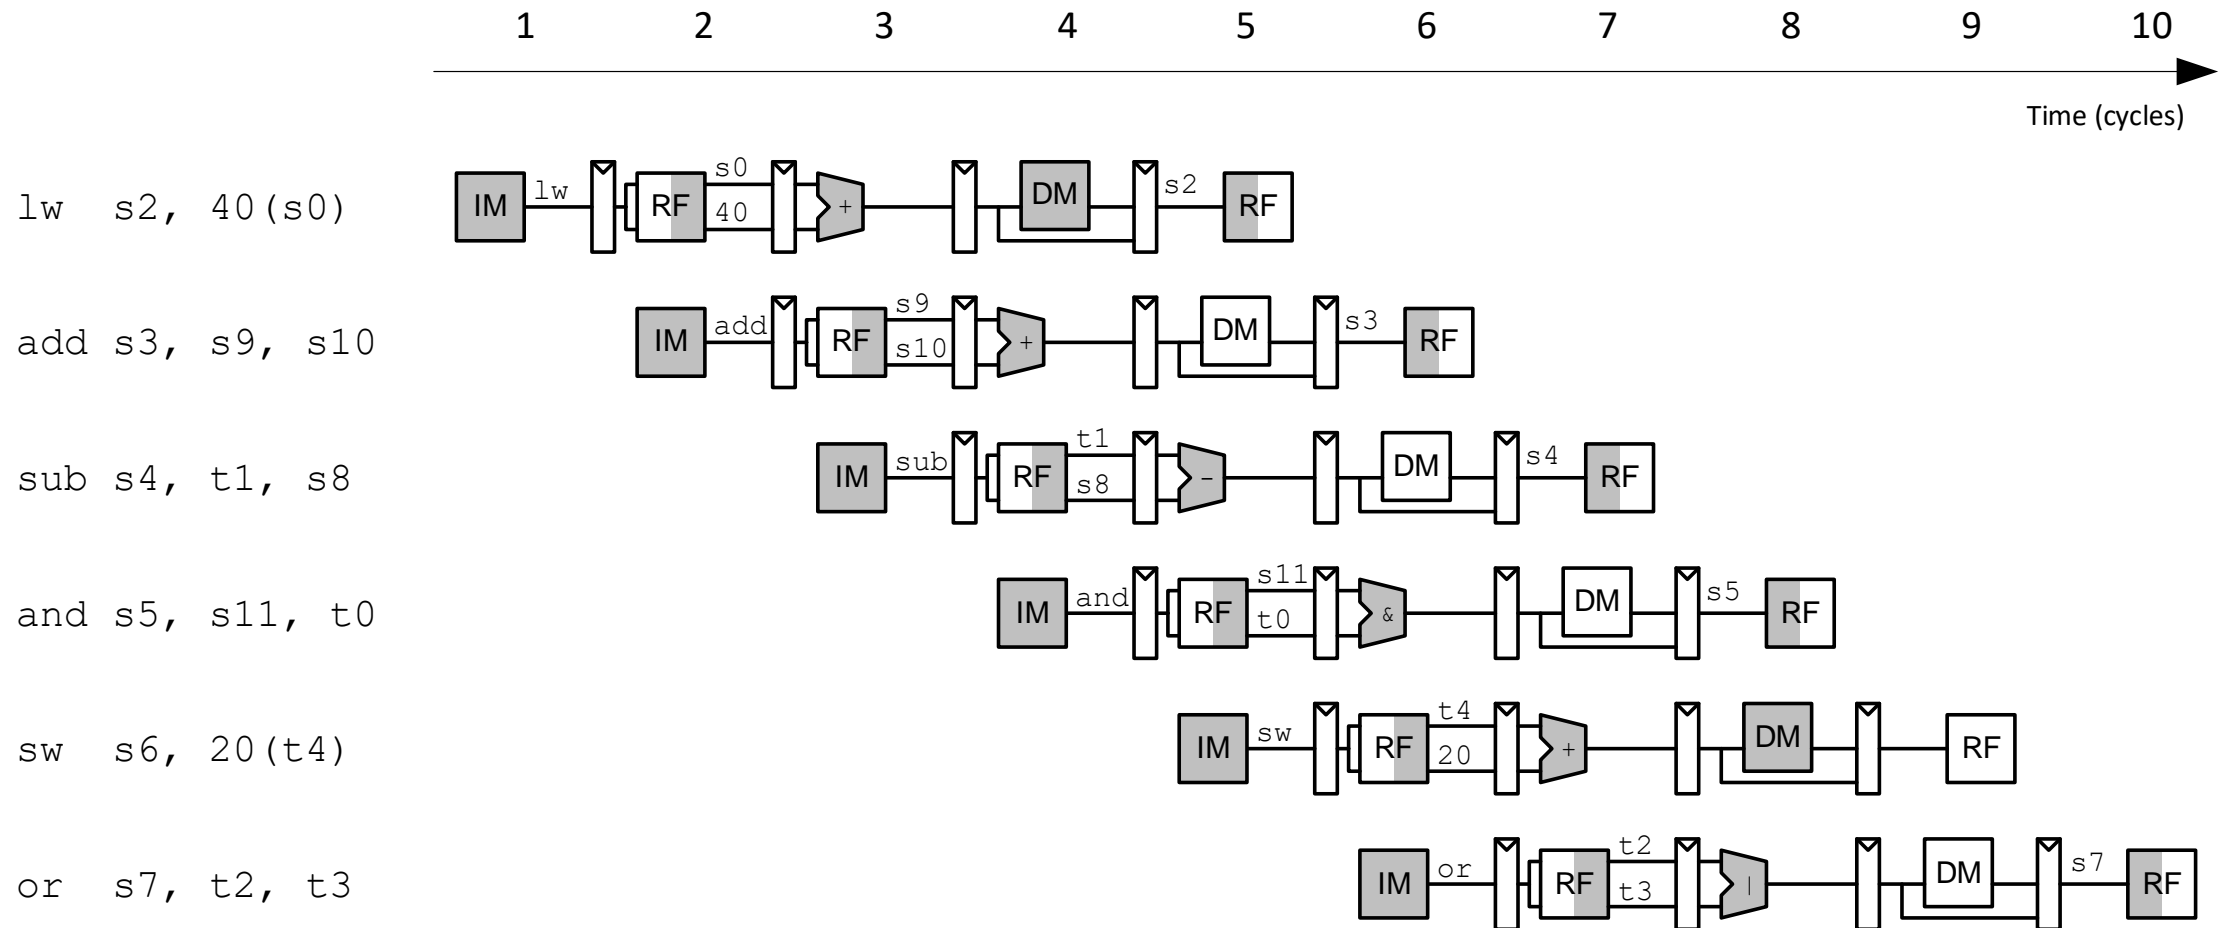
Single-Cycle



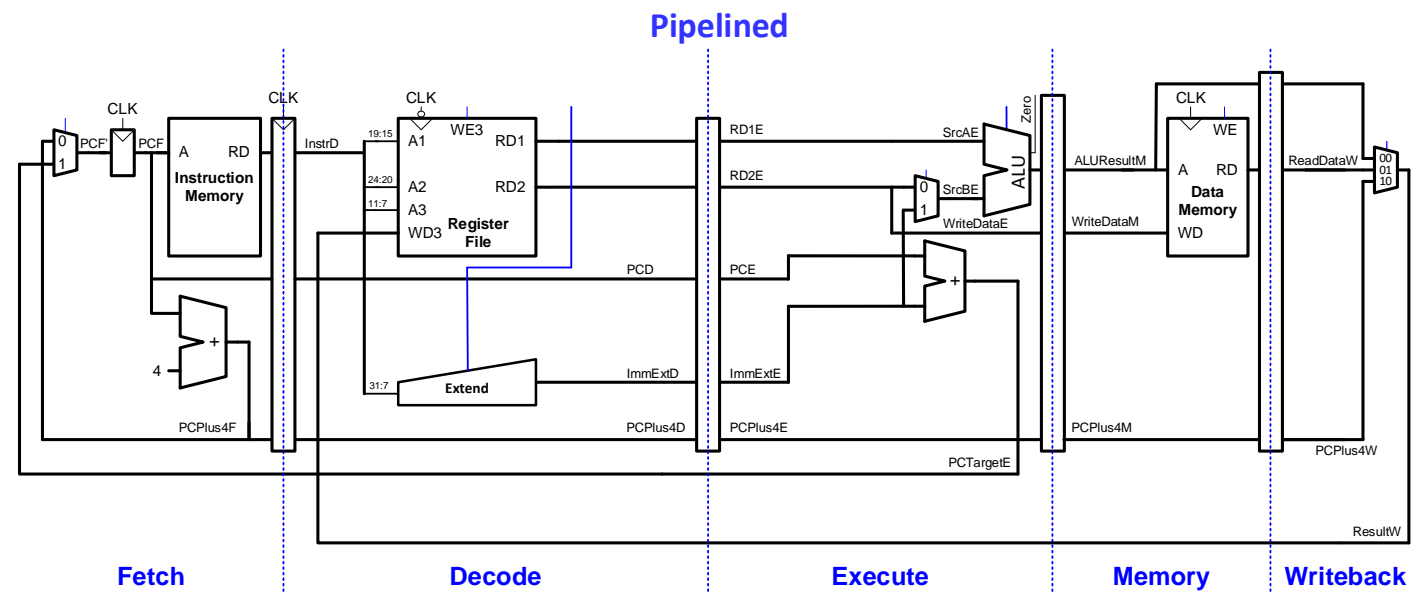
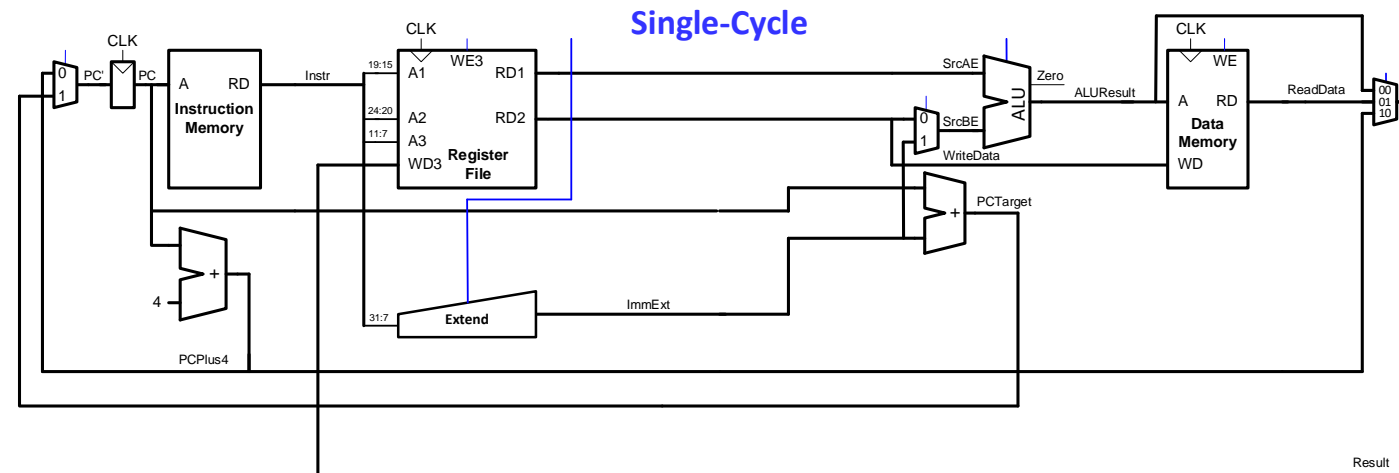
Pipelined



Pipelined Processor Abstraction

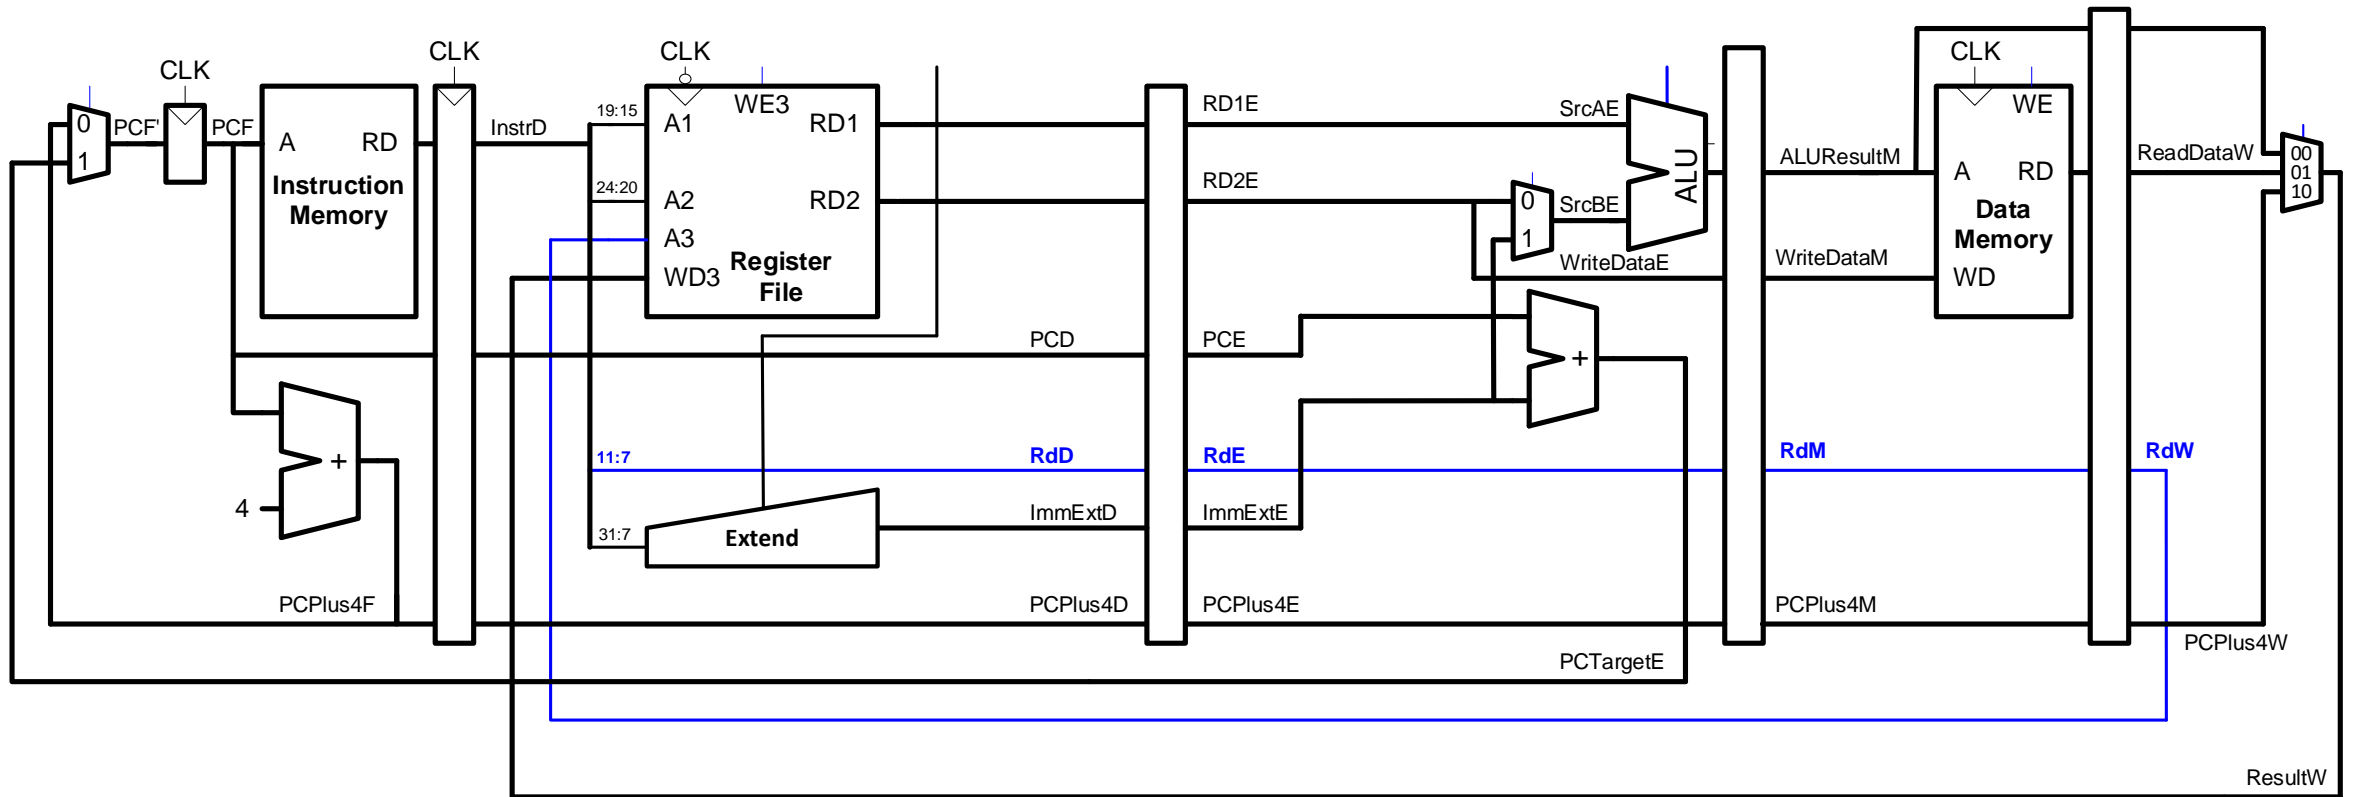


Single-Cycle & Pipelined Datapaths



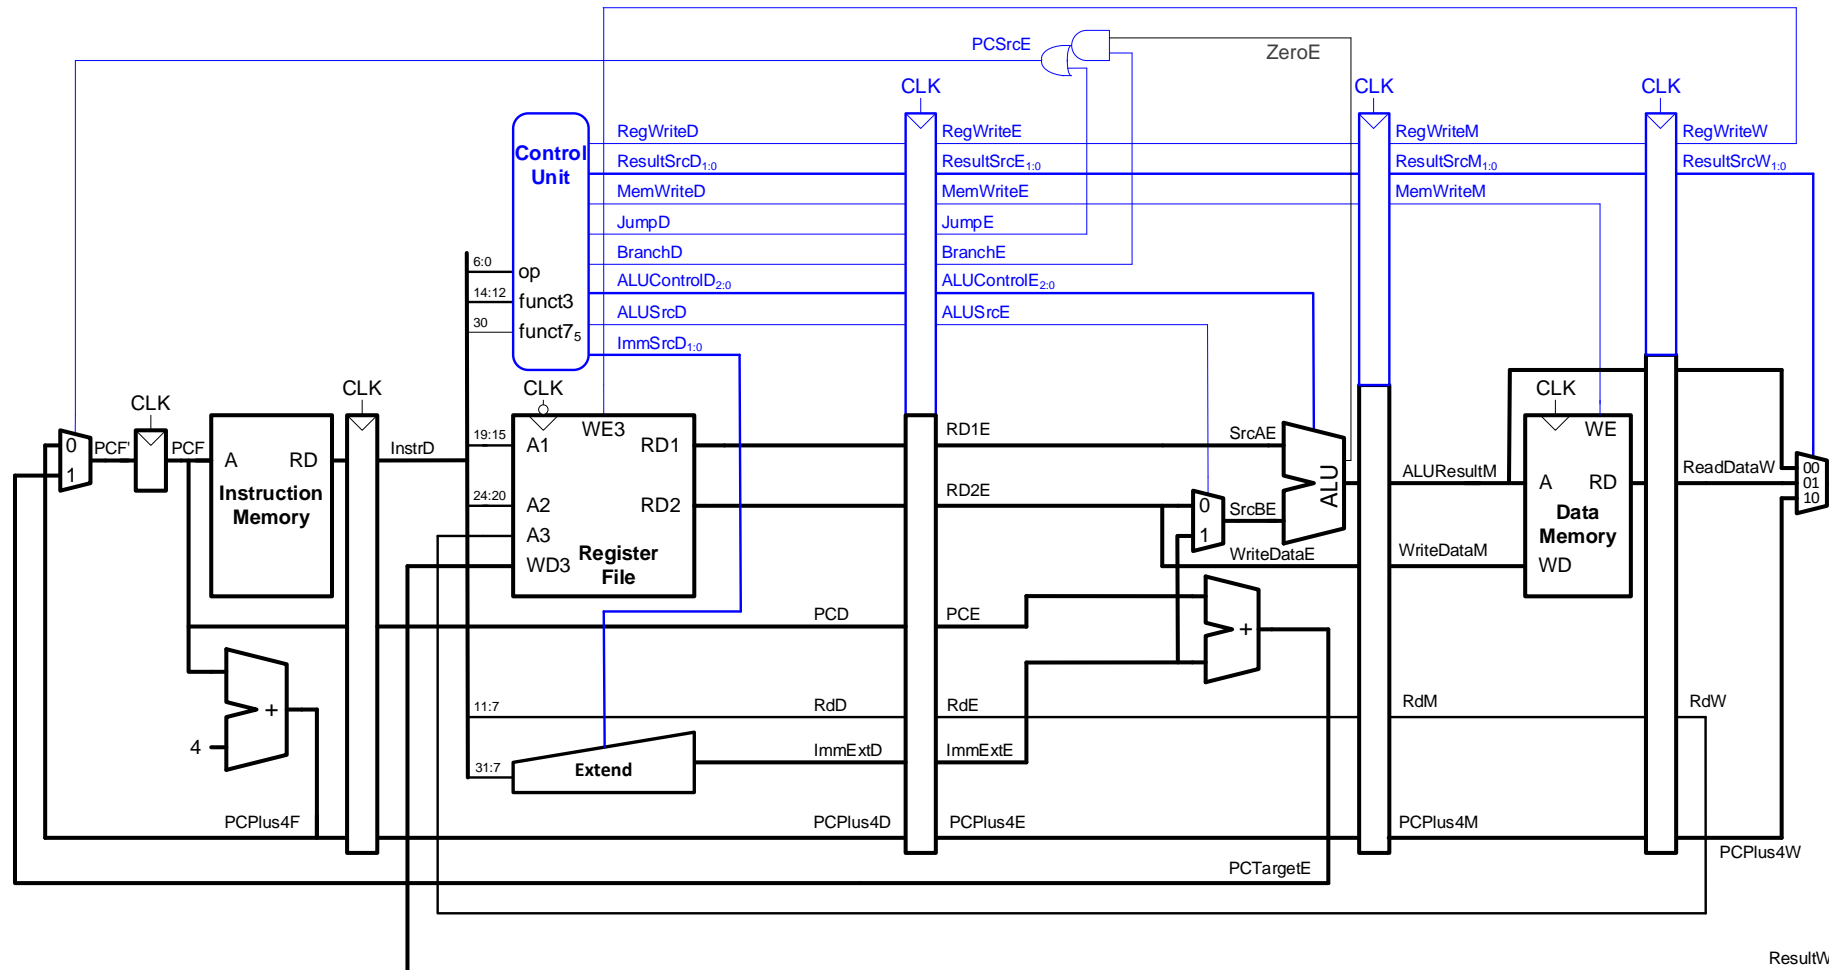
Signals in Pipelined Processor are appended with first letter of stage (i.e., PC**F**, PC**D**, PC**E**, PC**M**, PC**W**).

Corrected Pipelined Datapath



- **Rd** must arrive at same time as **Result**
- Register file written on **falling edge** of **CLK**

Pipelined Processor with Control



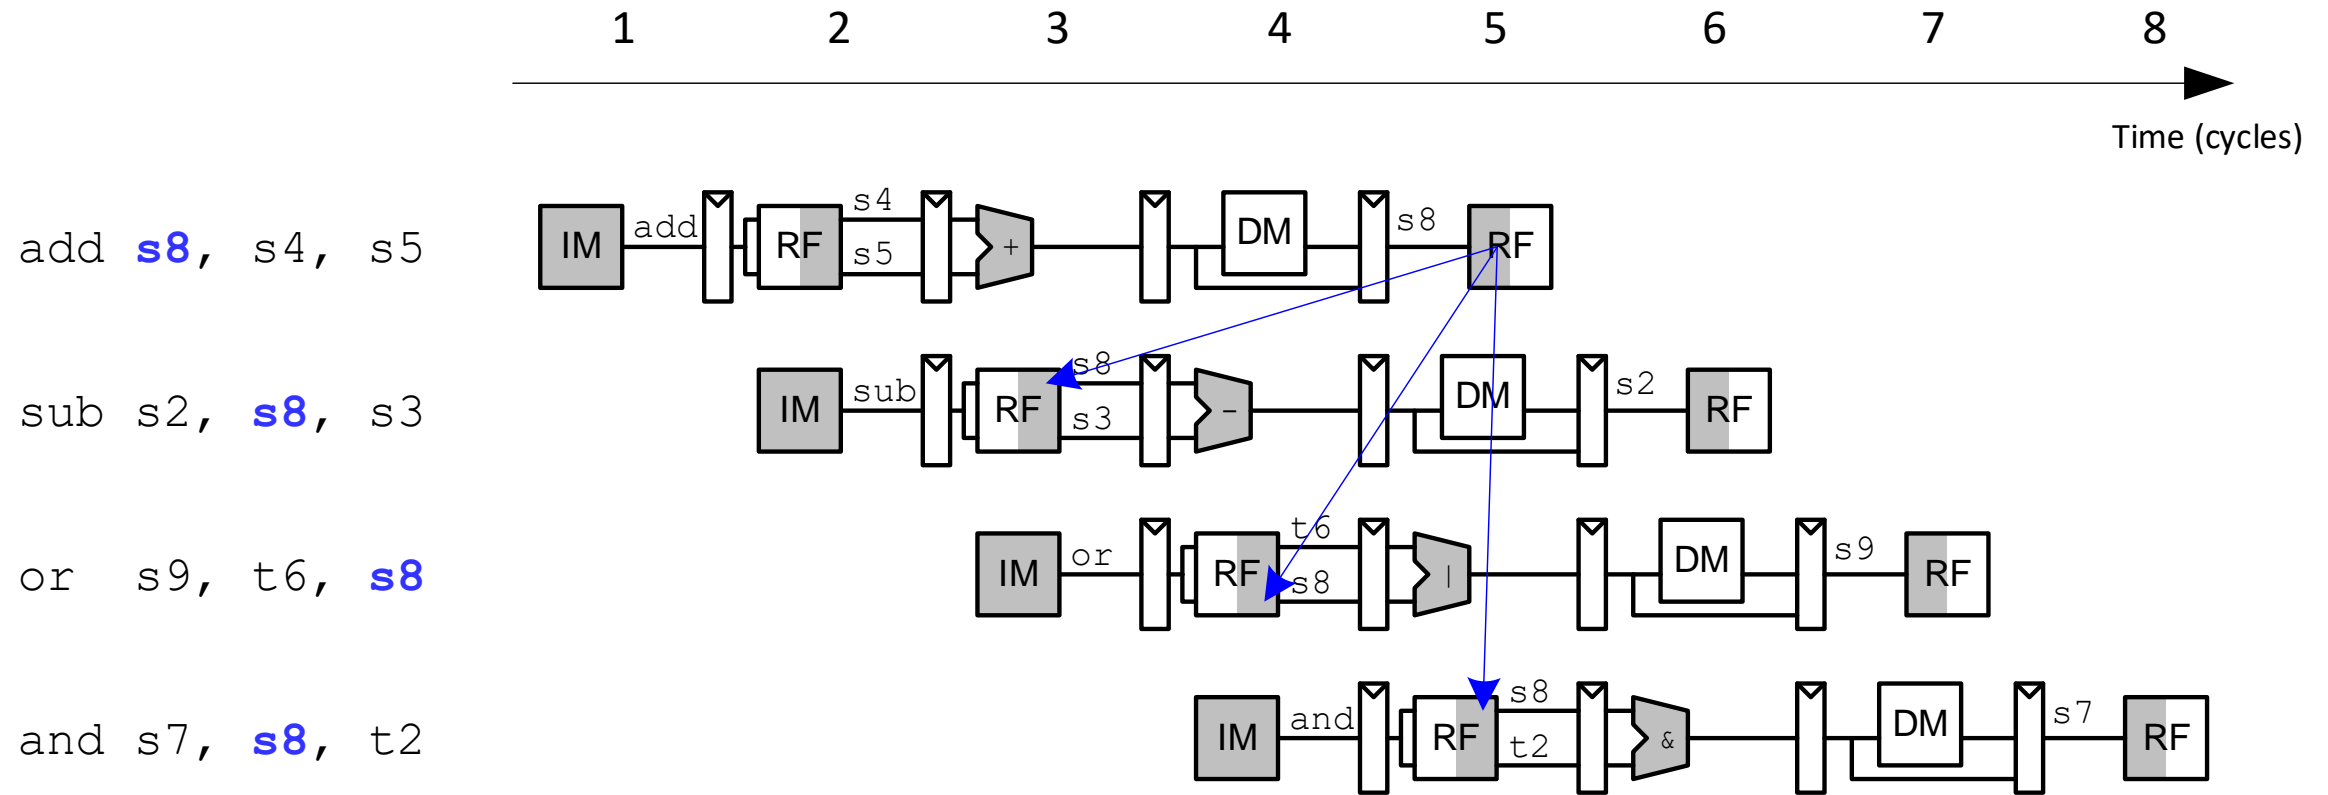
- **Same control unit** as single-cycle processor
- **Control signals travel with** the instruction (drop off when used)

Pipelined Processor Hazards

DDCA Ch7 - Part 14: Pipelined Processor Data Hazards <https://www.youtube.com/watch?v=zuegcg6ZSFQ>

- When an instruction depends on result from instruction that hasn't completed
- Types:
 - **Data hazard:**
register value not yet written back to register file
 - **Control hazard:**
next instruction not decided yet (caused by branch)

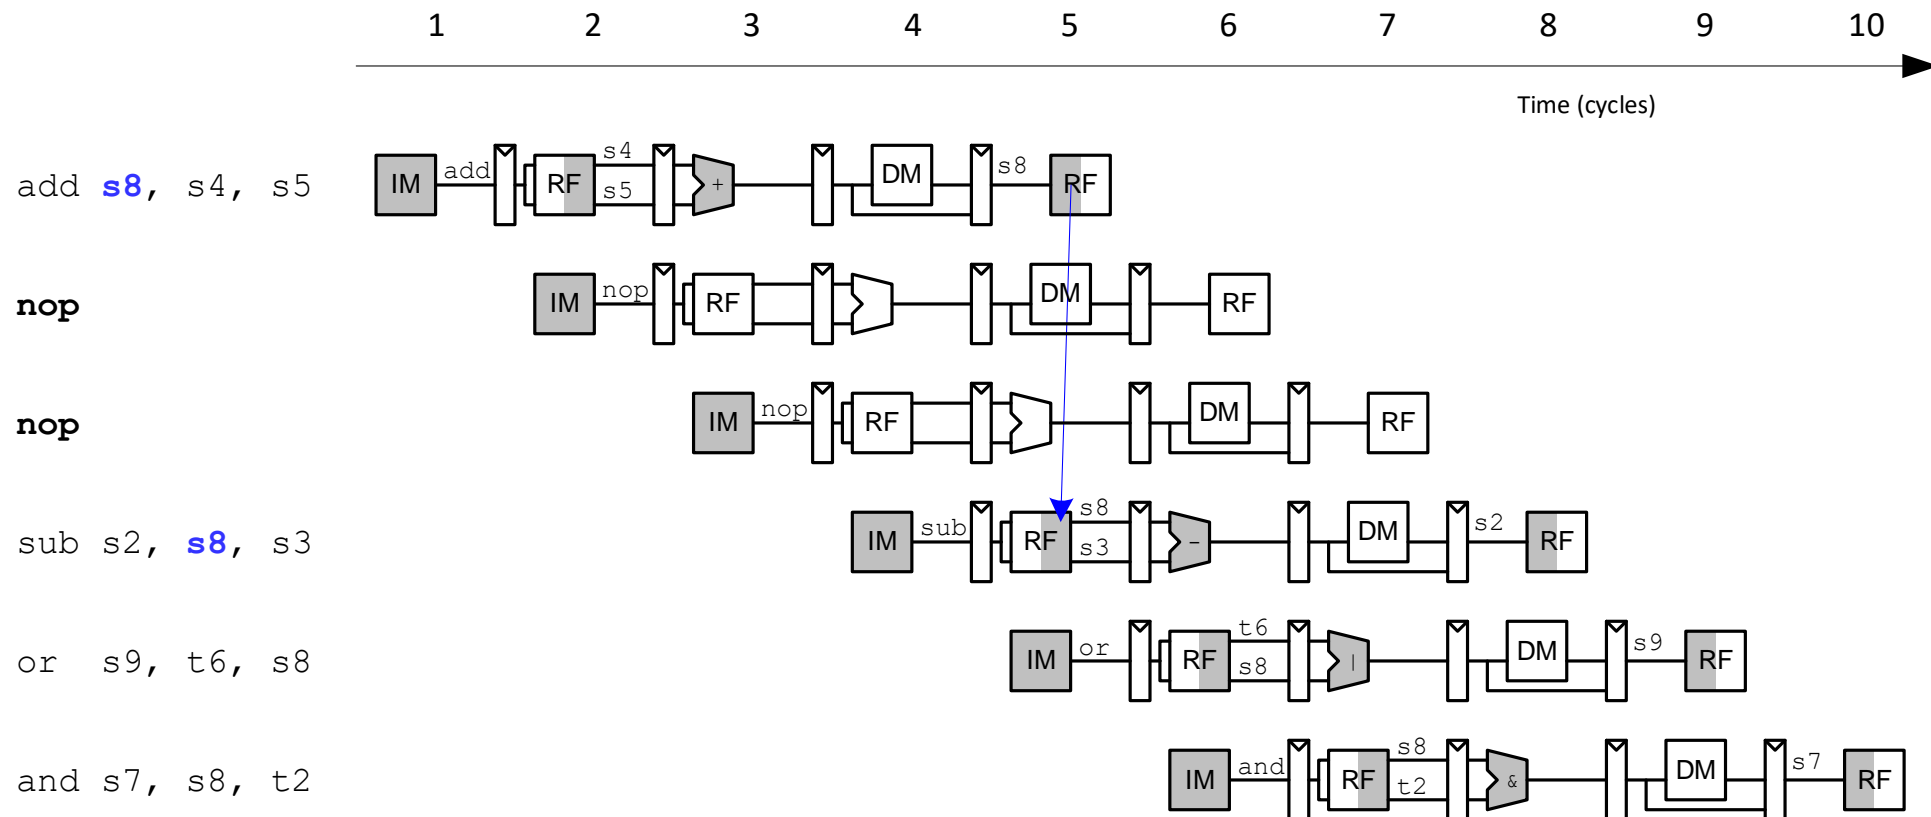
Data Hazard



- Insert **nops** in code at compile time
- **Rearrange** code at compile time
- **Forward** data at run time
- **Stall** the processor at run time

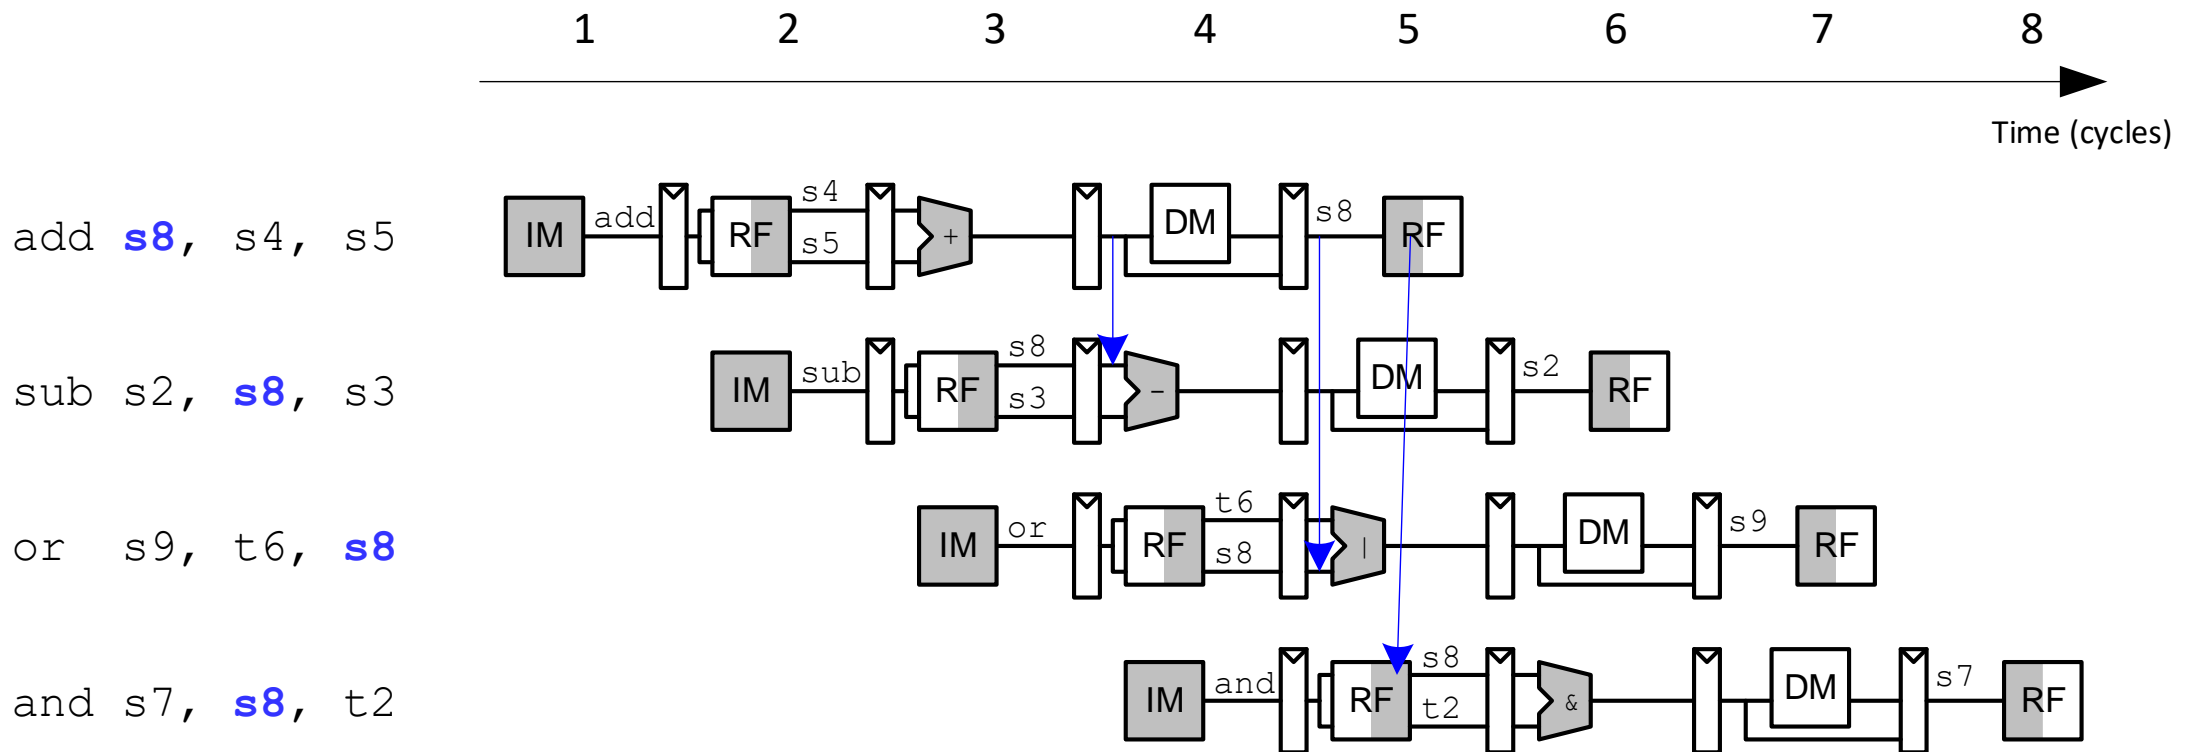
Handling Data Hazards

- **Insert** enough **nops** for result to be ready
- Or move independent useful instructions forward



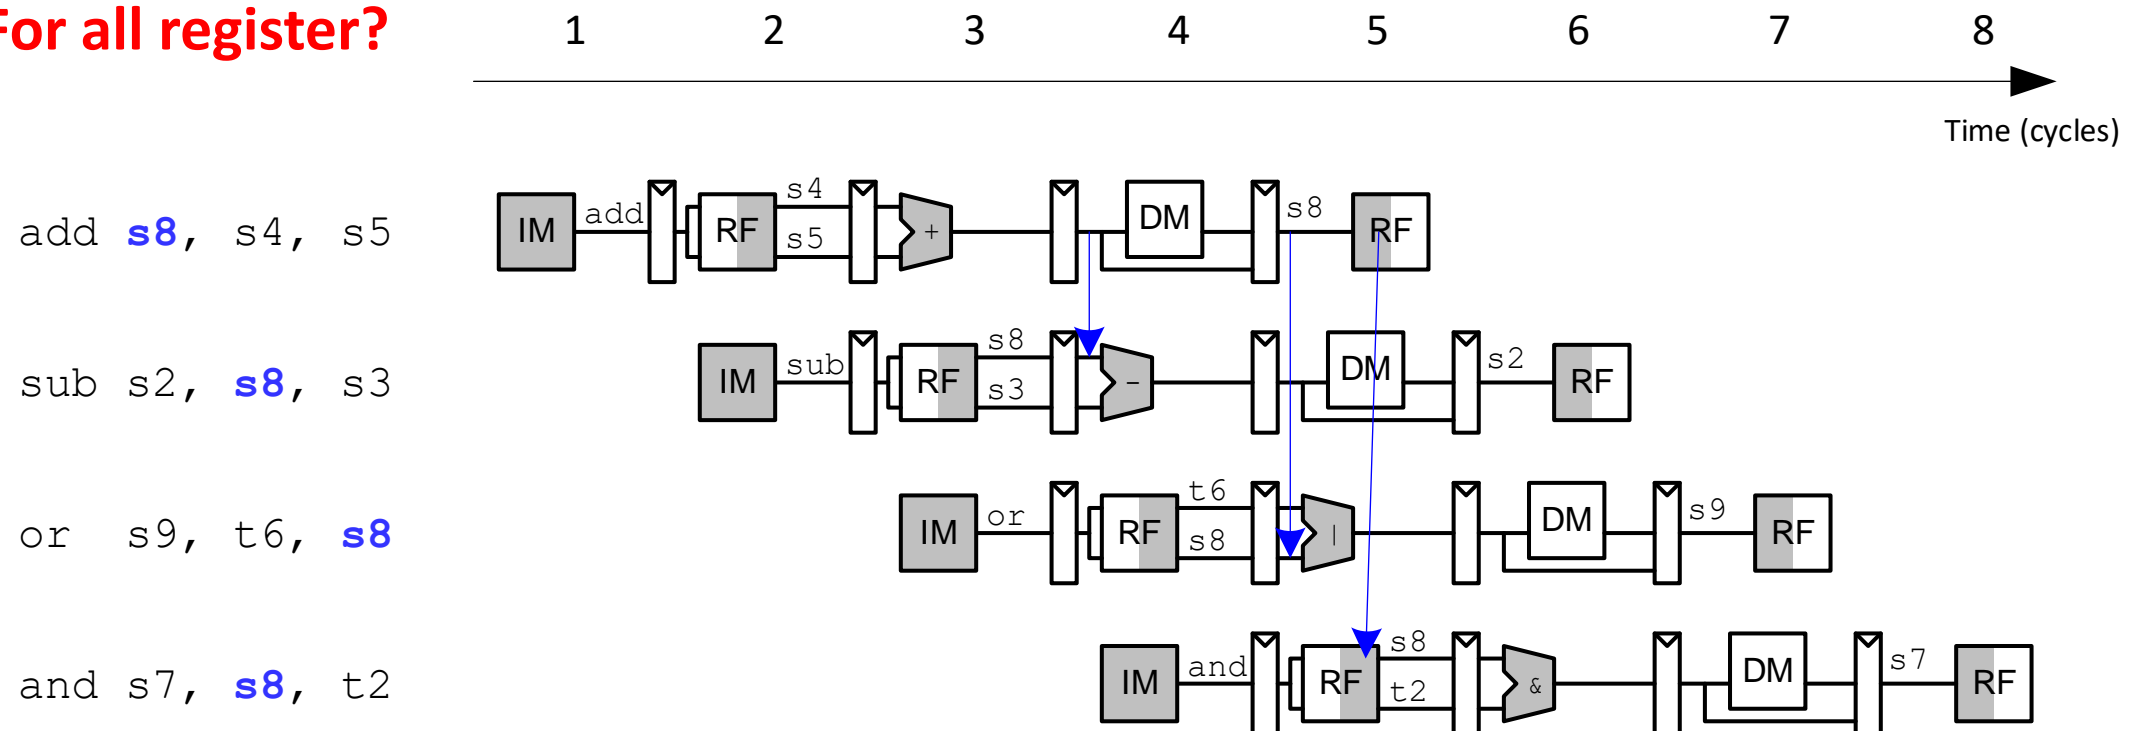
Data Forwarding

- Data is **available on internal busses** before it is written back to the register file (RF).
- **Forward data** from internal busses **to Execute stage**.

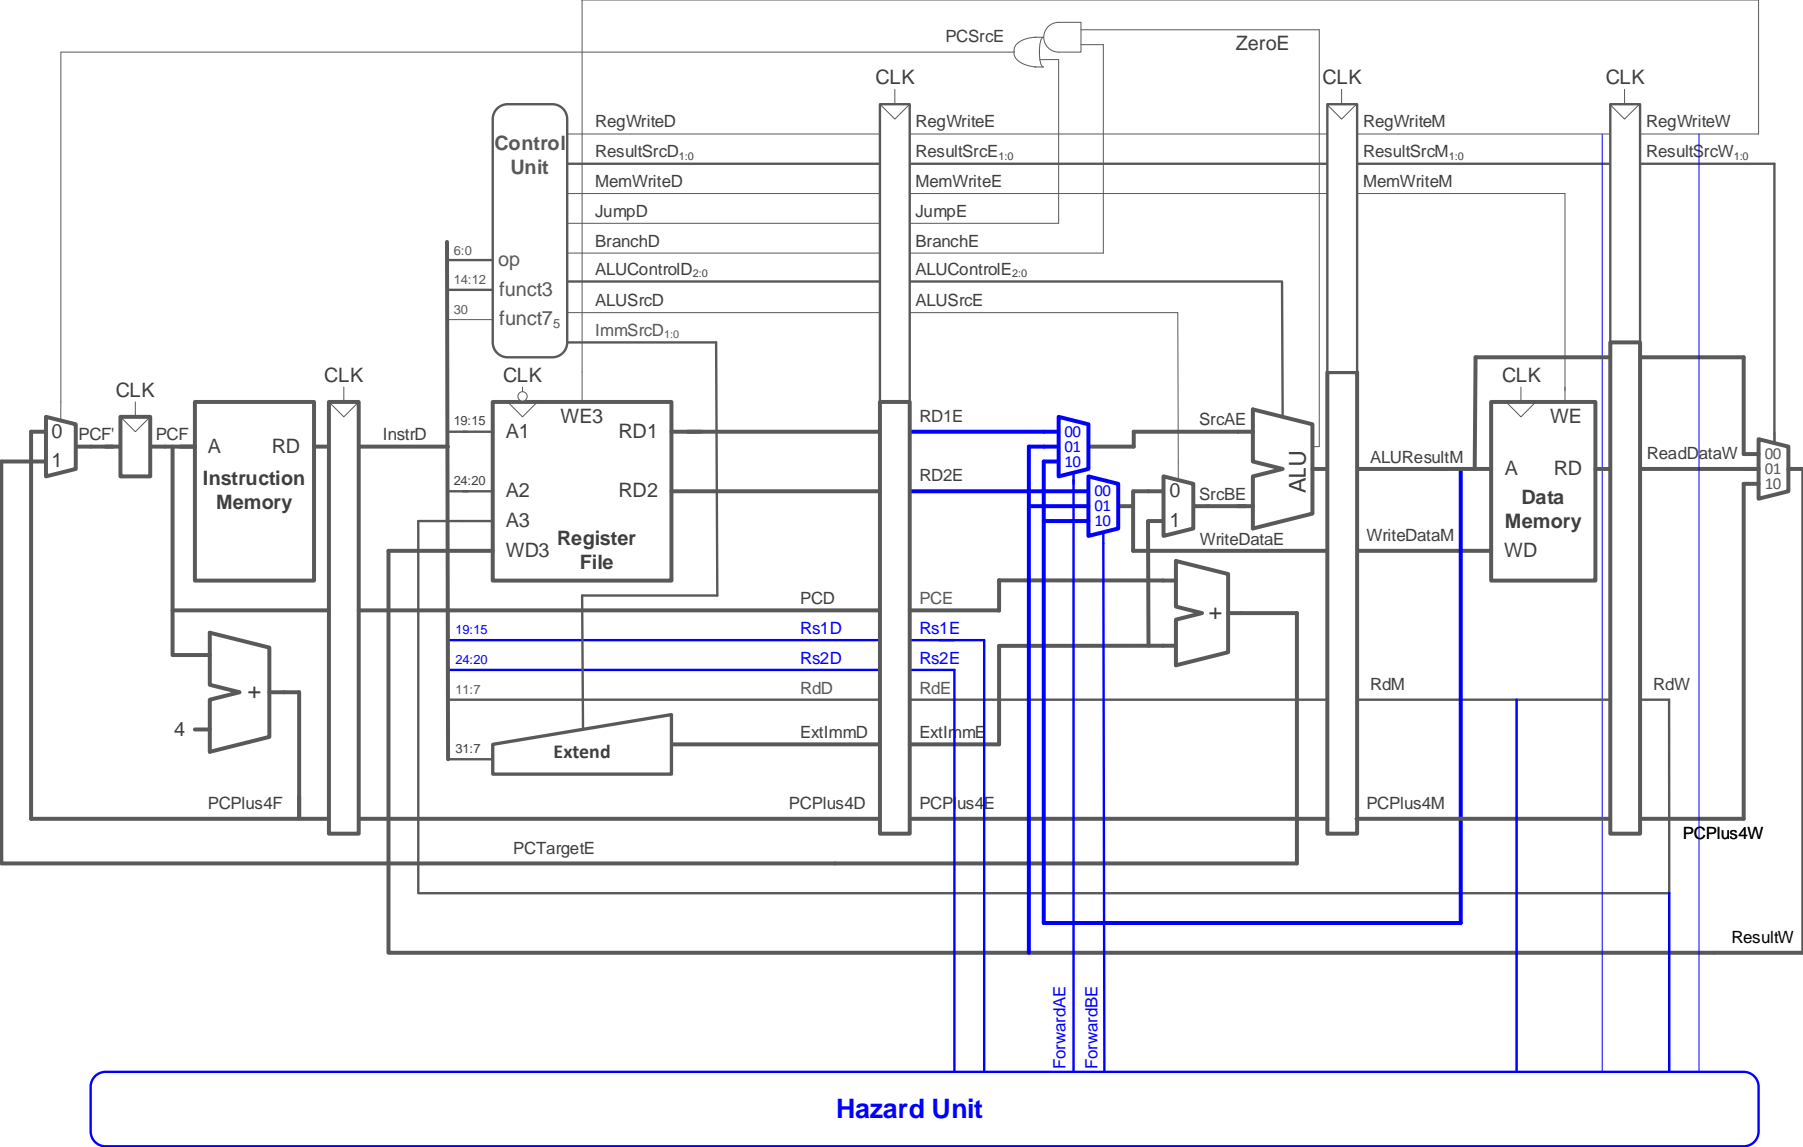


Data Forwarding

- Check if source register in **Execute stage** **matches** destination register of instruction in **Memory or Writeback stage**.
- If so, forward result.
- **For all register?**

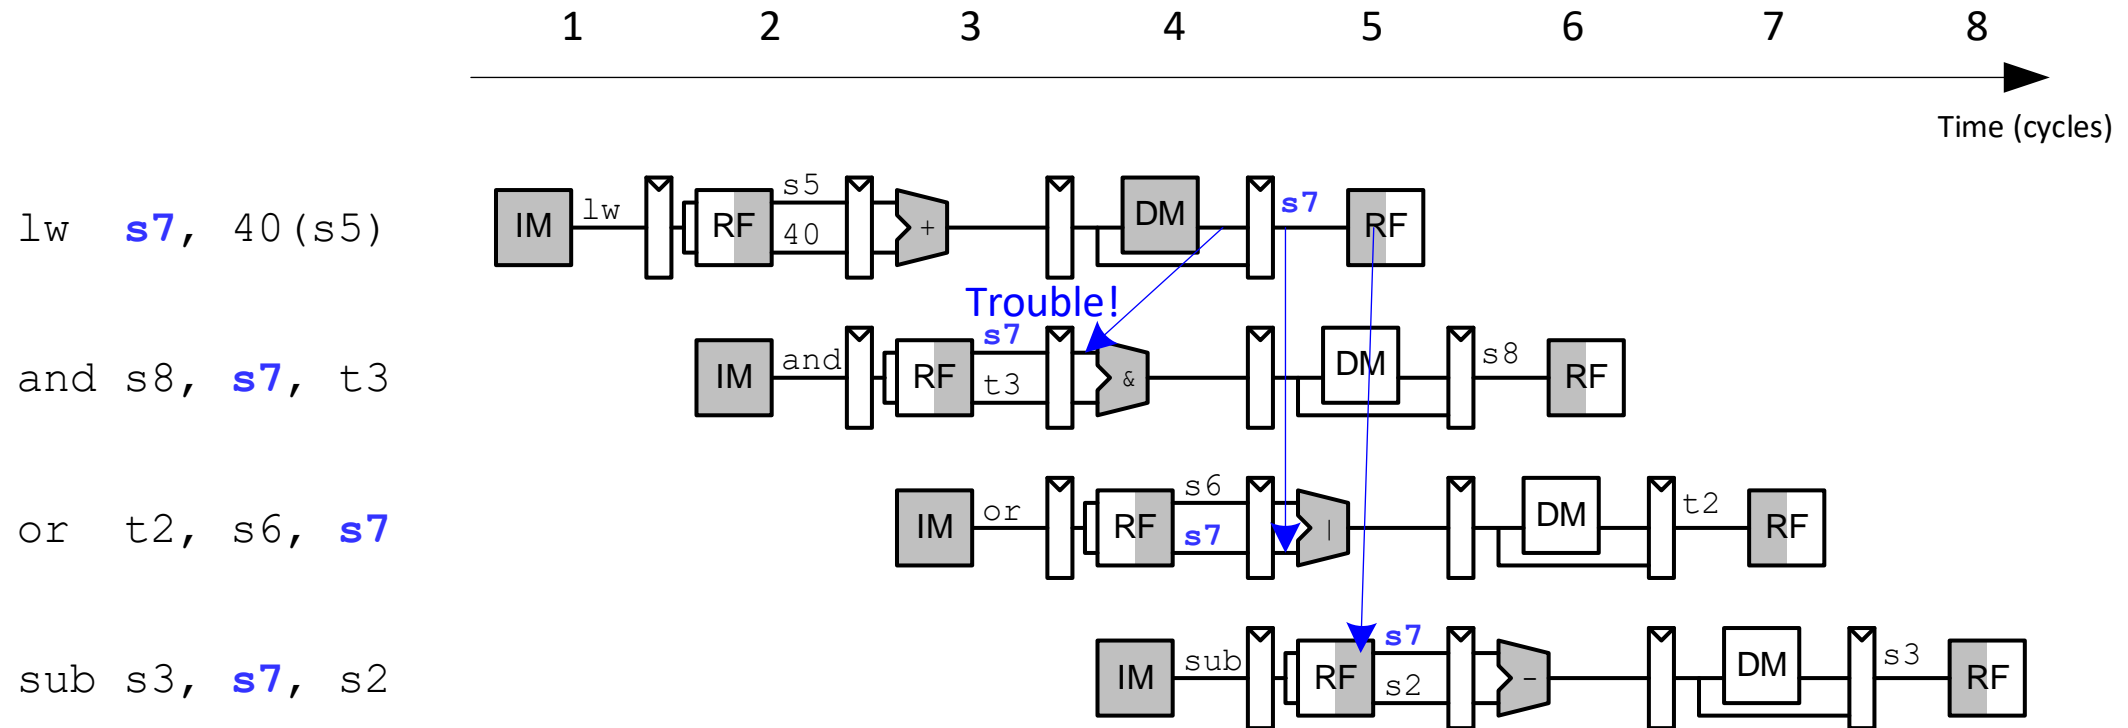


Data Forwarding: Hazard Unit

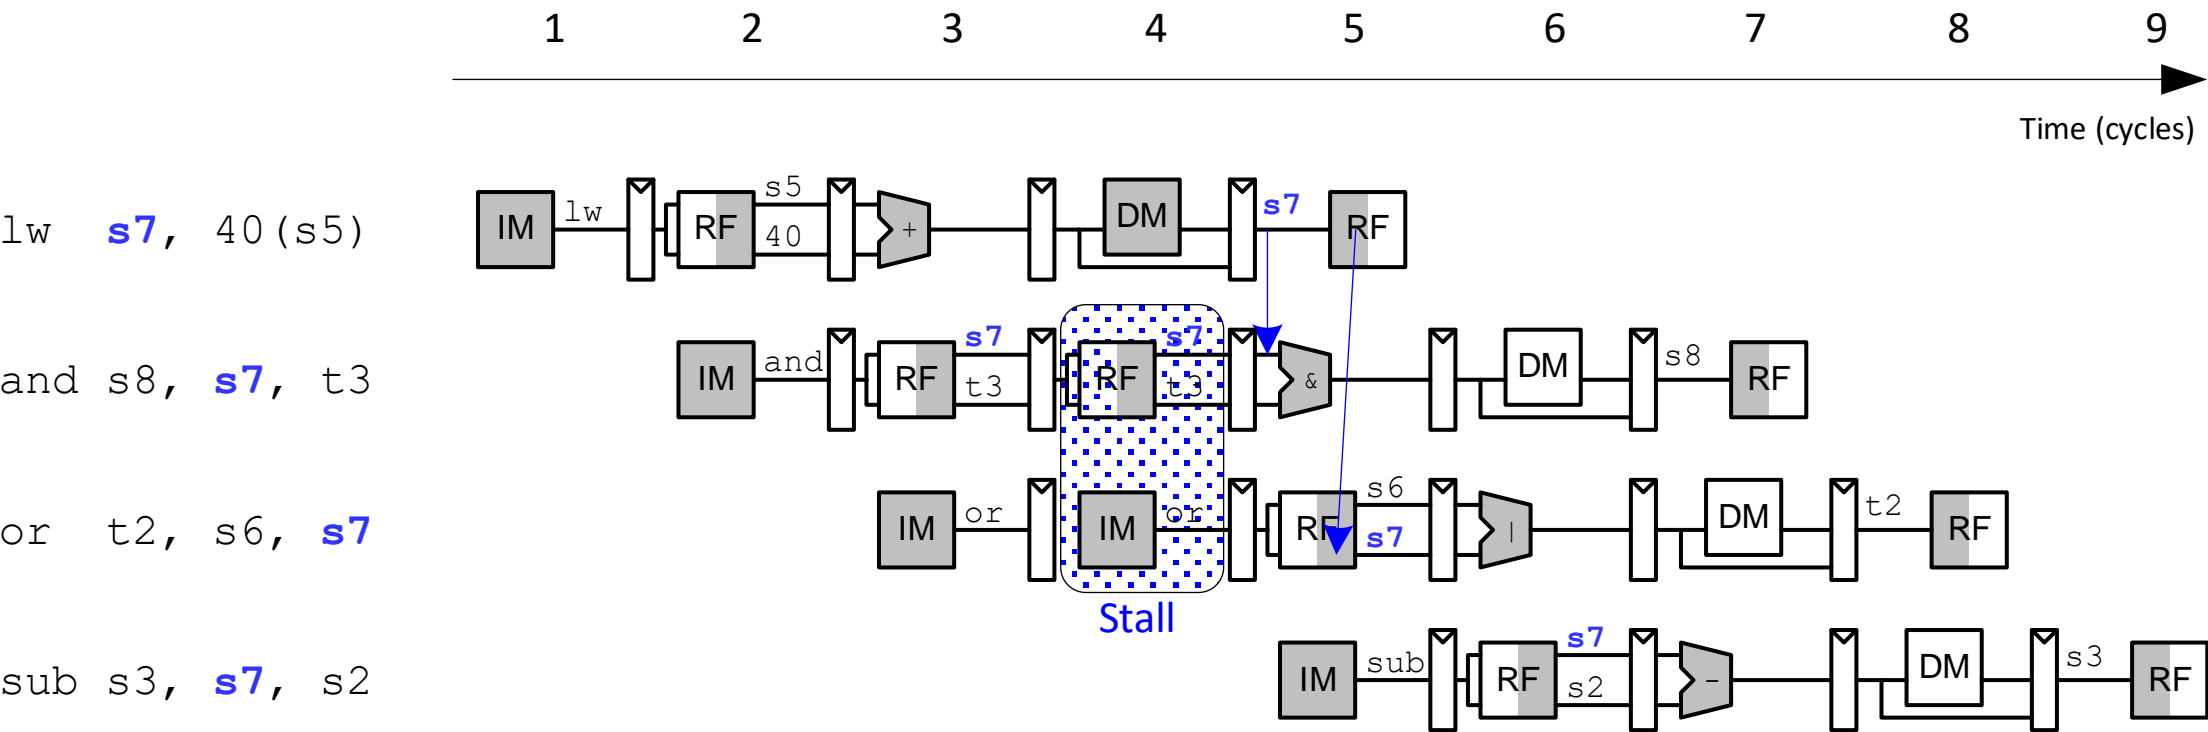


- **Case 1: Execute** stage $Rs1$ or $Rs2$ matches **Memory** stage Rd ?
Forward from Memory stage
- **Case 2: Execute** stage $Rs1$ or $Rs2$ matches **Writeback** stage Rd ?
Forward from Writeback stage
- **Case 3:** Otherwise use value read from register file (as usual)
- **Equations for ForwardAE - $Rs1$:**
if $((Rs1E == RdM) \text{ AND } RegWriteM) \text{ AND } (Rs1E != 0) // \text{ Case 1}$
 $ForwardAE = 10$
else if $((Rs1E == RdW) \text{ AND } RegWriteW) \text{ AND } (Rs1E != 0) // \text{ Case 2}$
 $ForwardAE = 01$
else $ForwardAE = 00 // \text{ Case 3}$
- **ForwardBE** equations are similar (replace $Rs1E$ with $Rs2E$)

Data Hazard due to l_w Dependency



Data Hazard due to l_w Dependency



Stalling Logic

- Is either **source register in the Decode stage** the same as the **destination register in the Execute stage**?

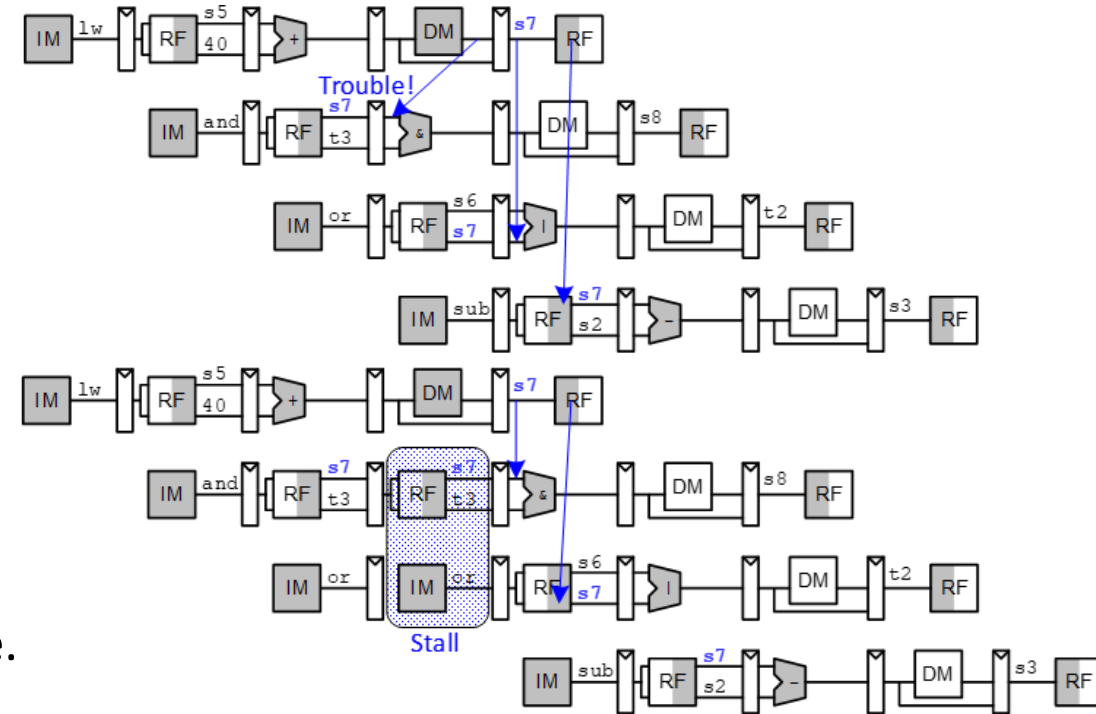
AND

- Is the instruction in the **Execute stage a lw**?

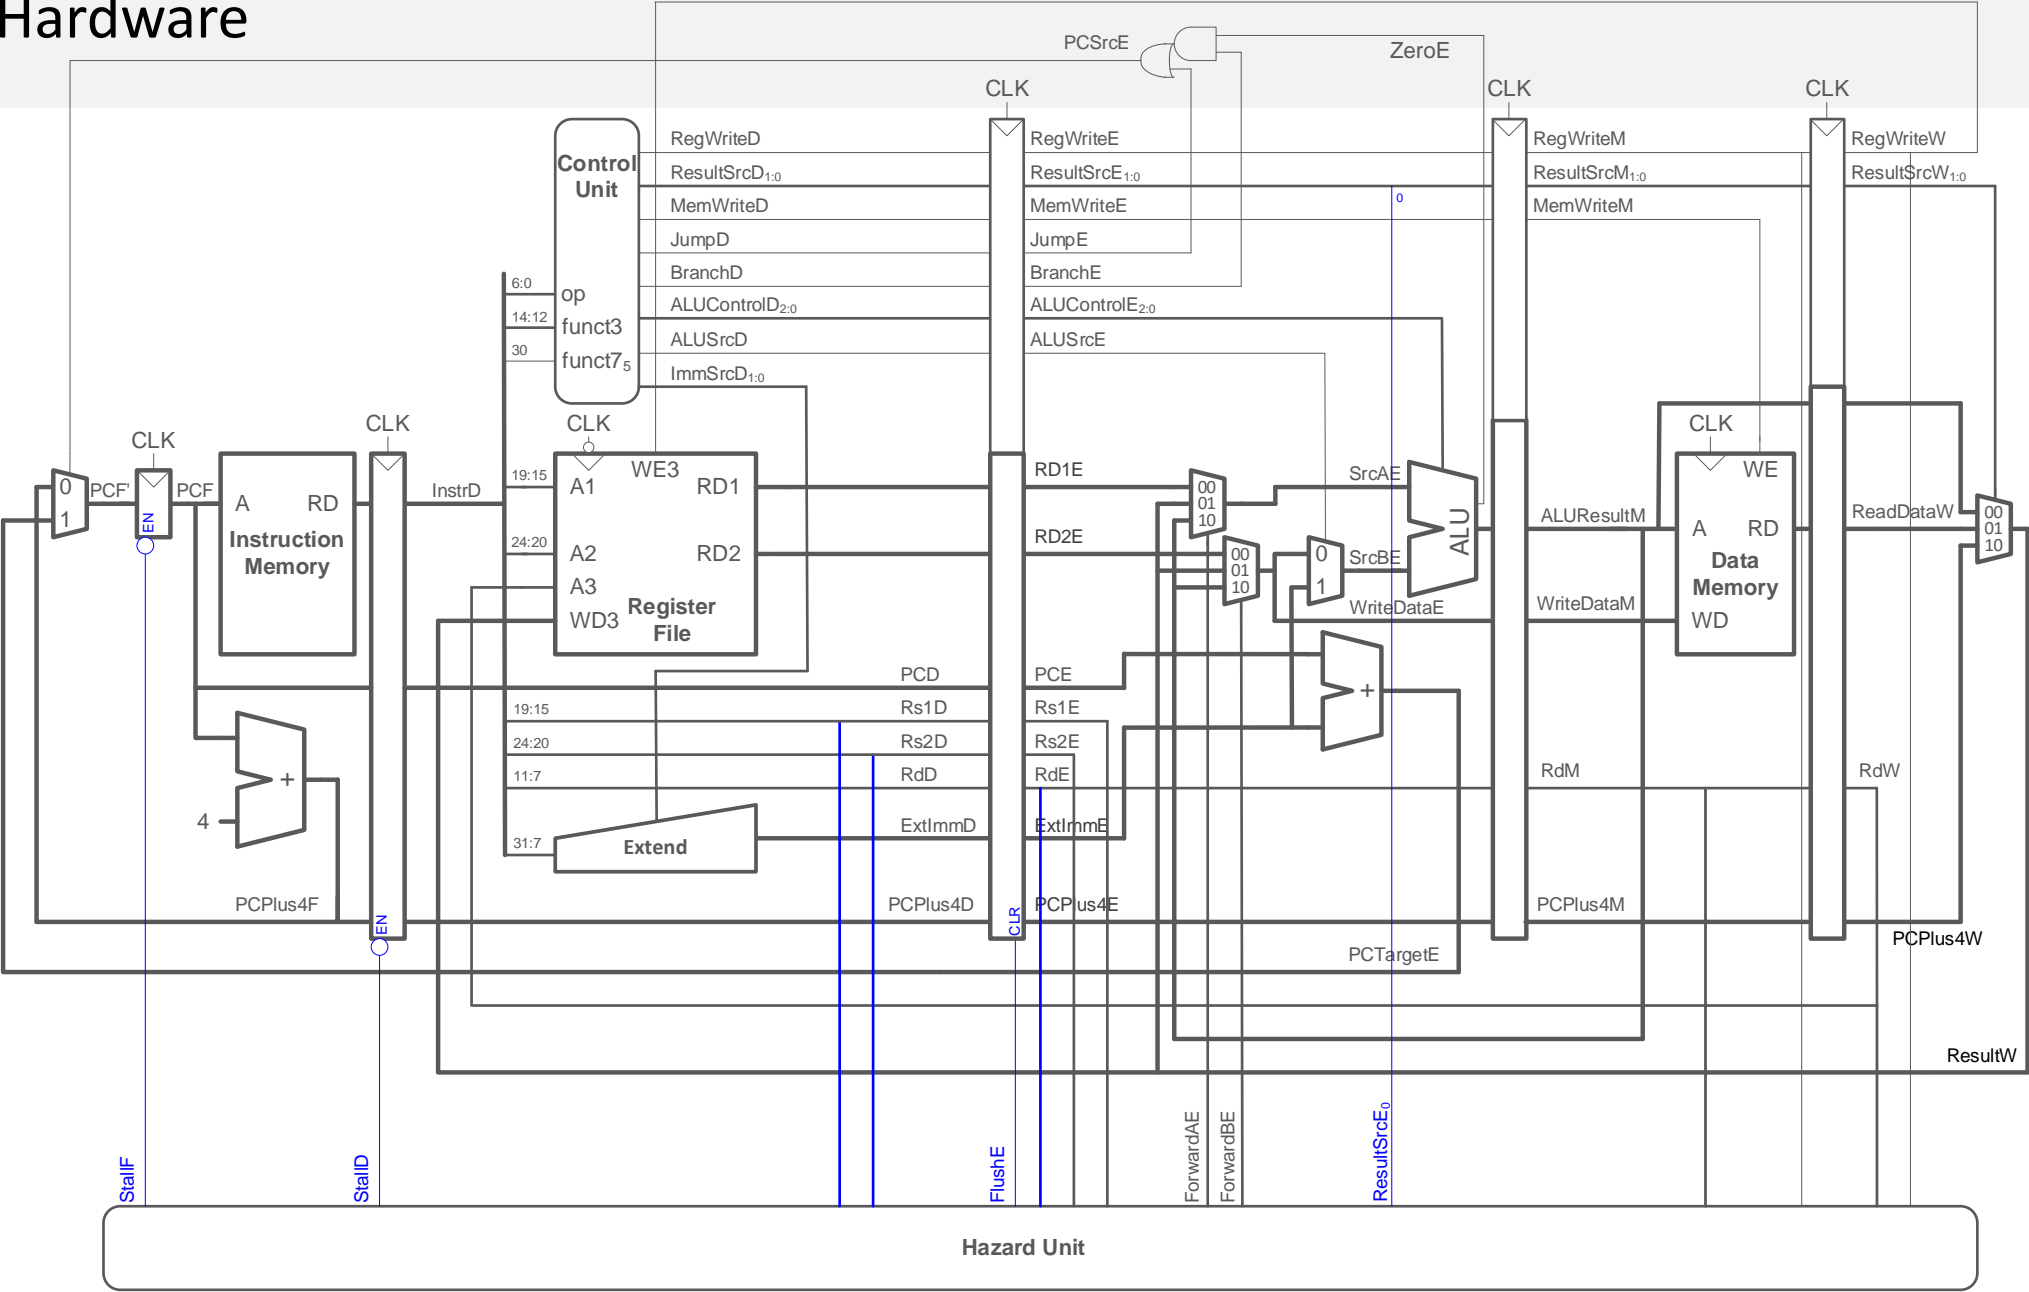
$lwStall = ((Rs1D == RdE) \text{ OR } (Rs2D == RdE)) \text{ AND } ResultSrcE_0$

$StallF = StallD = FlushE = lwStall$

(Stall the Fetch and Decode stages, and flush the Execute stage.)



Stalling Hardware



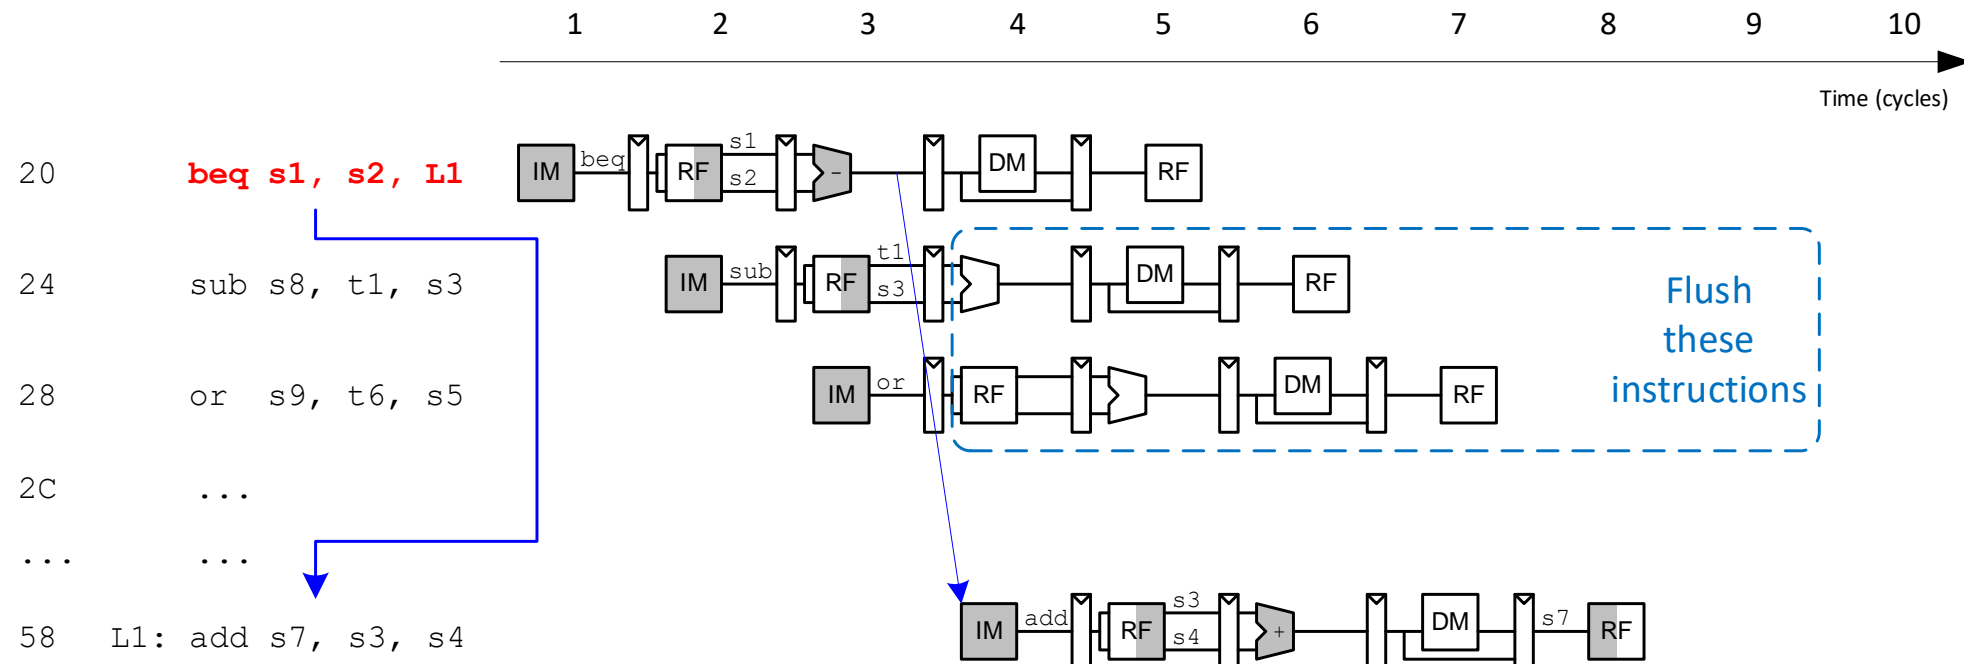
Pipelined Processor Control Hazards

DDCA Ch7 - Part 15: Pipelined Processor Control Hazards <https://www.youtube.com/watch?v=VcnwVxD4LAc>

Control Hazards

- **beq:**

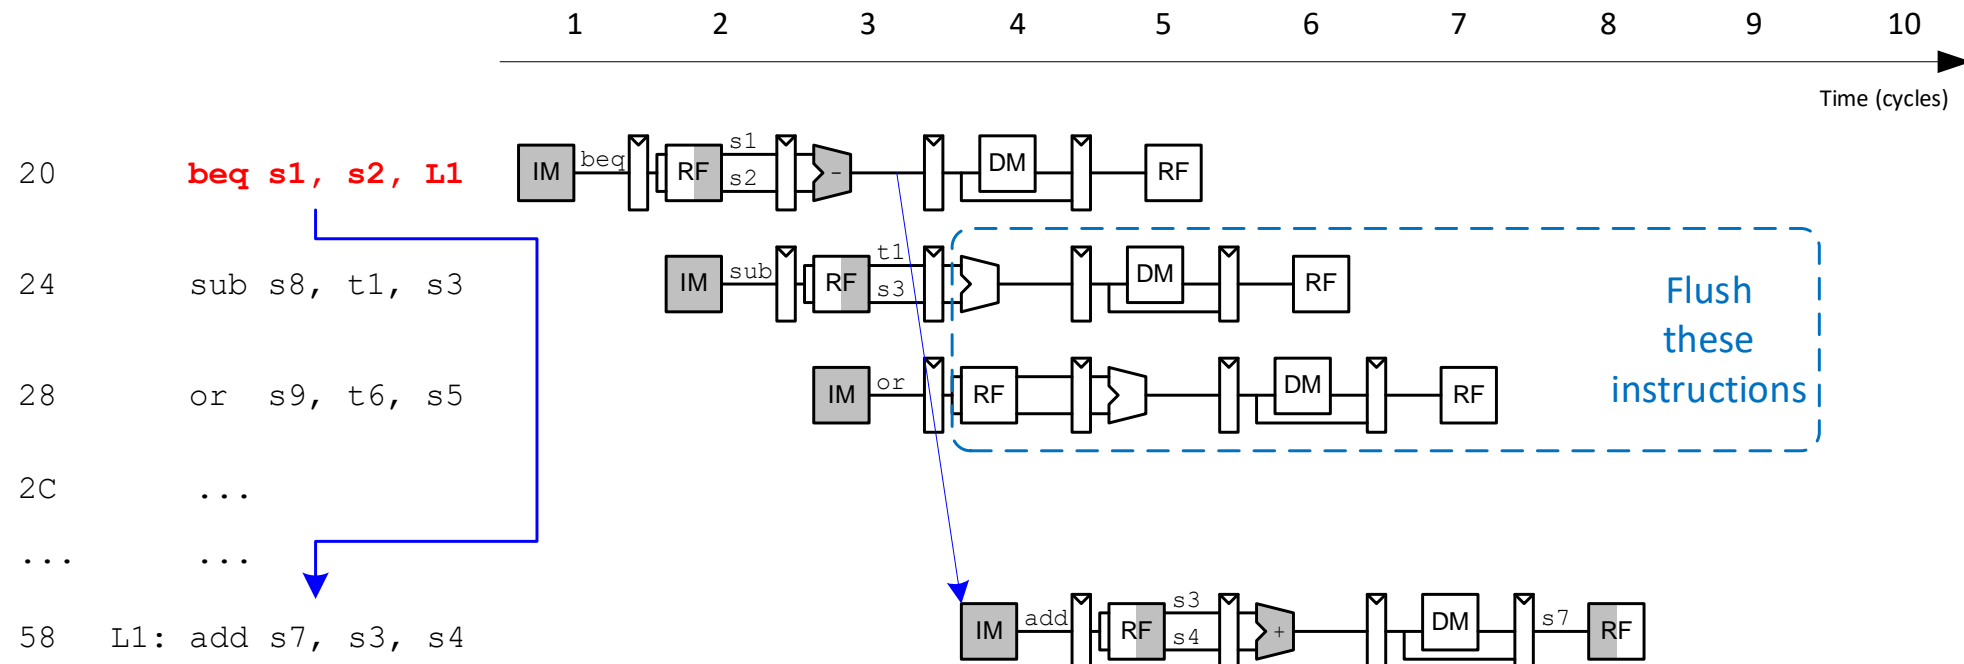
- Branch **not determined** until the **Execute stage** of pipeline
- **Instructions** after branch **fetched** before branch occurs
- These **2 instructions must be flushed** if branch happens



Control Hazards

Branch misprediction penalty:

- The number of instructions flushed when a branch is taken (in this case, 2 instructions)



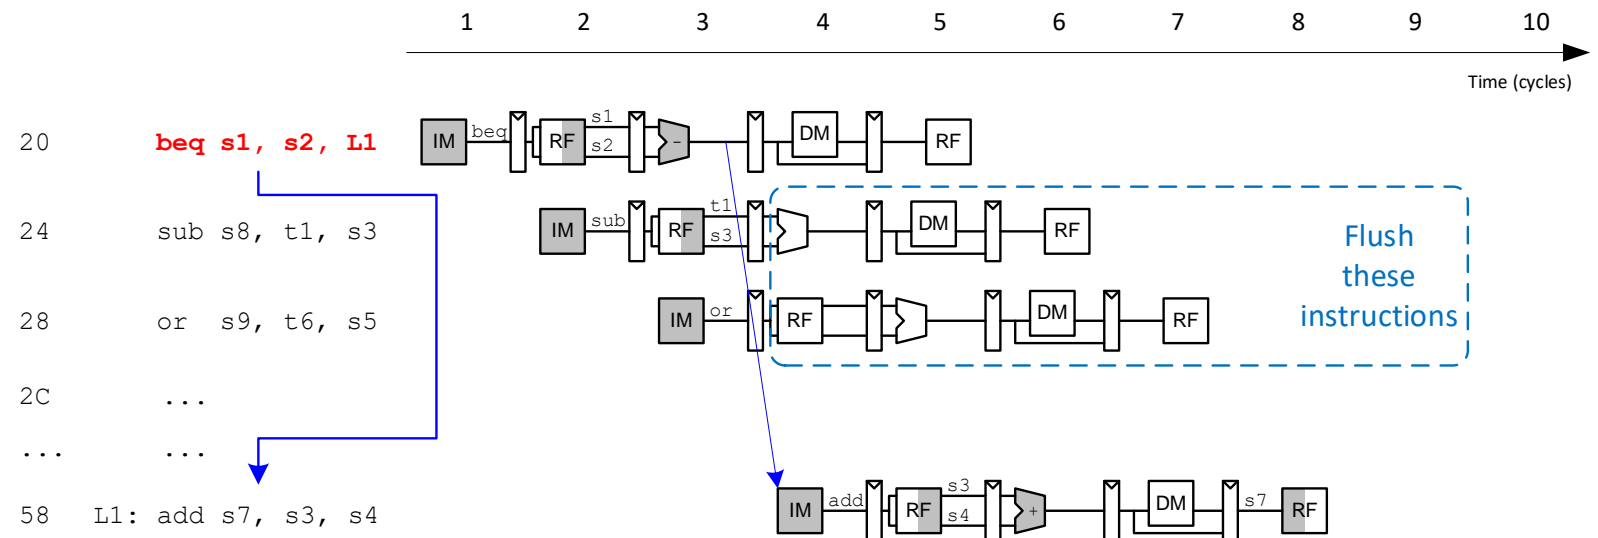
Control Hazards: Flushing Logic

- If branch is taken in execute stage, need to flush the instructions in the Fetch and Decode stages
 - Do this by clearing Decode and Execute Pipeline registers using *FlushD* and *FlushE*

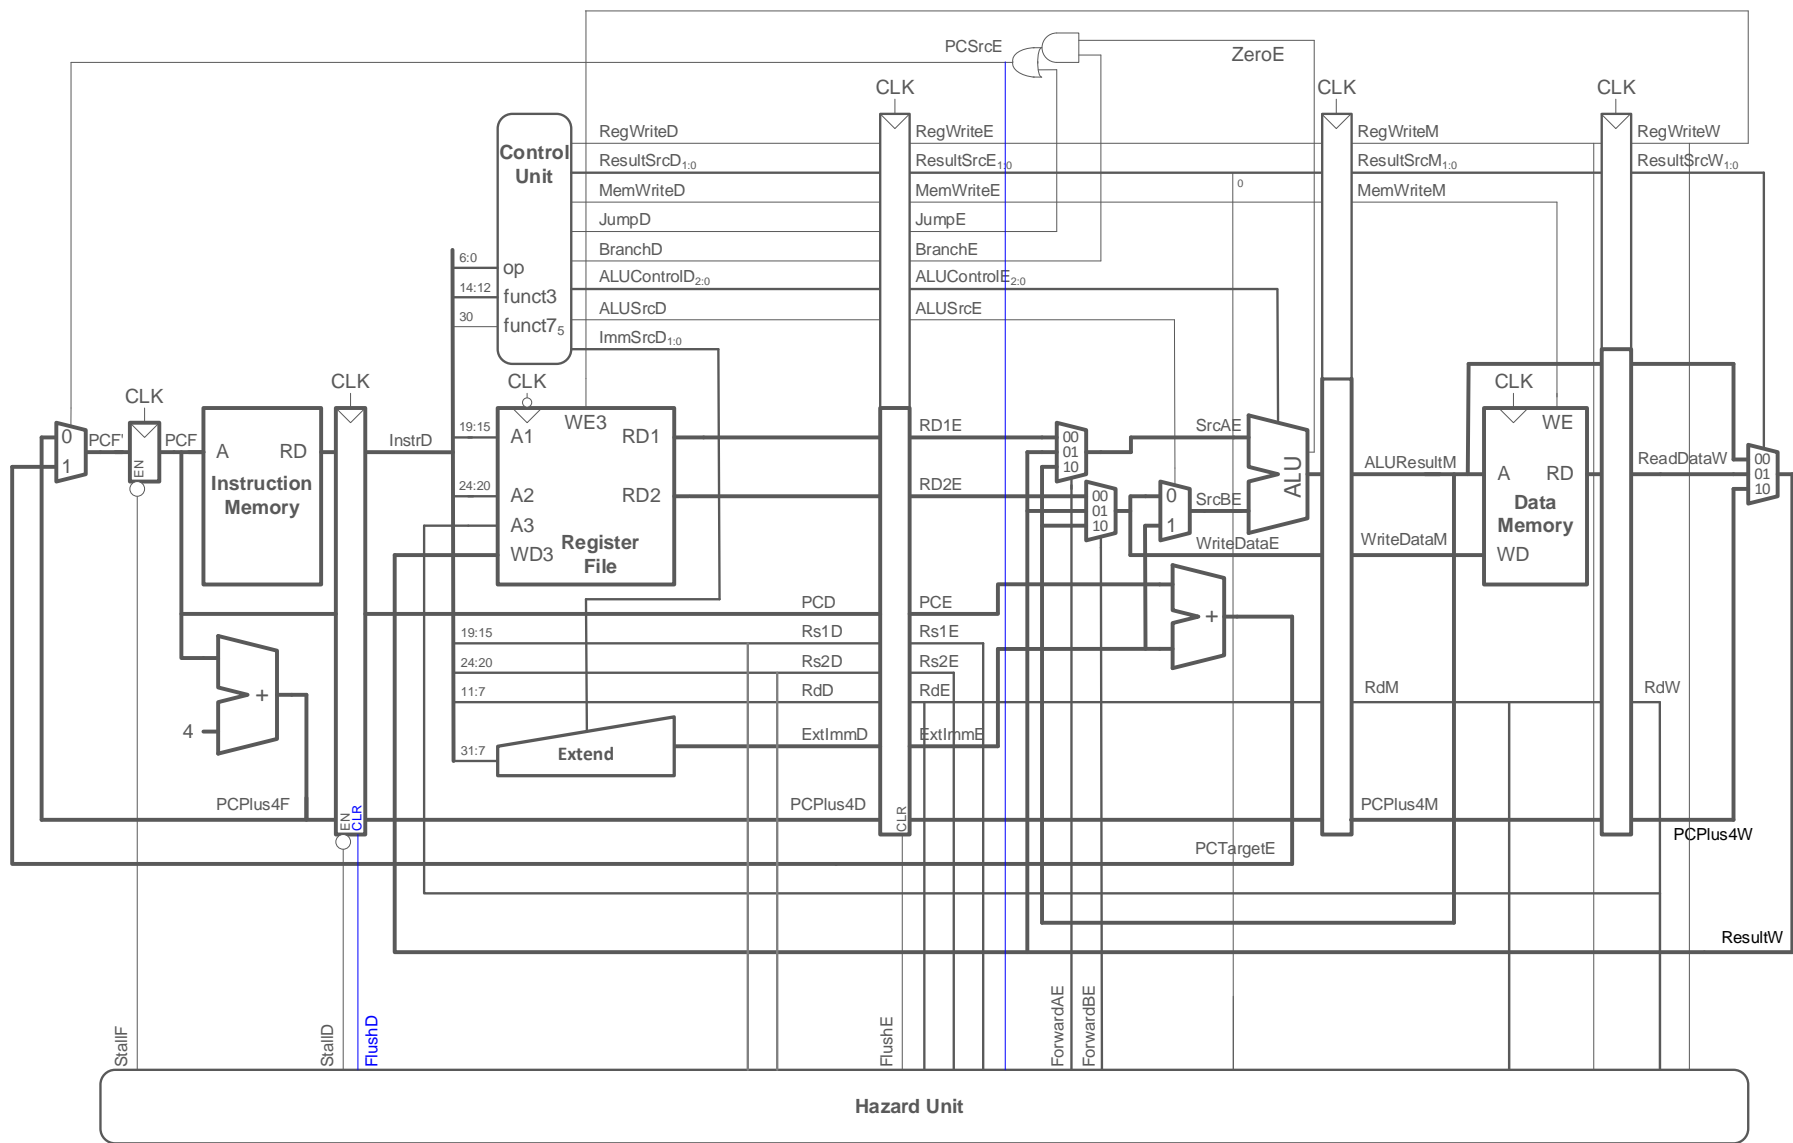
- **Equations:**

$$FlushD = PCSrcE$$

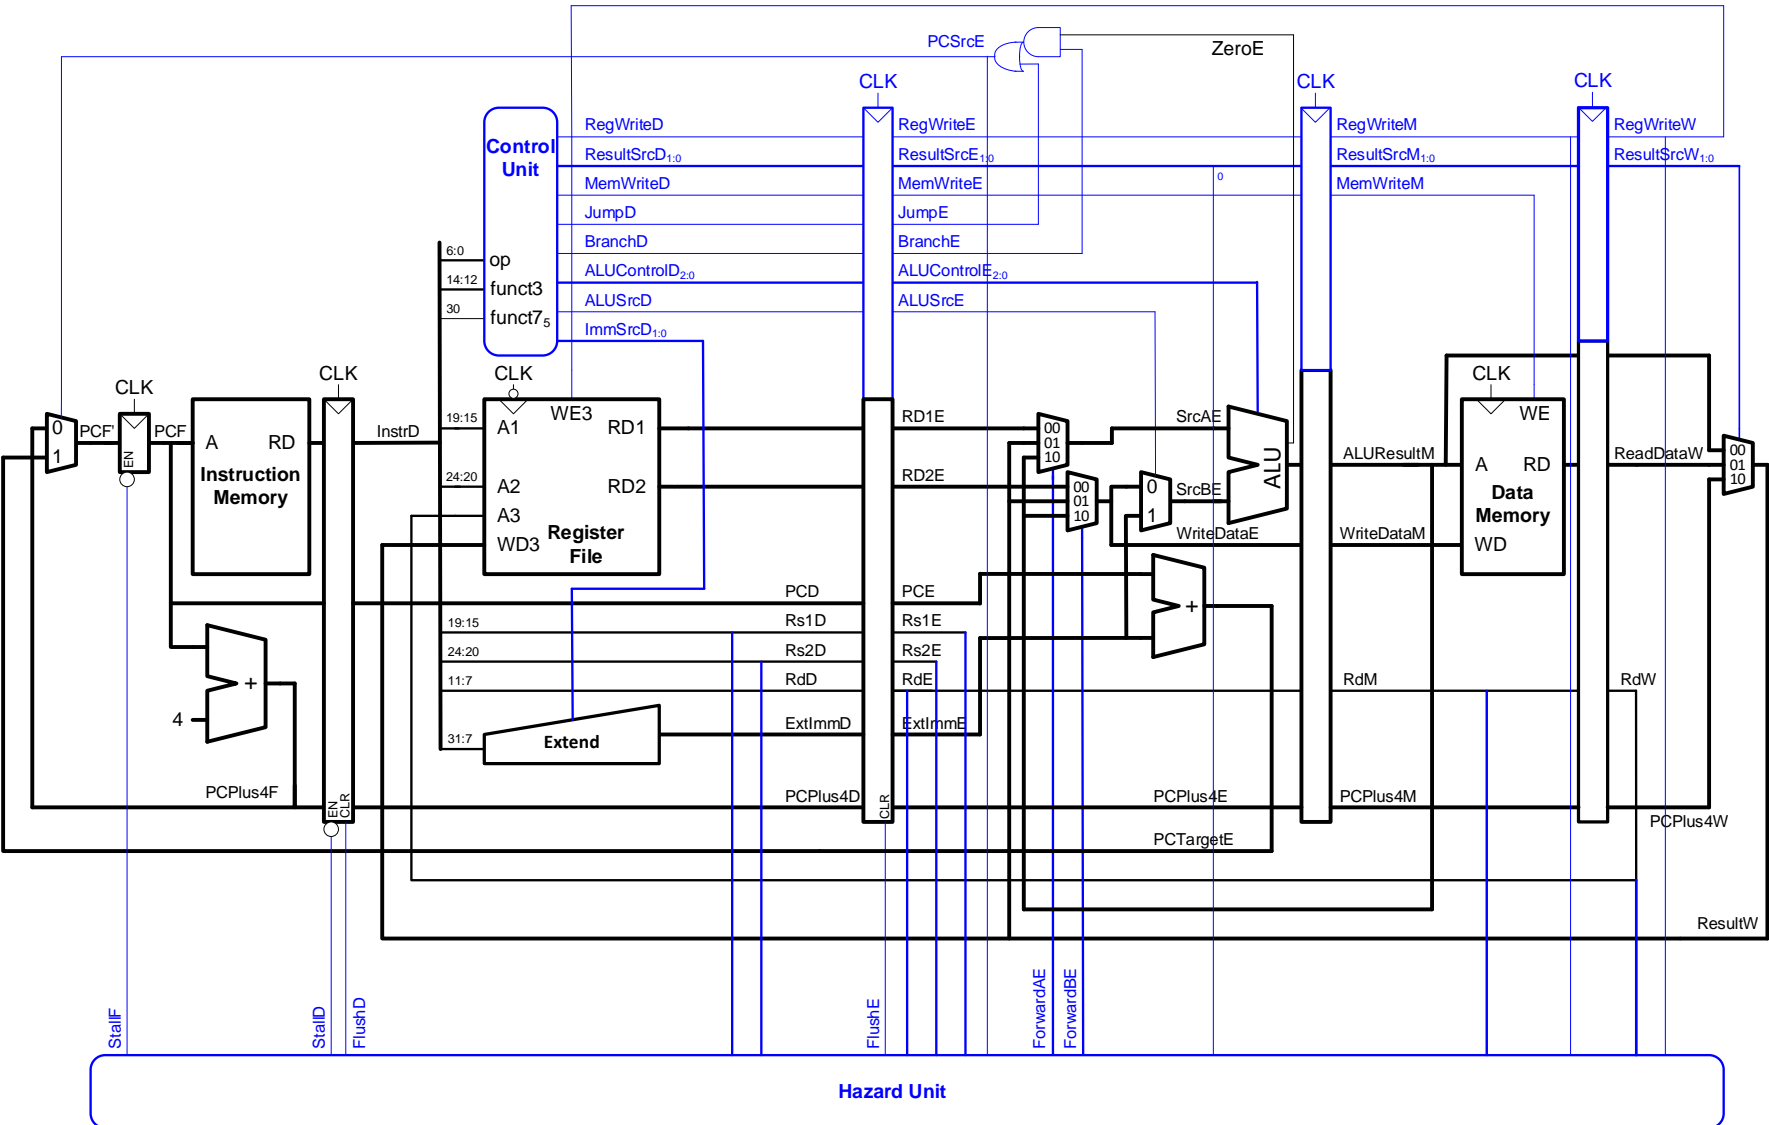
$$FlushE = lwStall \text{ OR } PCSrcE$$



Control Hazards: Flushing Hardware



RISC-V Pipelined Processor with Hazard Unit



Summary of Hazard Logic

Data hazard logic (shown for SrcA of ALU):

```
if ((Rs1E == RdM) AND RegWriteM) AND (Rs1E != 0) // Case 1
```

ForwardAE = 10

```
else if ((Rs1E == RdW) AND RegWriteW) AND (Rs1E != 0) // Case 2
```

***ForwardAE* = 01**

```
else ForwardAE = 00 // Case 3
```

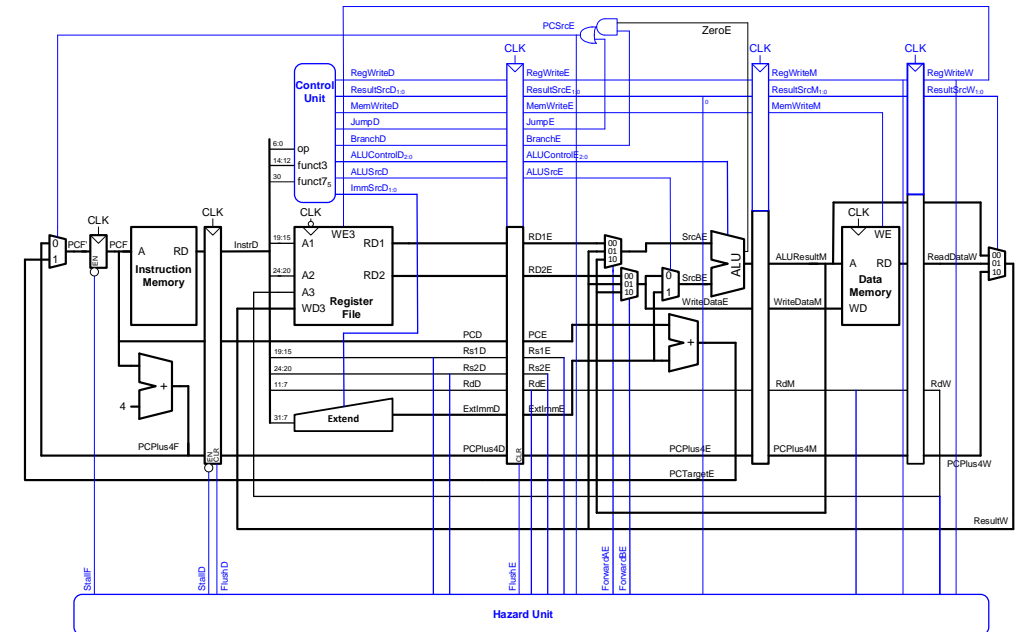
Load word stall logic:

$$lwStall = ((Rs1D == RdE) \text{ OR } (Rs2D == RdE)) \text{ AND } ResultSrcE_0$$
$$StallF = StallD = lwStall$$

Control hazard flush:

$$\mathbf{FlushD} = \mathbf{PCSrcE}$$

FlushE = lwStall OR PCSrcE



Pipelined Performance

DDCA Ch7 - Part 14: Pipelined Processor Data Hazards <https://www.youtube.com/watch?v=zuegcg6ZSFQ>

Pipelined Processor Performance Example

- **SPECINT2000 benchmark:**

- 25% loads
- 10% stores
- 13% branches
- 52% R-type

- **Suppose:**

- 40% of loads used by next instruction
- 50% of branches mispredicted

- **What is the average CPI?**

(Ideally it's 1, but...)

- Load CPI = 1 when not stalling, 2 when stalling
- Branch CPI = 1 when not stalling, 3 when stalling

$$\rightarrow \text{So, } \text{CPI}_{\text{lw}} = 1(0.6) + 2(0.4) = 1.4$$

$$\rightarrow \text{So, } \text{CPI}_{\text{beq}} = 1(0.5) + 3(0.5) = 2$$

- **Average CPI** = $(0.25)(1.4) + (0.1)(1) + (0.13)(2) + (0.52)(1) = \mathbf{1.23}$

Pipelined Processor Performance Example

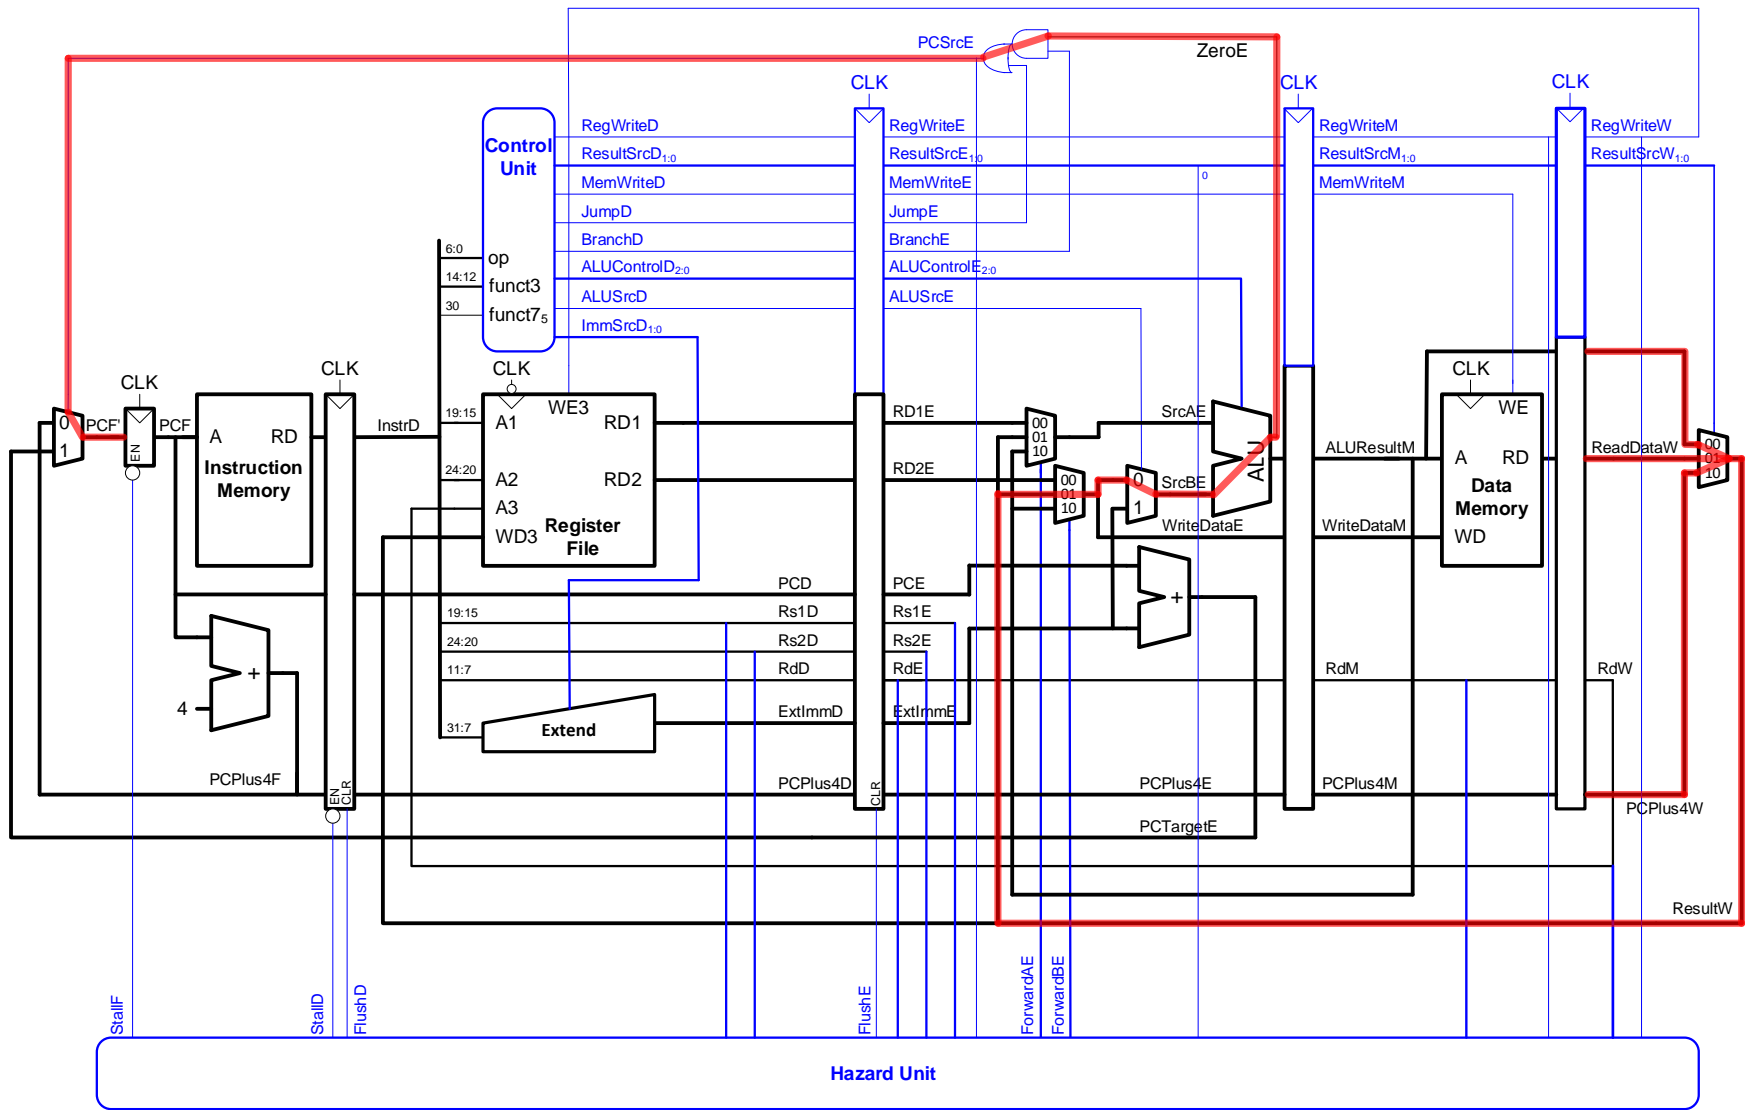
- Pipelined processor critical path:

$$T_{c_pipelined} = \max \left[\begin{array}{l} t_{pcq} + t_{mem} + t_{setup} \\ 2(t_{RFread} + t_{setup}) \\ t_{pcq} + 4t_{mux} + t_{ALU} + t_{AND-OR} + t_{setup} \\ t_{pcq} + t_{mem} + t_{setup} \\ 2(t_{pcq} + t_{mux} + t_{RFwrite}) \end{array} \right]$$

Fetch
Decode
Execute
Memory
Writeback

- Decode and Writeback stages **both use the register file** in each cycle
- So each stage gets half of the cycle time ($T_c/2$) to do their work
- Or, stated a different way, **2x of their work** must fit in a cycle (T_c)

Pipelined Critical Path: Execute Stage



Pipelined Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	t_{pcq_PC}	40
Register setup	t_{setup}	50
Multiplexer	t_{mux}	30
AND-OR gate	t_{AND-OR}	20
ALU	t_{ALU}	120
Decoder (Control Unit)	t_{dec}	25
Extend unit	t_{dec}	35
Memory read	t_{mem}	200
Register file read	t_{RFread}	100
Register file setup	$t_{RFsetup}$	60

$$\begin{aligned} T_{c_pipelined} &= t_{pcq} + 4t_{mux} + t_{ALU} + t_{AND-OR} + t_{setup} \\ &= (40 + 4*30 + 120 + 20 + 50) \text{ ps} = \mathbf{350 \text{ ps}} \end{aligned}$$

Pipelined Performance Example

Program with 100 billion instructions

$$\begin{aligned}\text{Execution Time} &= (\# \text{ instructions}) \times \text{CPI} \times T_c \\ &= (100 \times 10^9)(1.23)(350 \times 10^{-12}) \\ &= 43 \text{ seconds}\end{aligned}$$

Processor Performance Comparison

Processor	Execution Time (seconds)	Speedup (single-cycle as baseline)
Single-cycle	75	1
Multi-cycle	155	0.5
Pipelined	43	1.7